

AI made easy with Apache Flink + AI Flow

Jiangjie (Becket) Qin

Staff Software Engineer & Senior Manager

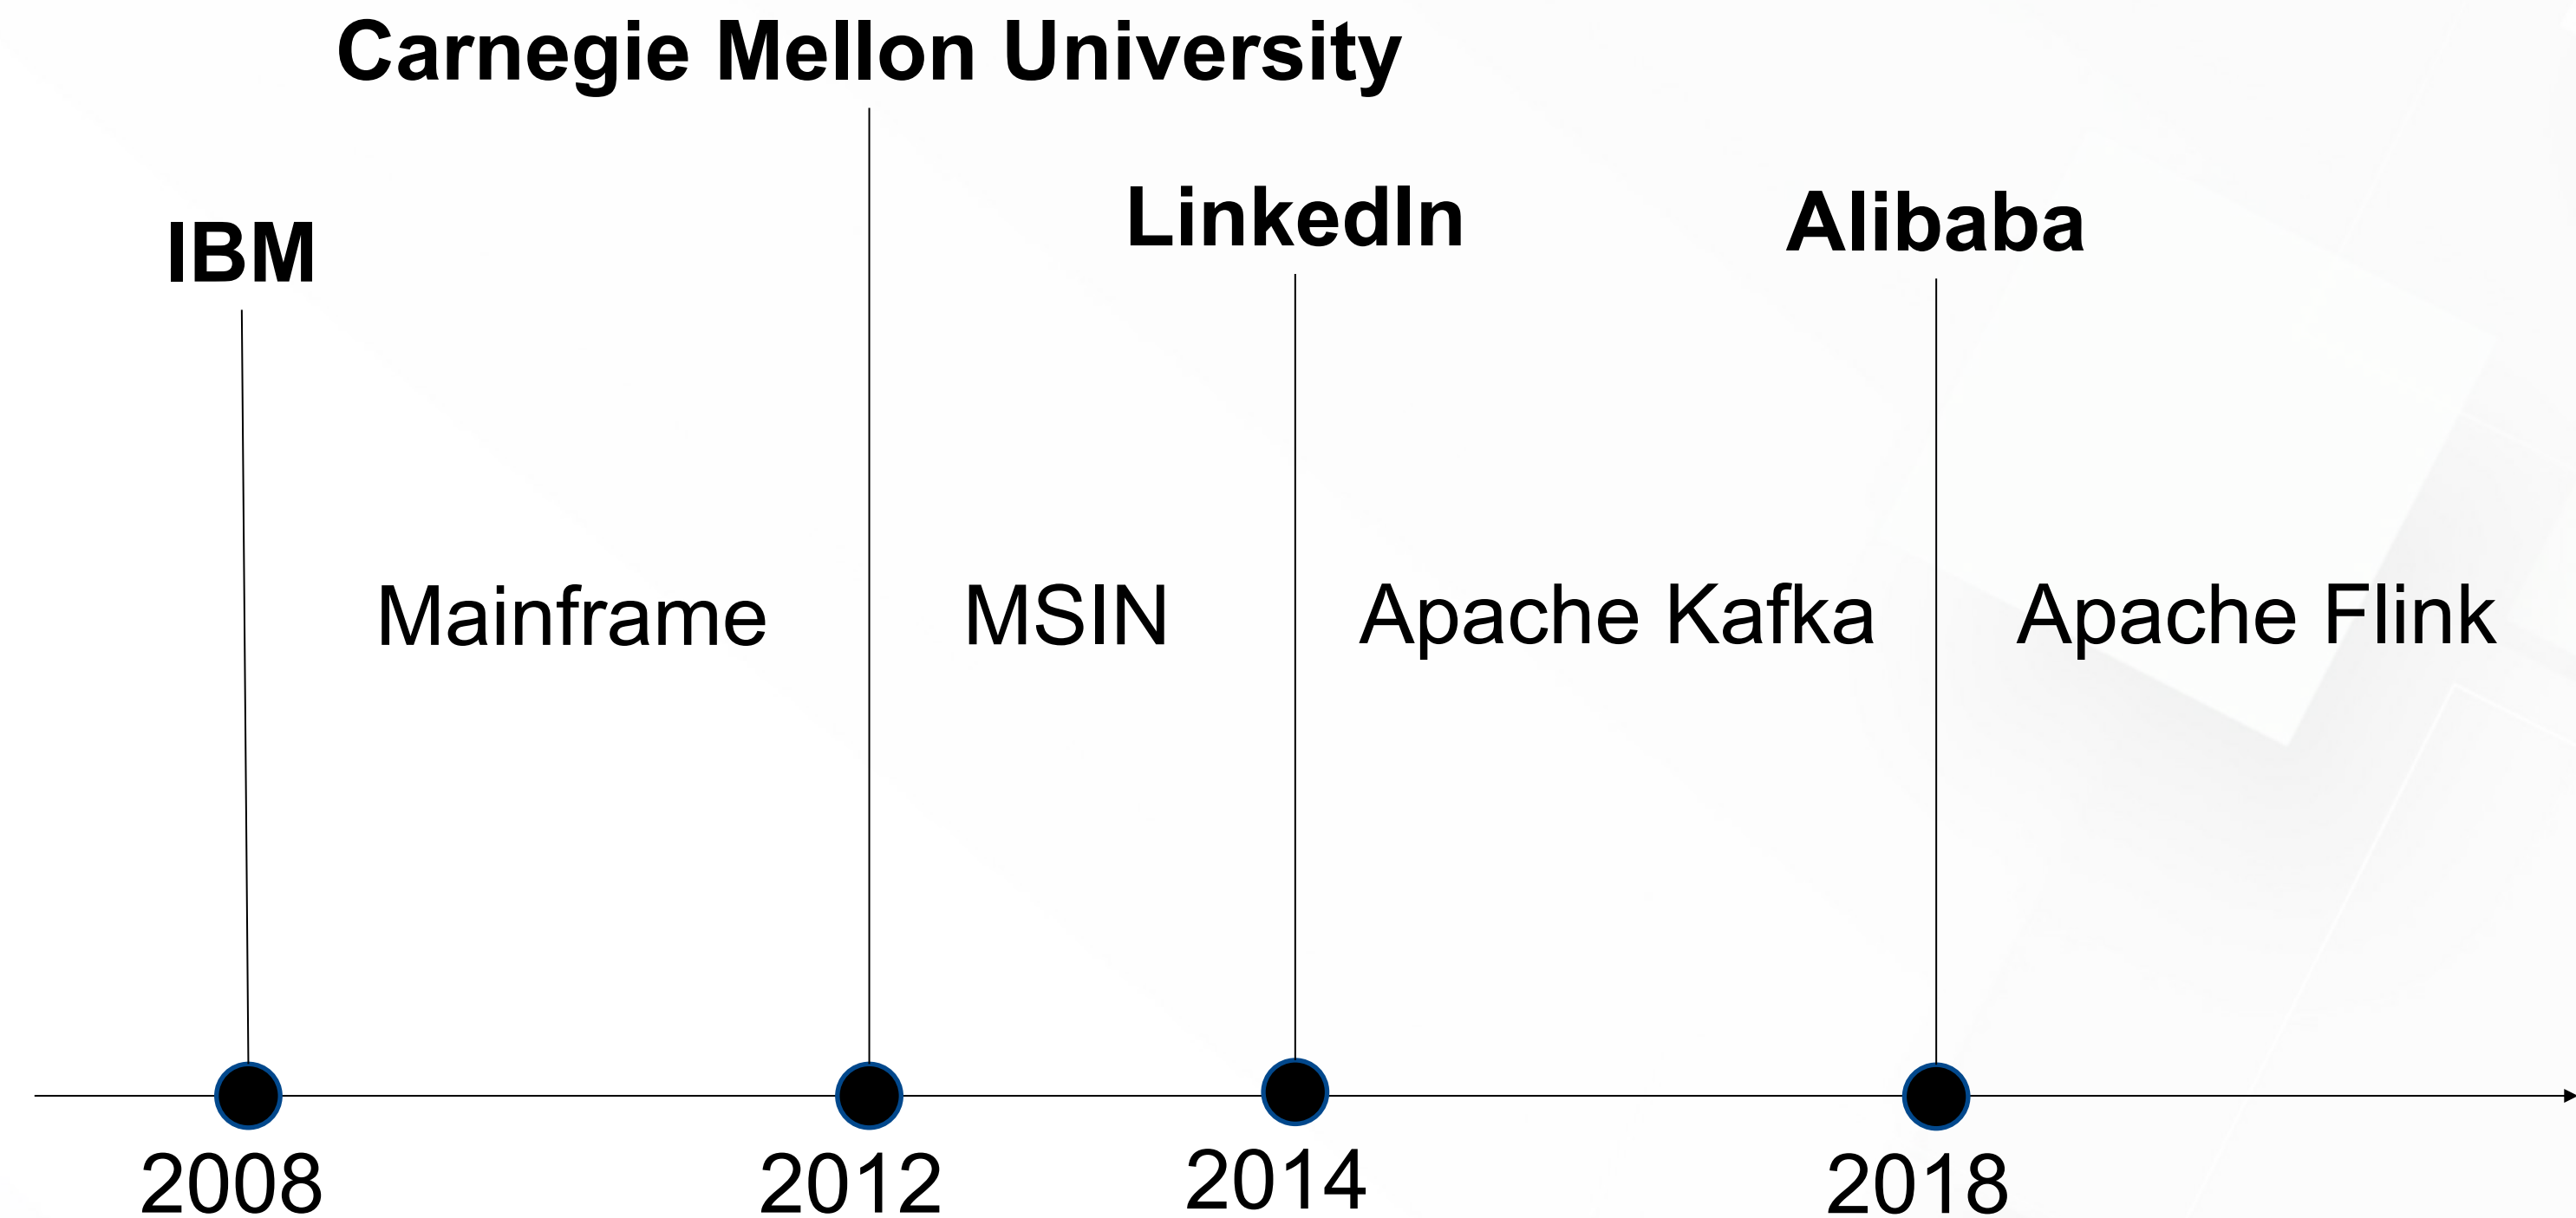
24/04/2020

About Me



Alibaba Inc.
Staff Software Engineer
Senior Manager

Apache Flink PMC Member
Apache Kafka PMC Member



Agenda

- What is AI Flow
- AI Flow Introduction
- Flink AI Flow
- Summary

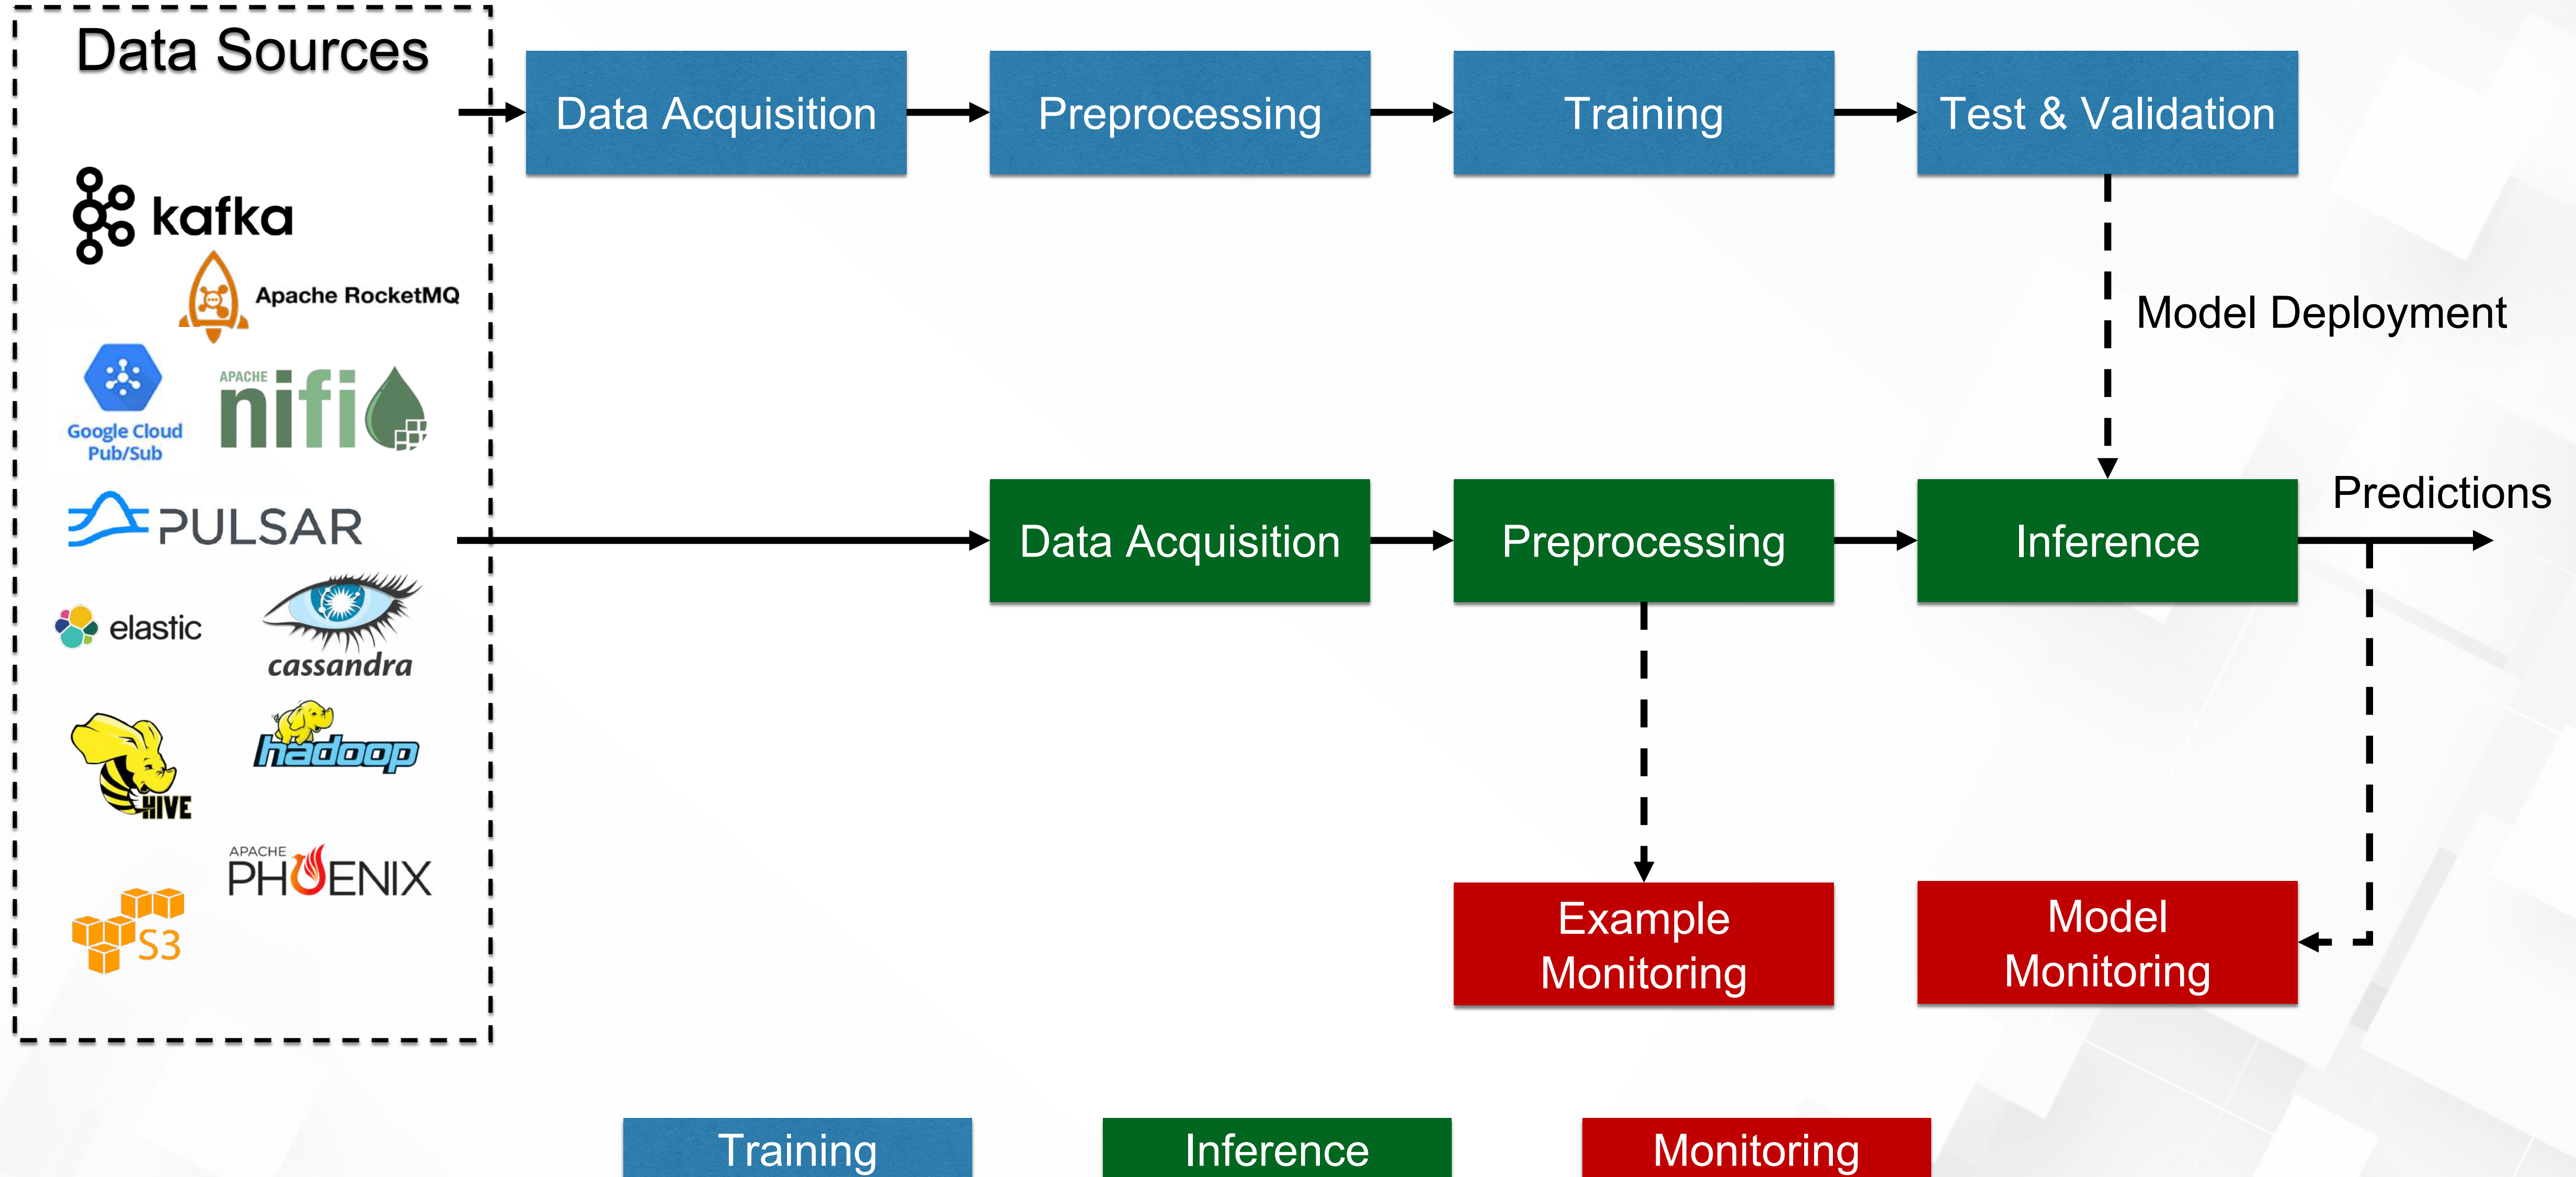


What is AI Flow

- An overview of AI workflow
- Challenges in AI workflow
- Meet AI Flow




An Overview of AI Workflow



Understand a Workflow

Workflow
Definition



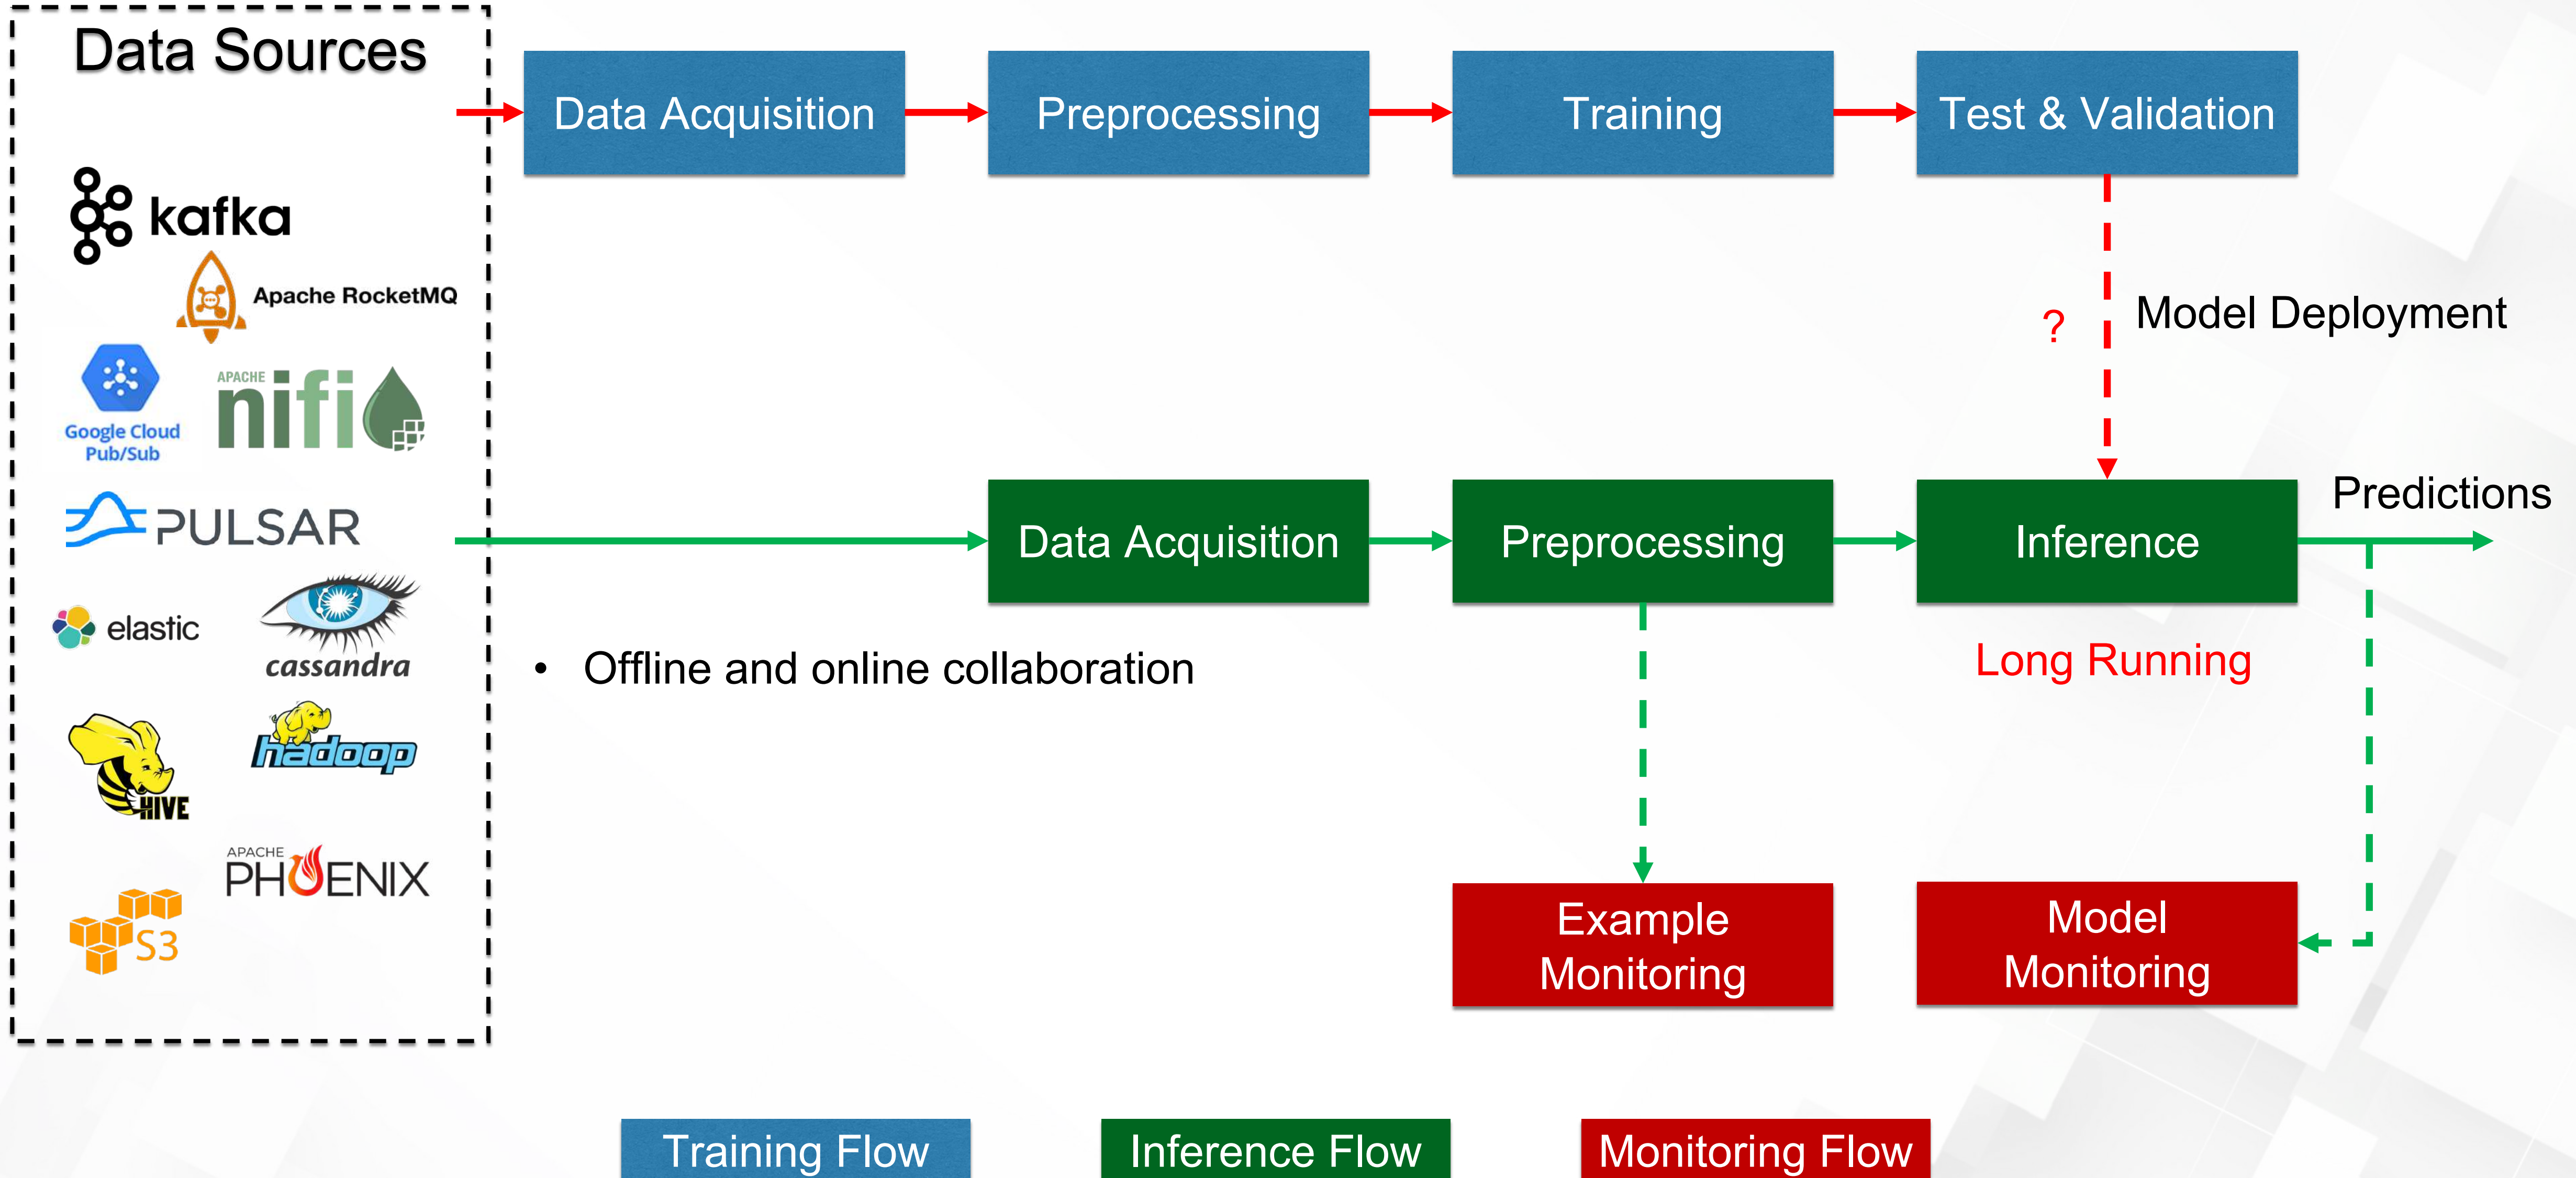
Workflow
Execution

Scheduler

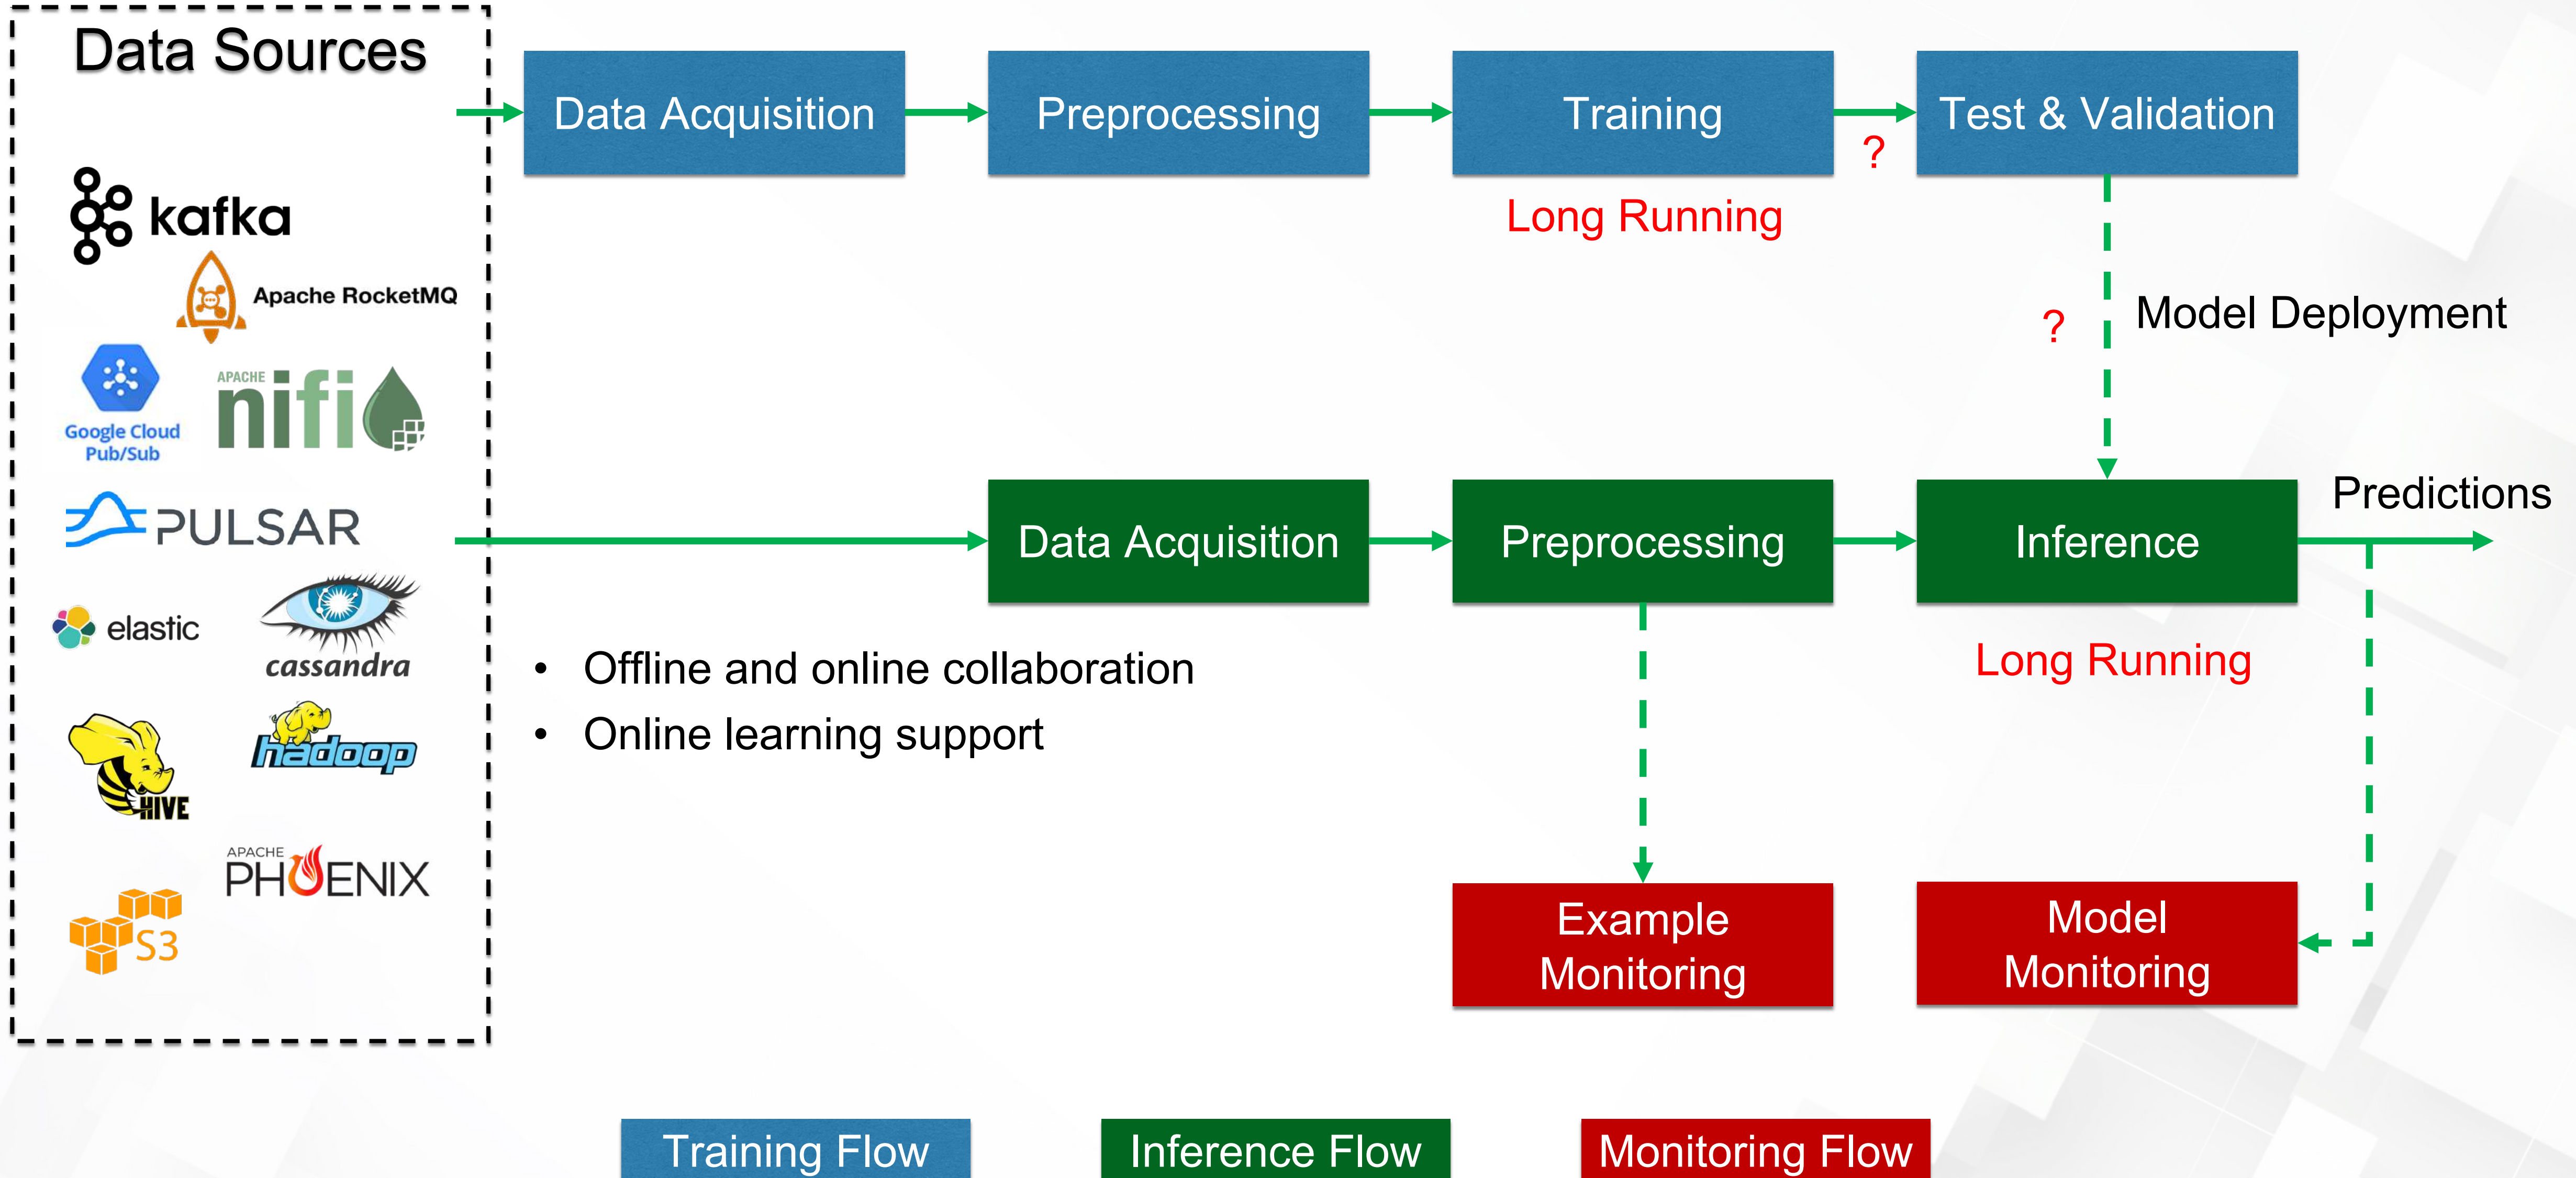


Job status based scheduling

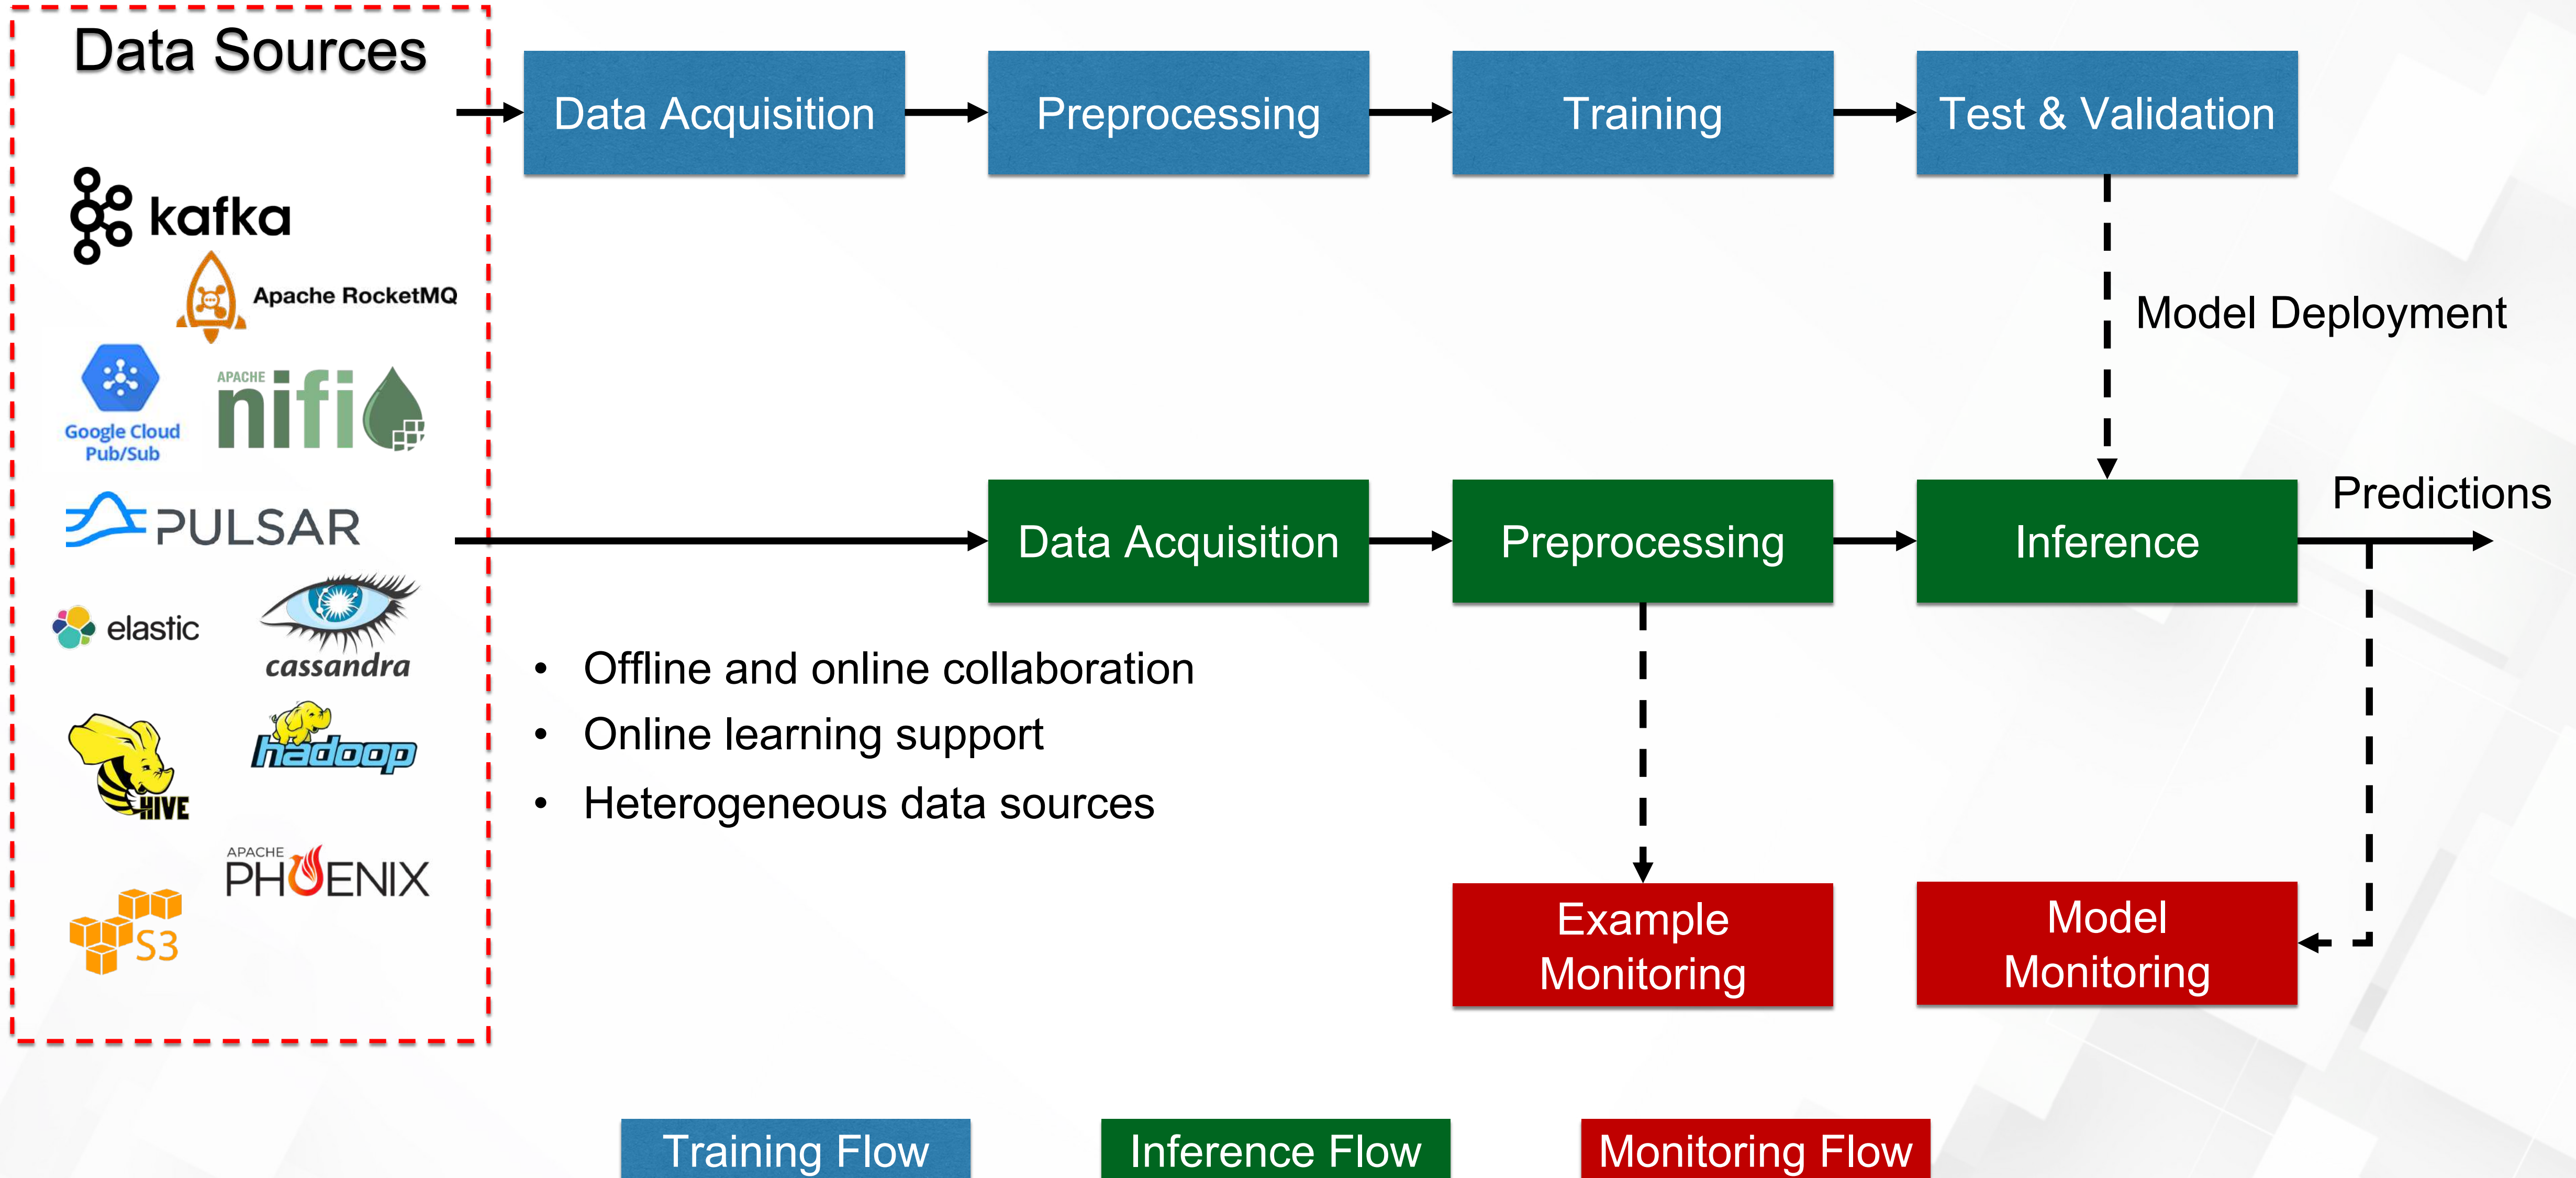
Challenges in AI Workflow



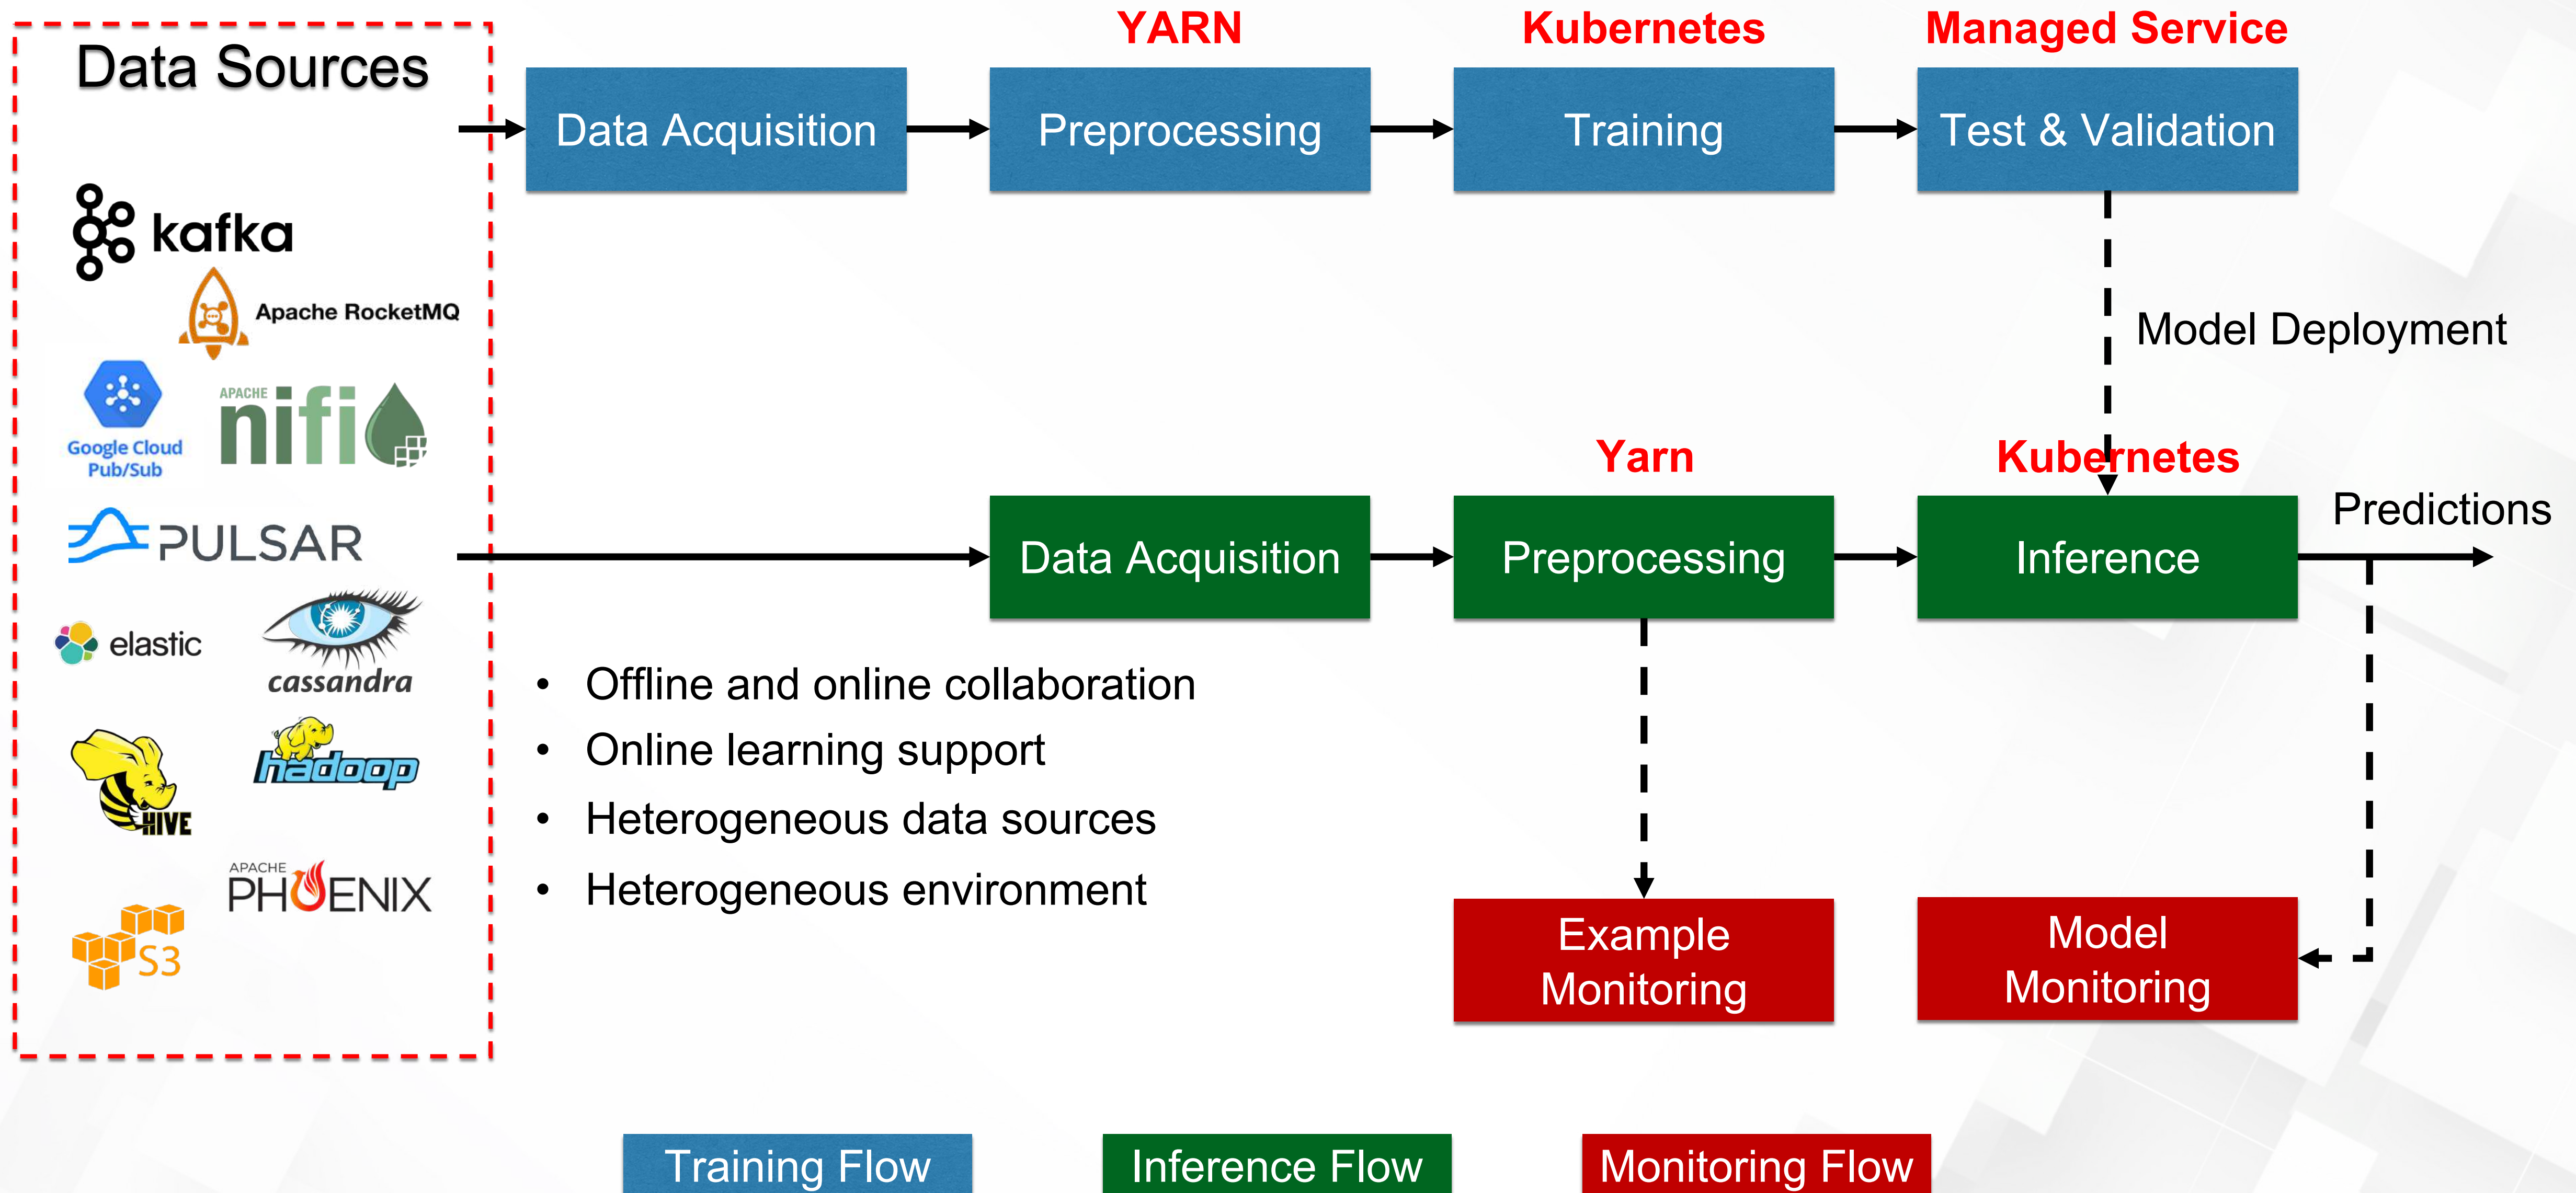
Challenges in AI Workflow



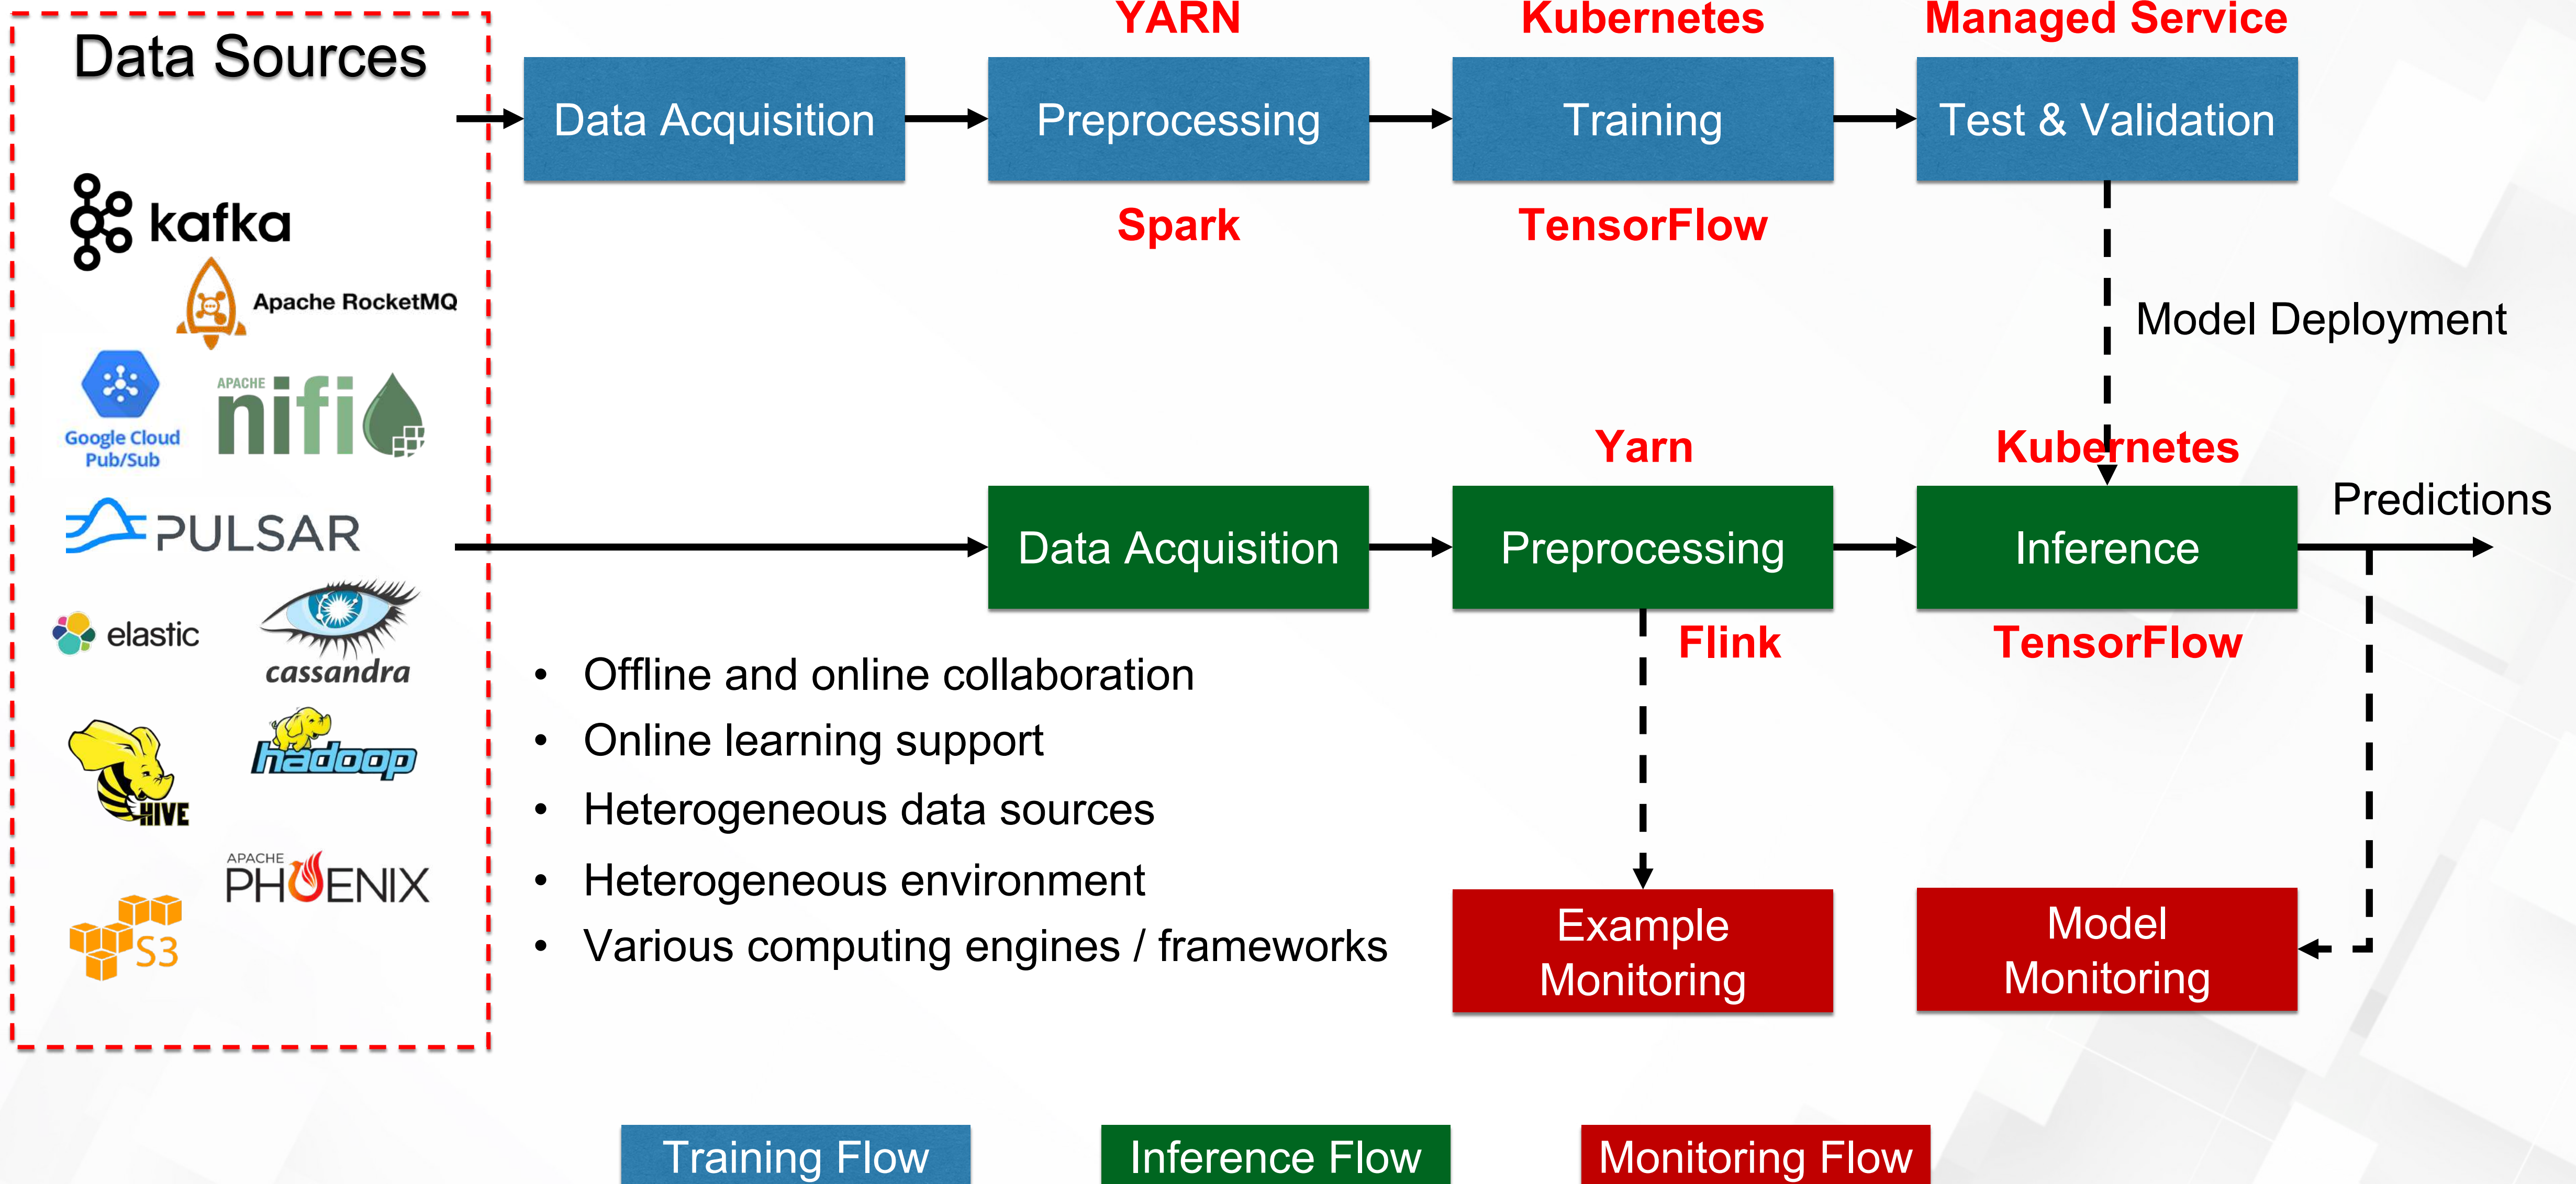
Challenges in AI workflow



Challenges in AI workflow



Challenges in AI workflow



Meet AI Flow

- Native support for real-time AI
 - Define workflows with streaming jobs
- Adaptive to various engines
 - Simple Python programs, Flink, Spark, TF, PyTorch, etc.
- Platform agnostic
 - Works with K8S, YARN, Cloud Services, etc

AI Flow Introduction


- Brief Introduction
- Architecture Overview
- Design Walkthrough
- Supporting Services



Brief Introduction

- A top level workflow abstraction for big data and AI (esp. real-time AI)

Workflow
Definition



Workflow
Execution

Scheduler




Event based scheduling


Brief Introduction

- A few supporting services helping fulfill the semantics
 - Metadata Service
 - Notification Service
 - Model Center

Metadata Service




Dataset



Project



Workflow & jobs



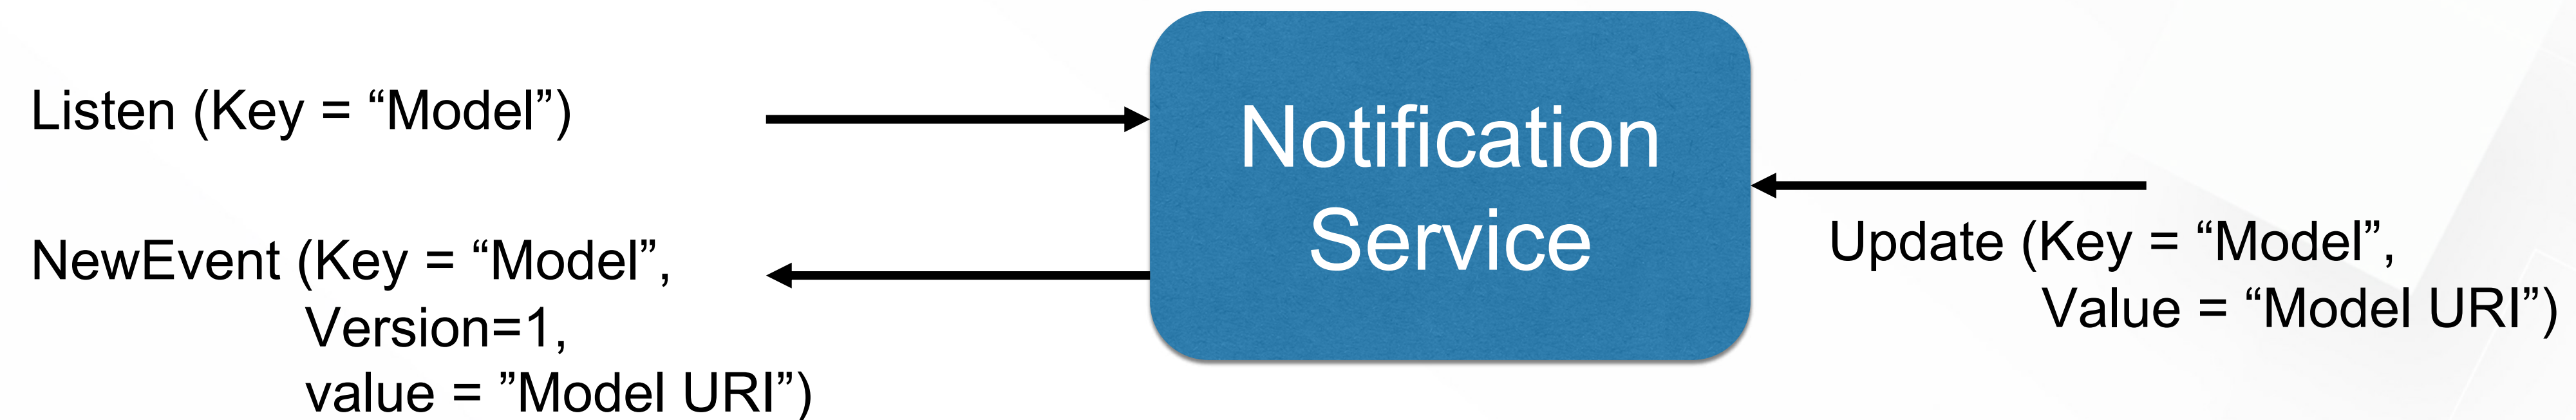
Model Lineage



Artifacts

Notification Service

- Keyed events and event listeners




Model Center




Versioning



Parameters



Metrics



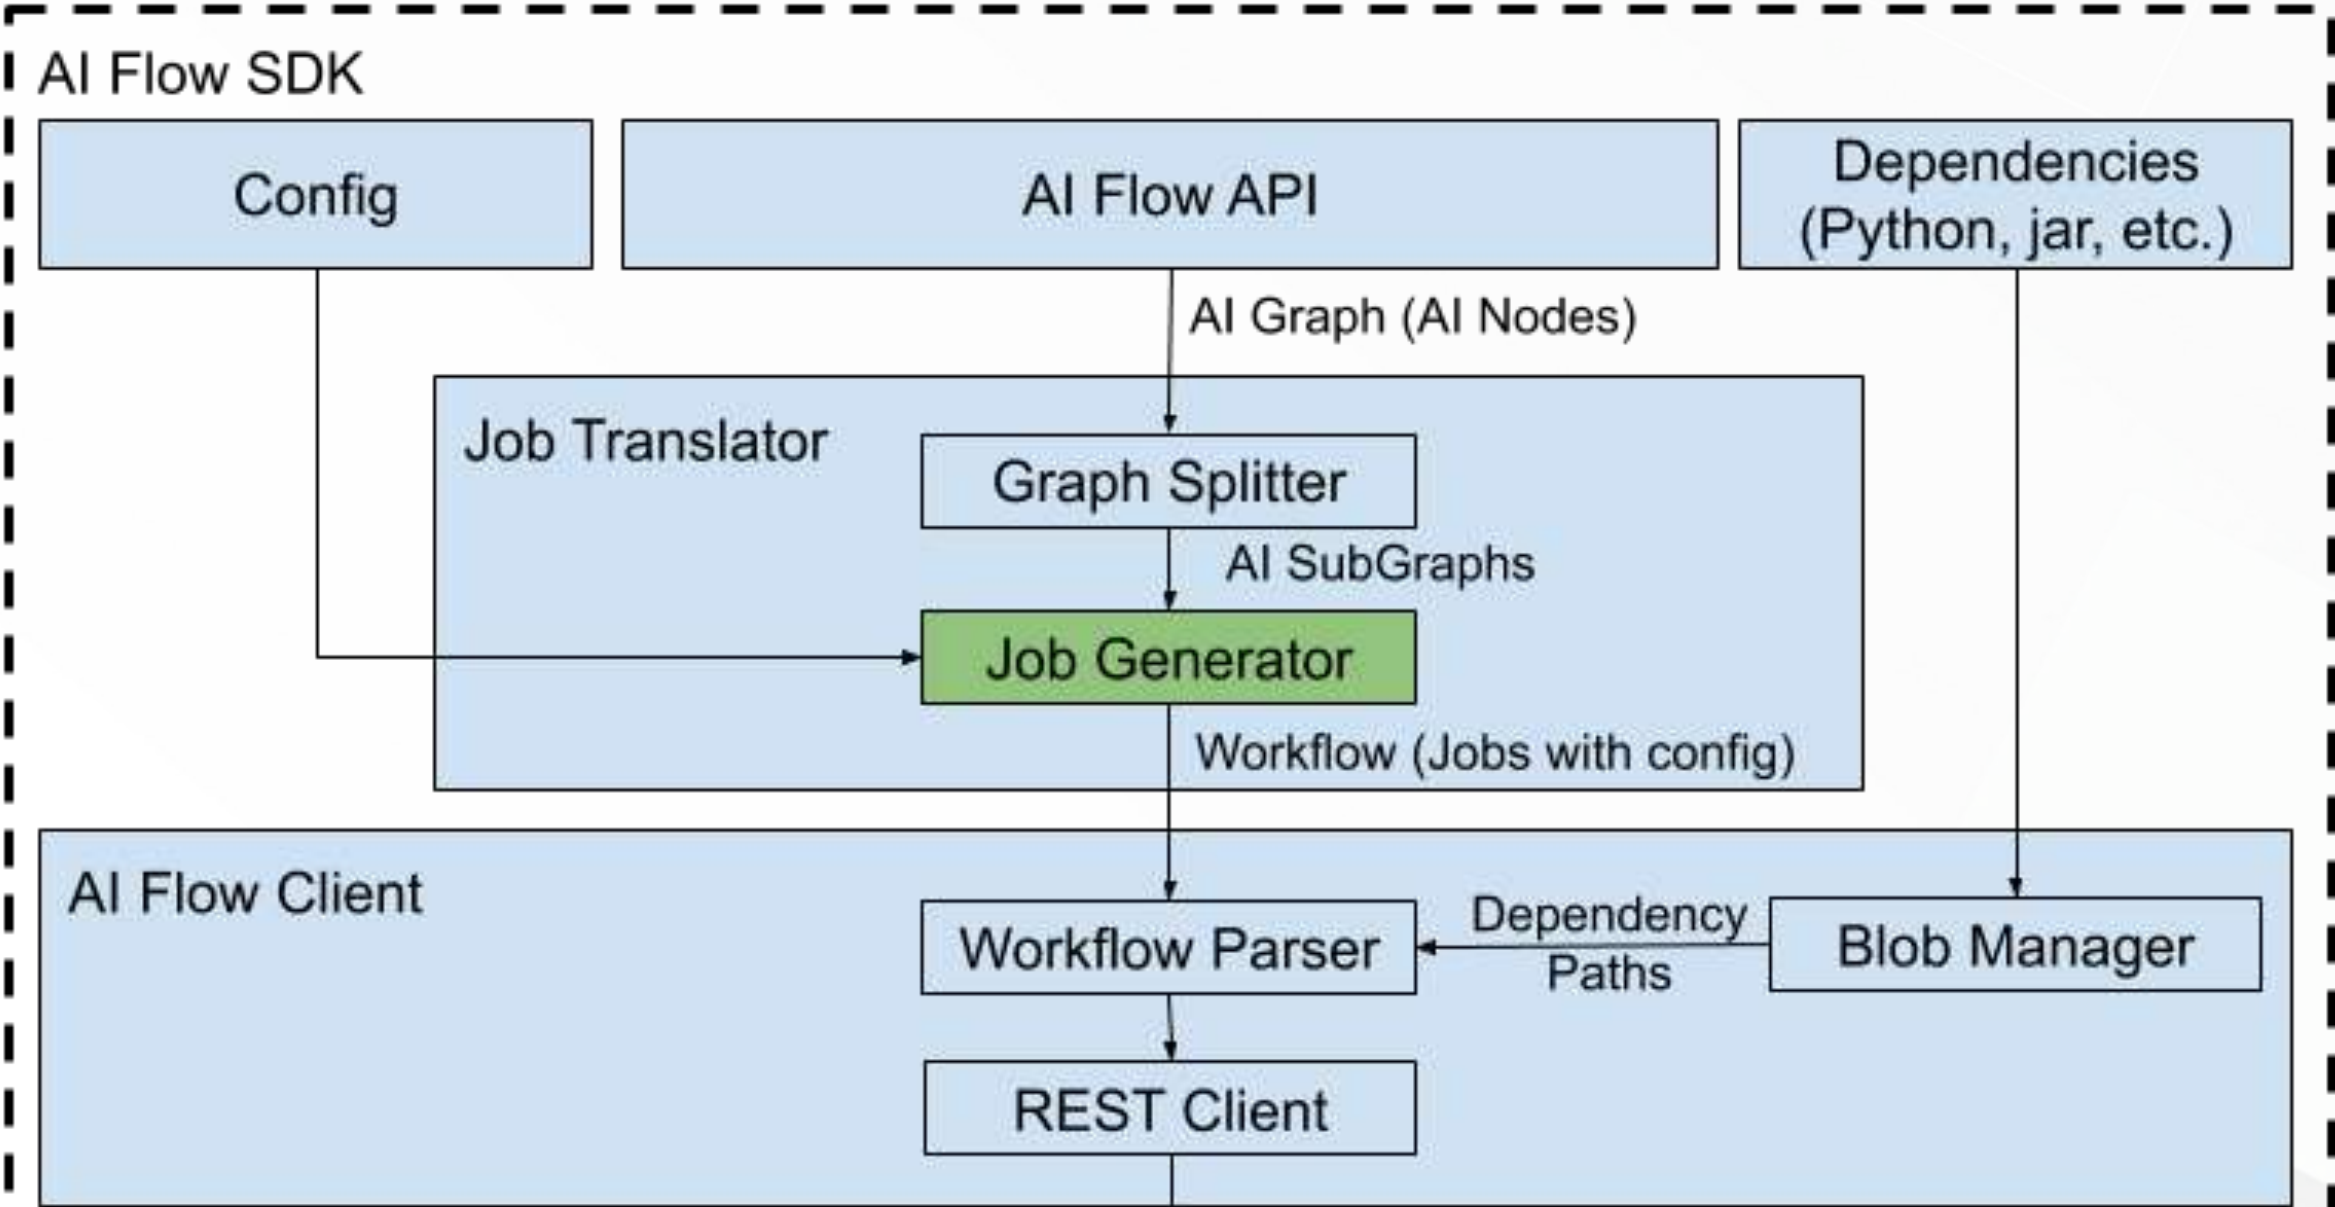
Lifecycle



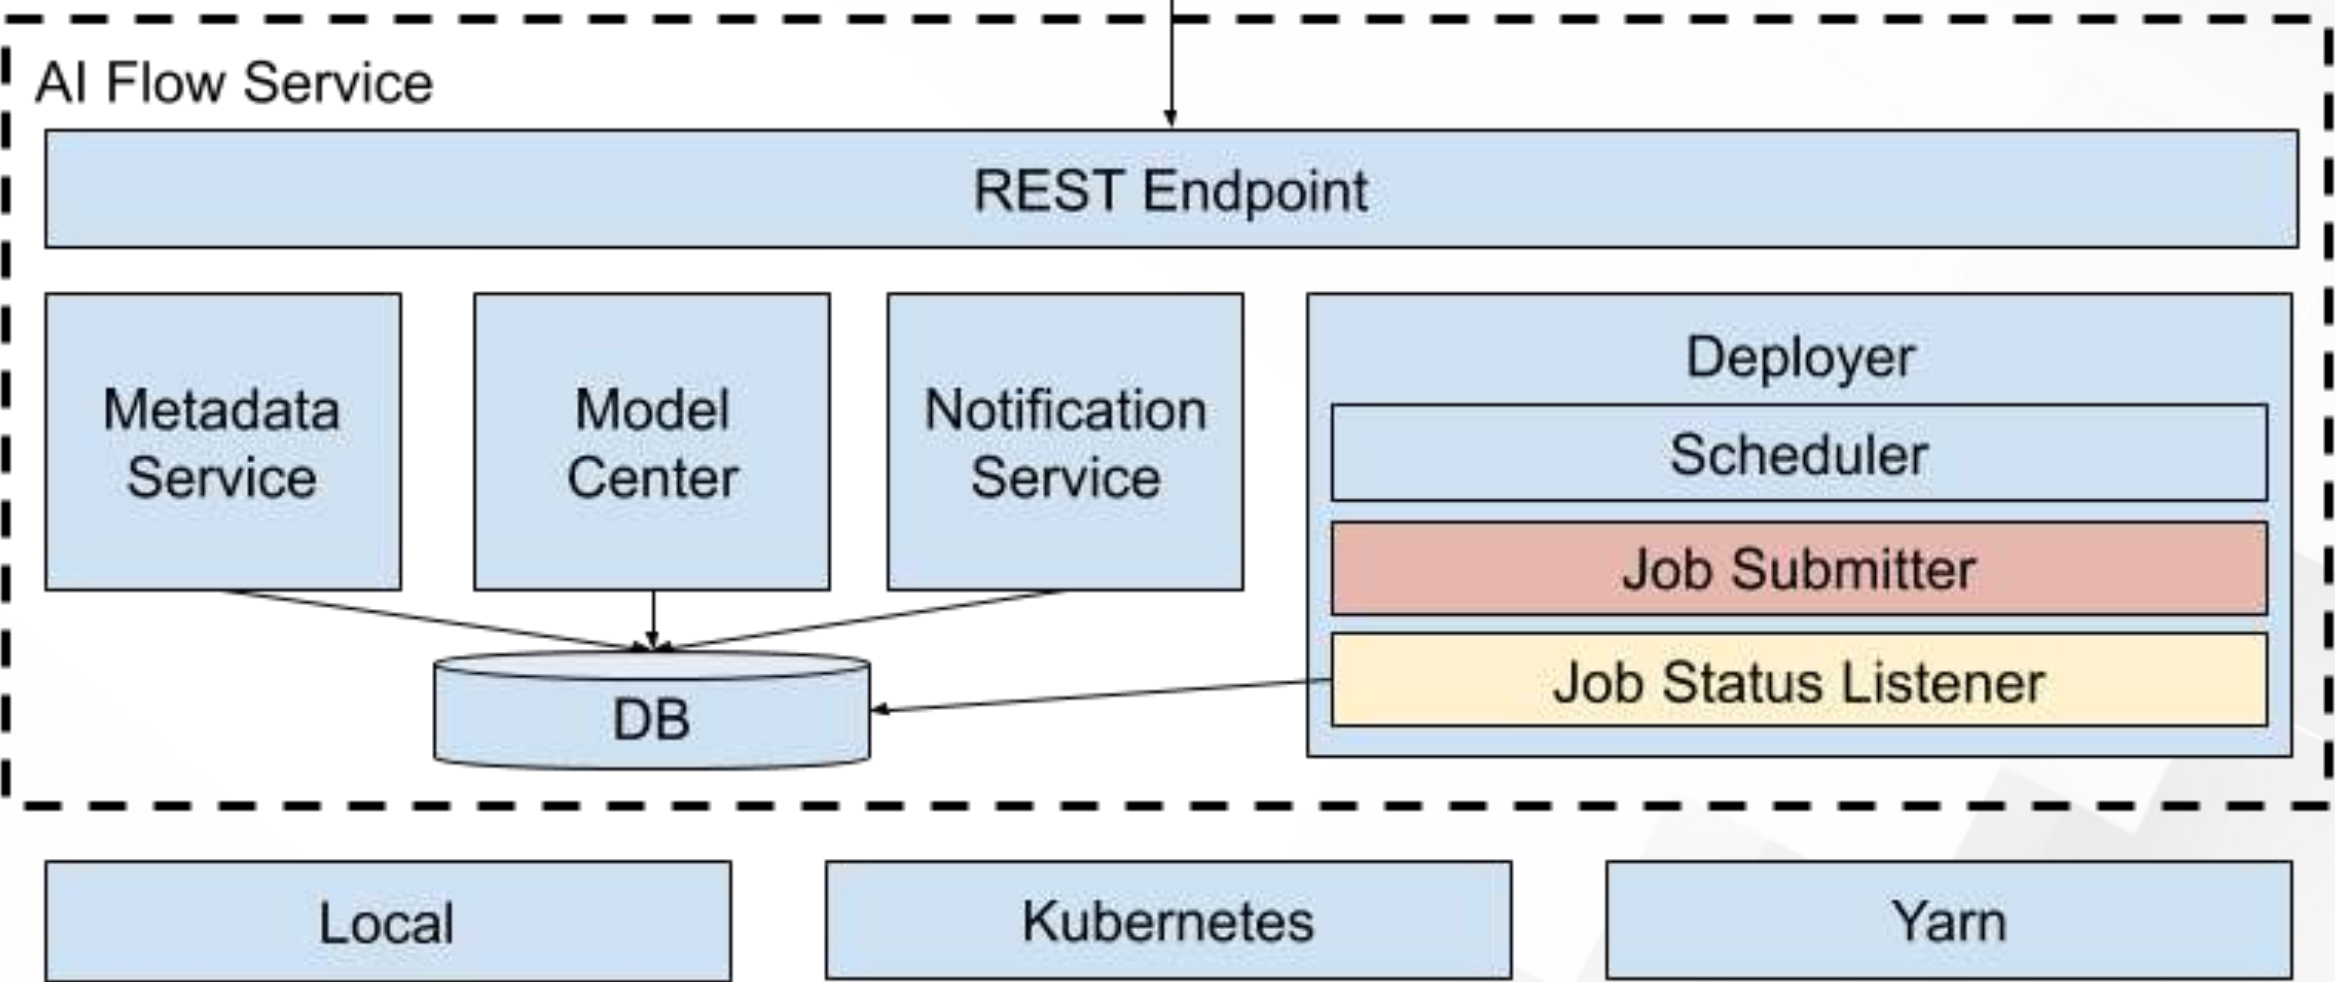
Visualization

Architecture Overview

Define & Compile



Execute

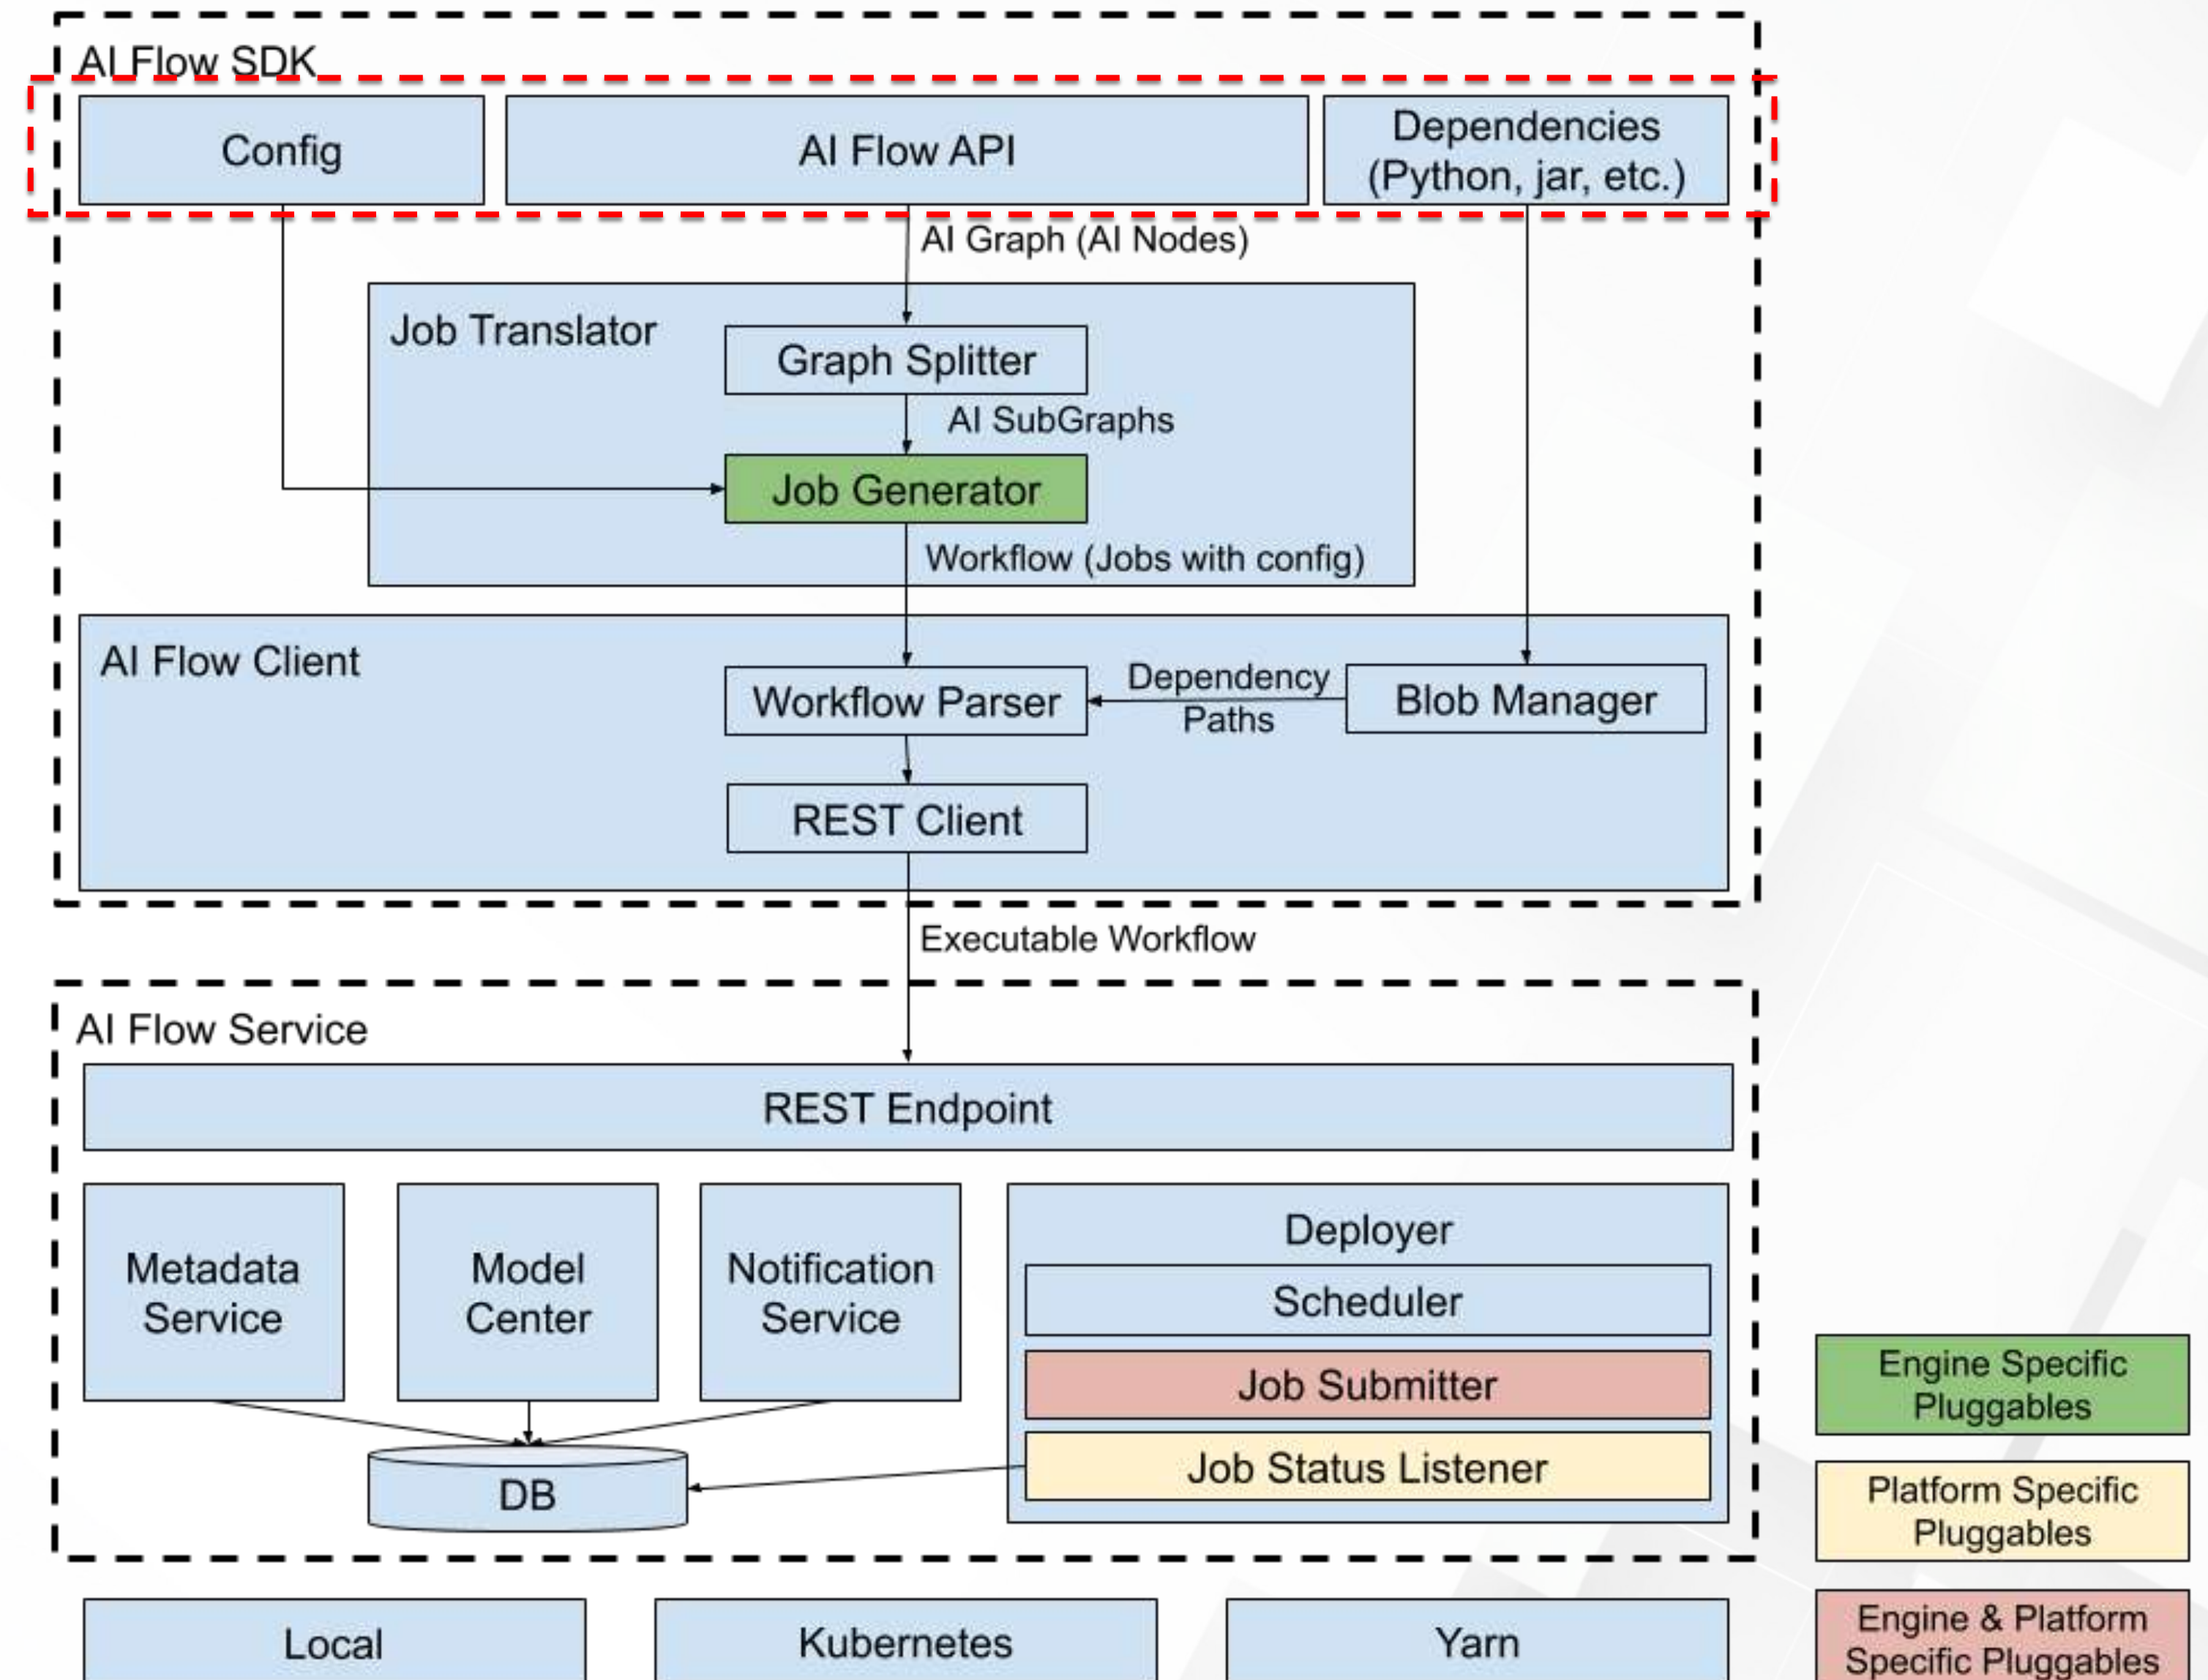


- Engine Specific Pluggables
- Platform Specific Pluggables
- Engine & Platform Specific Pluggables

API Abstraction

Define an **AI Graph**

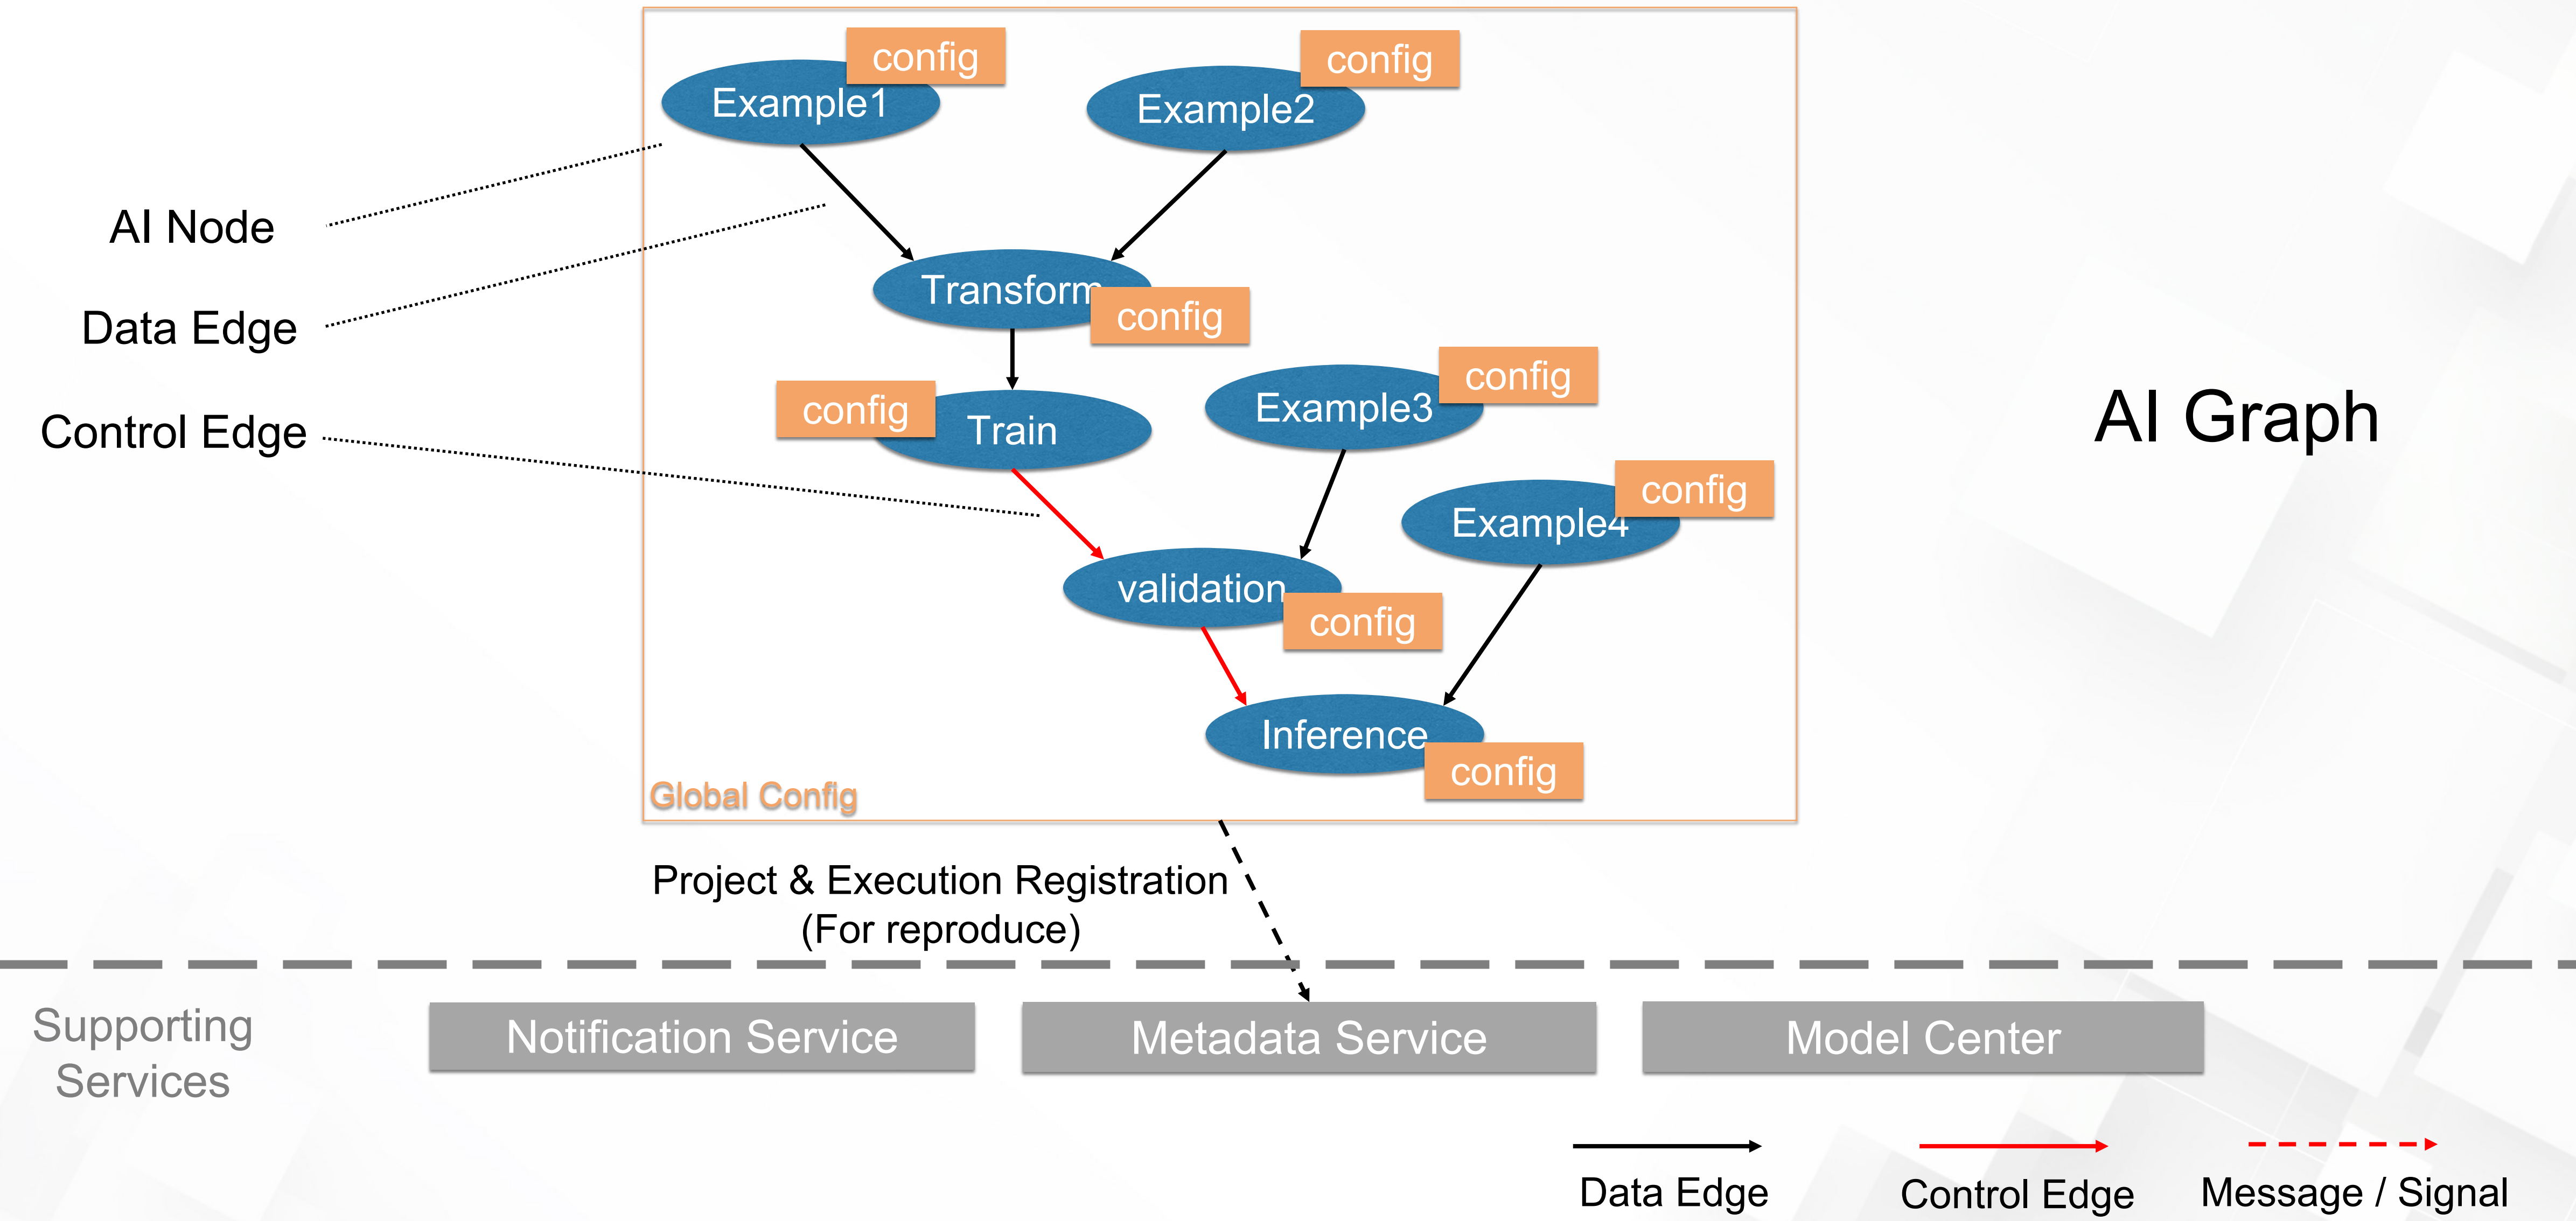
- AI Nodes & relations
- Configs
- Dependencies



Nodes and Relations

- AI Nodes
 - Named **logical** processing unit with multiple inputs and outputs
 - To be interpreted by the ***Job Generator***
- Relation
 - Data dependency (intra-job)
 - Data flows to be interpreted by the ***Job Generator***
 - Control dependency (inter-job)
 - Event driven control flows to be interpreted by the ***Scheduler***

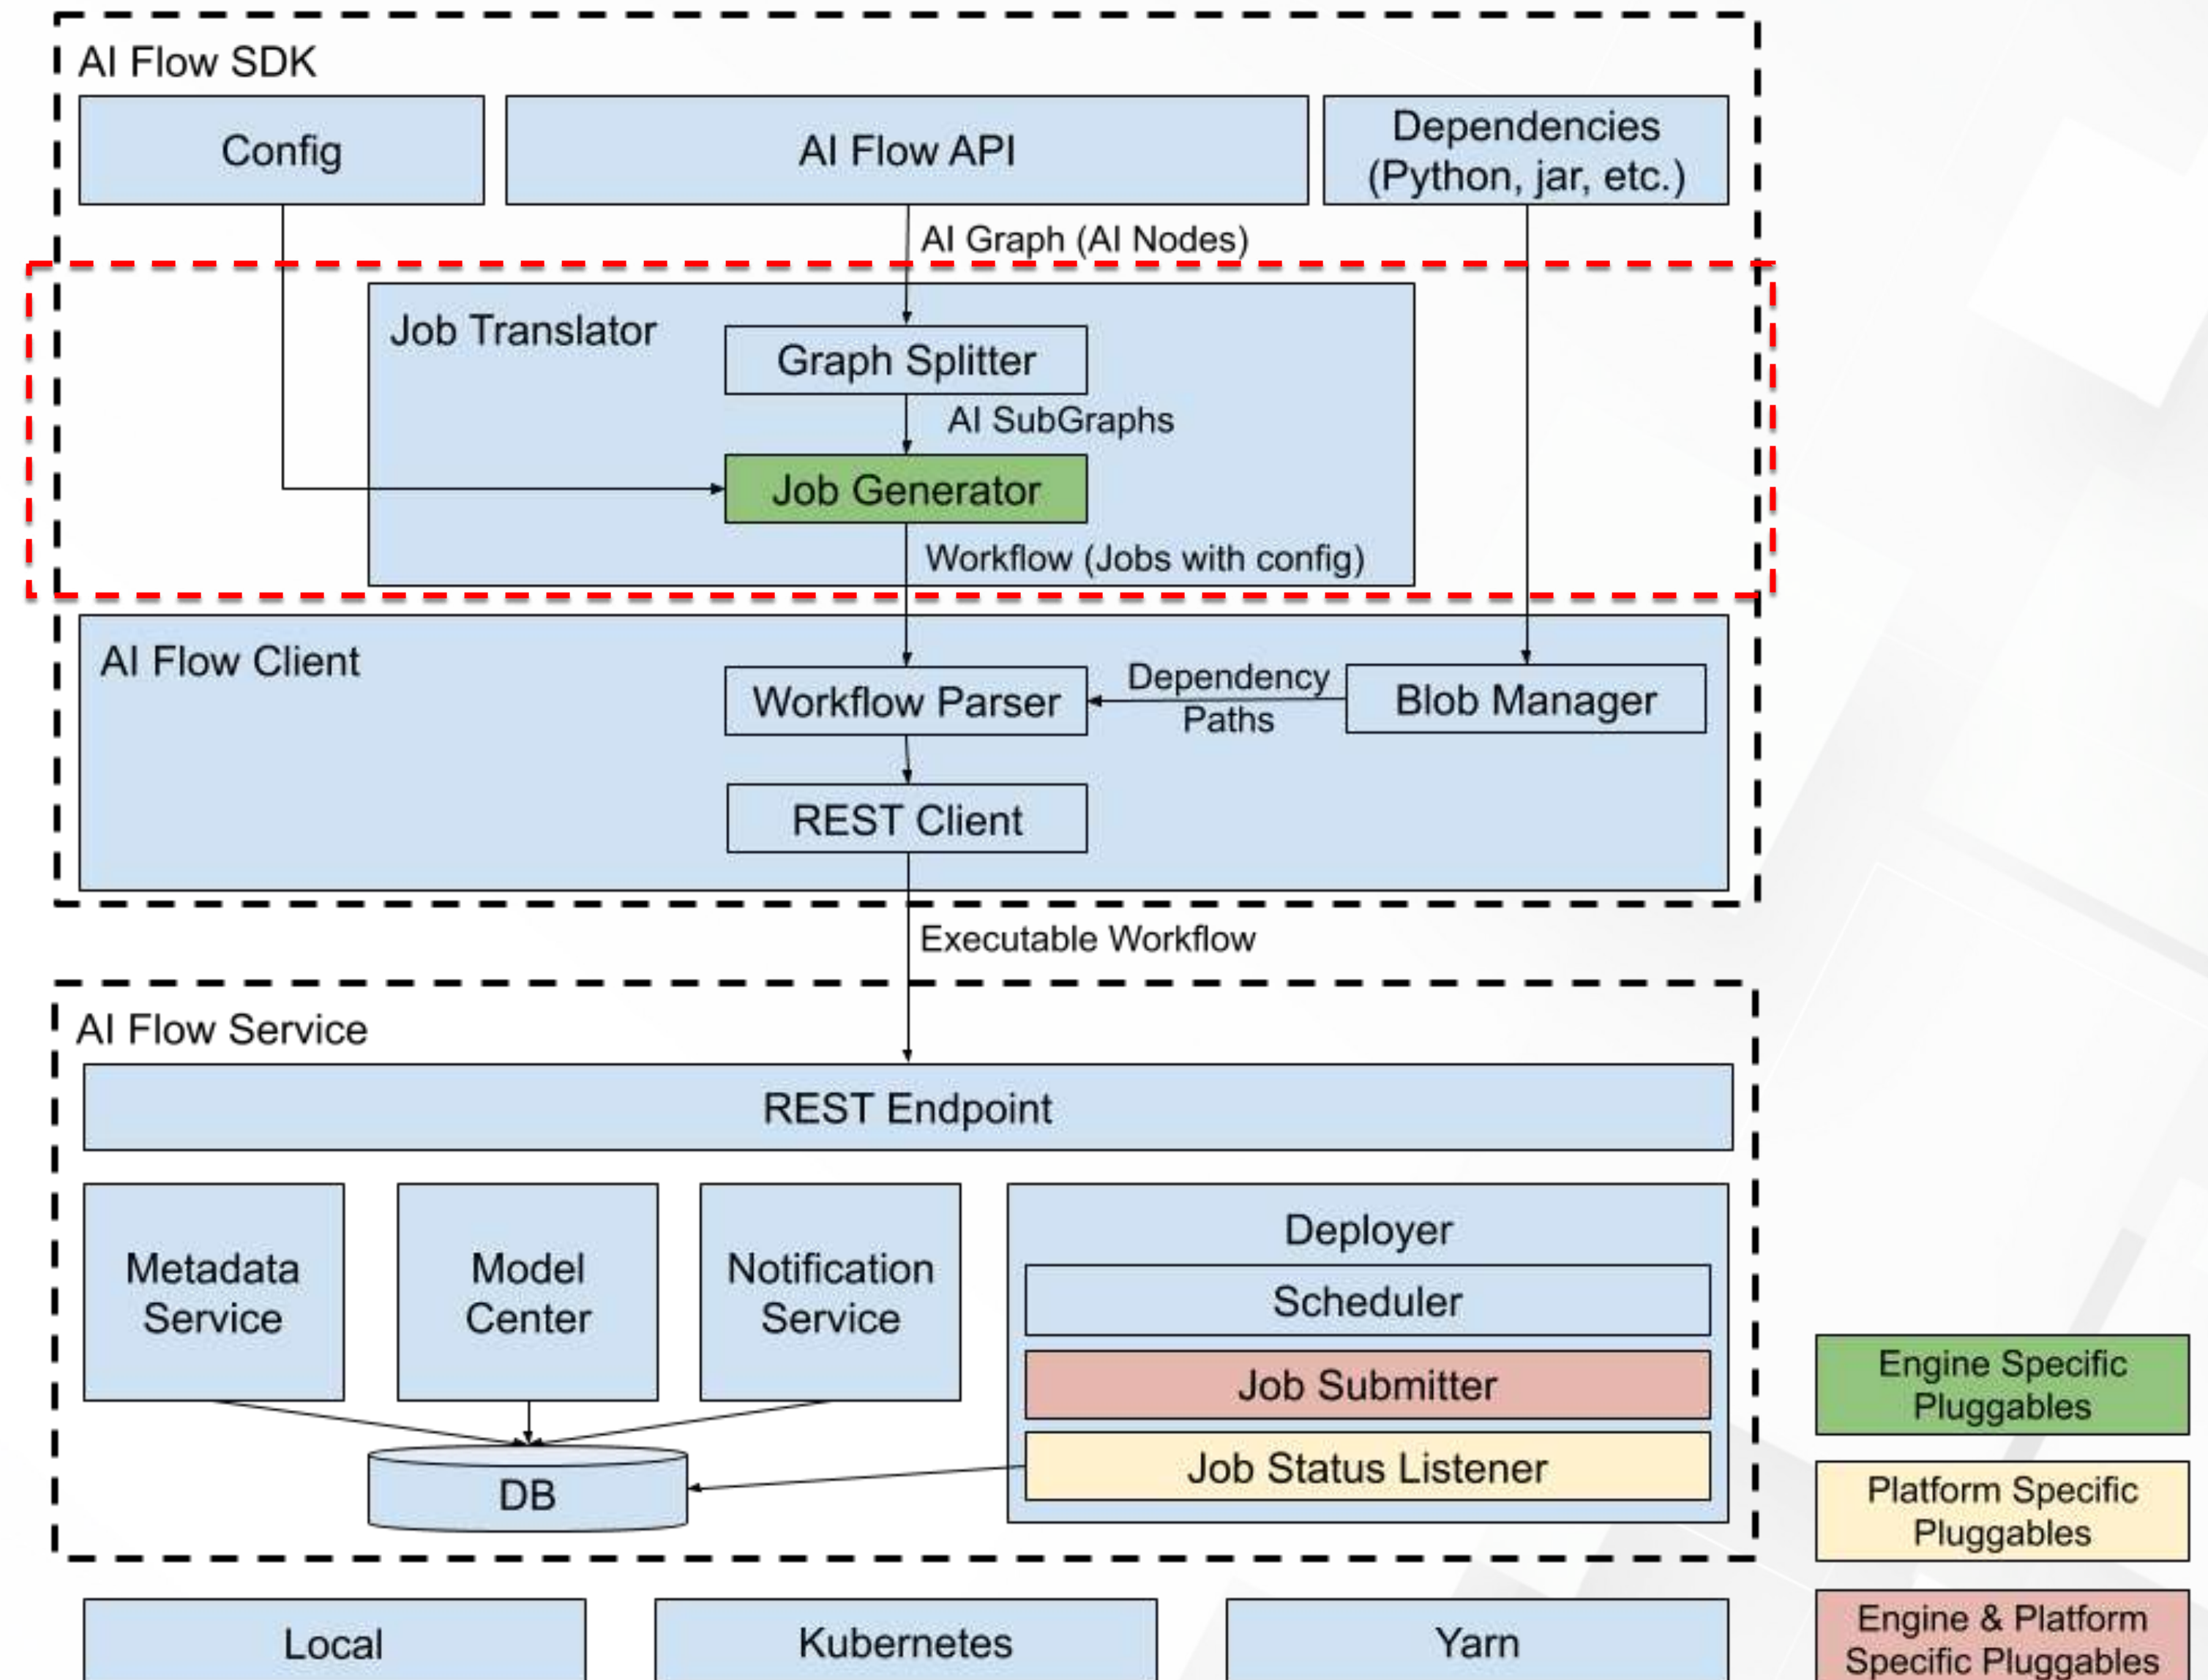
AI Graph Description



Workflow Compilation

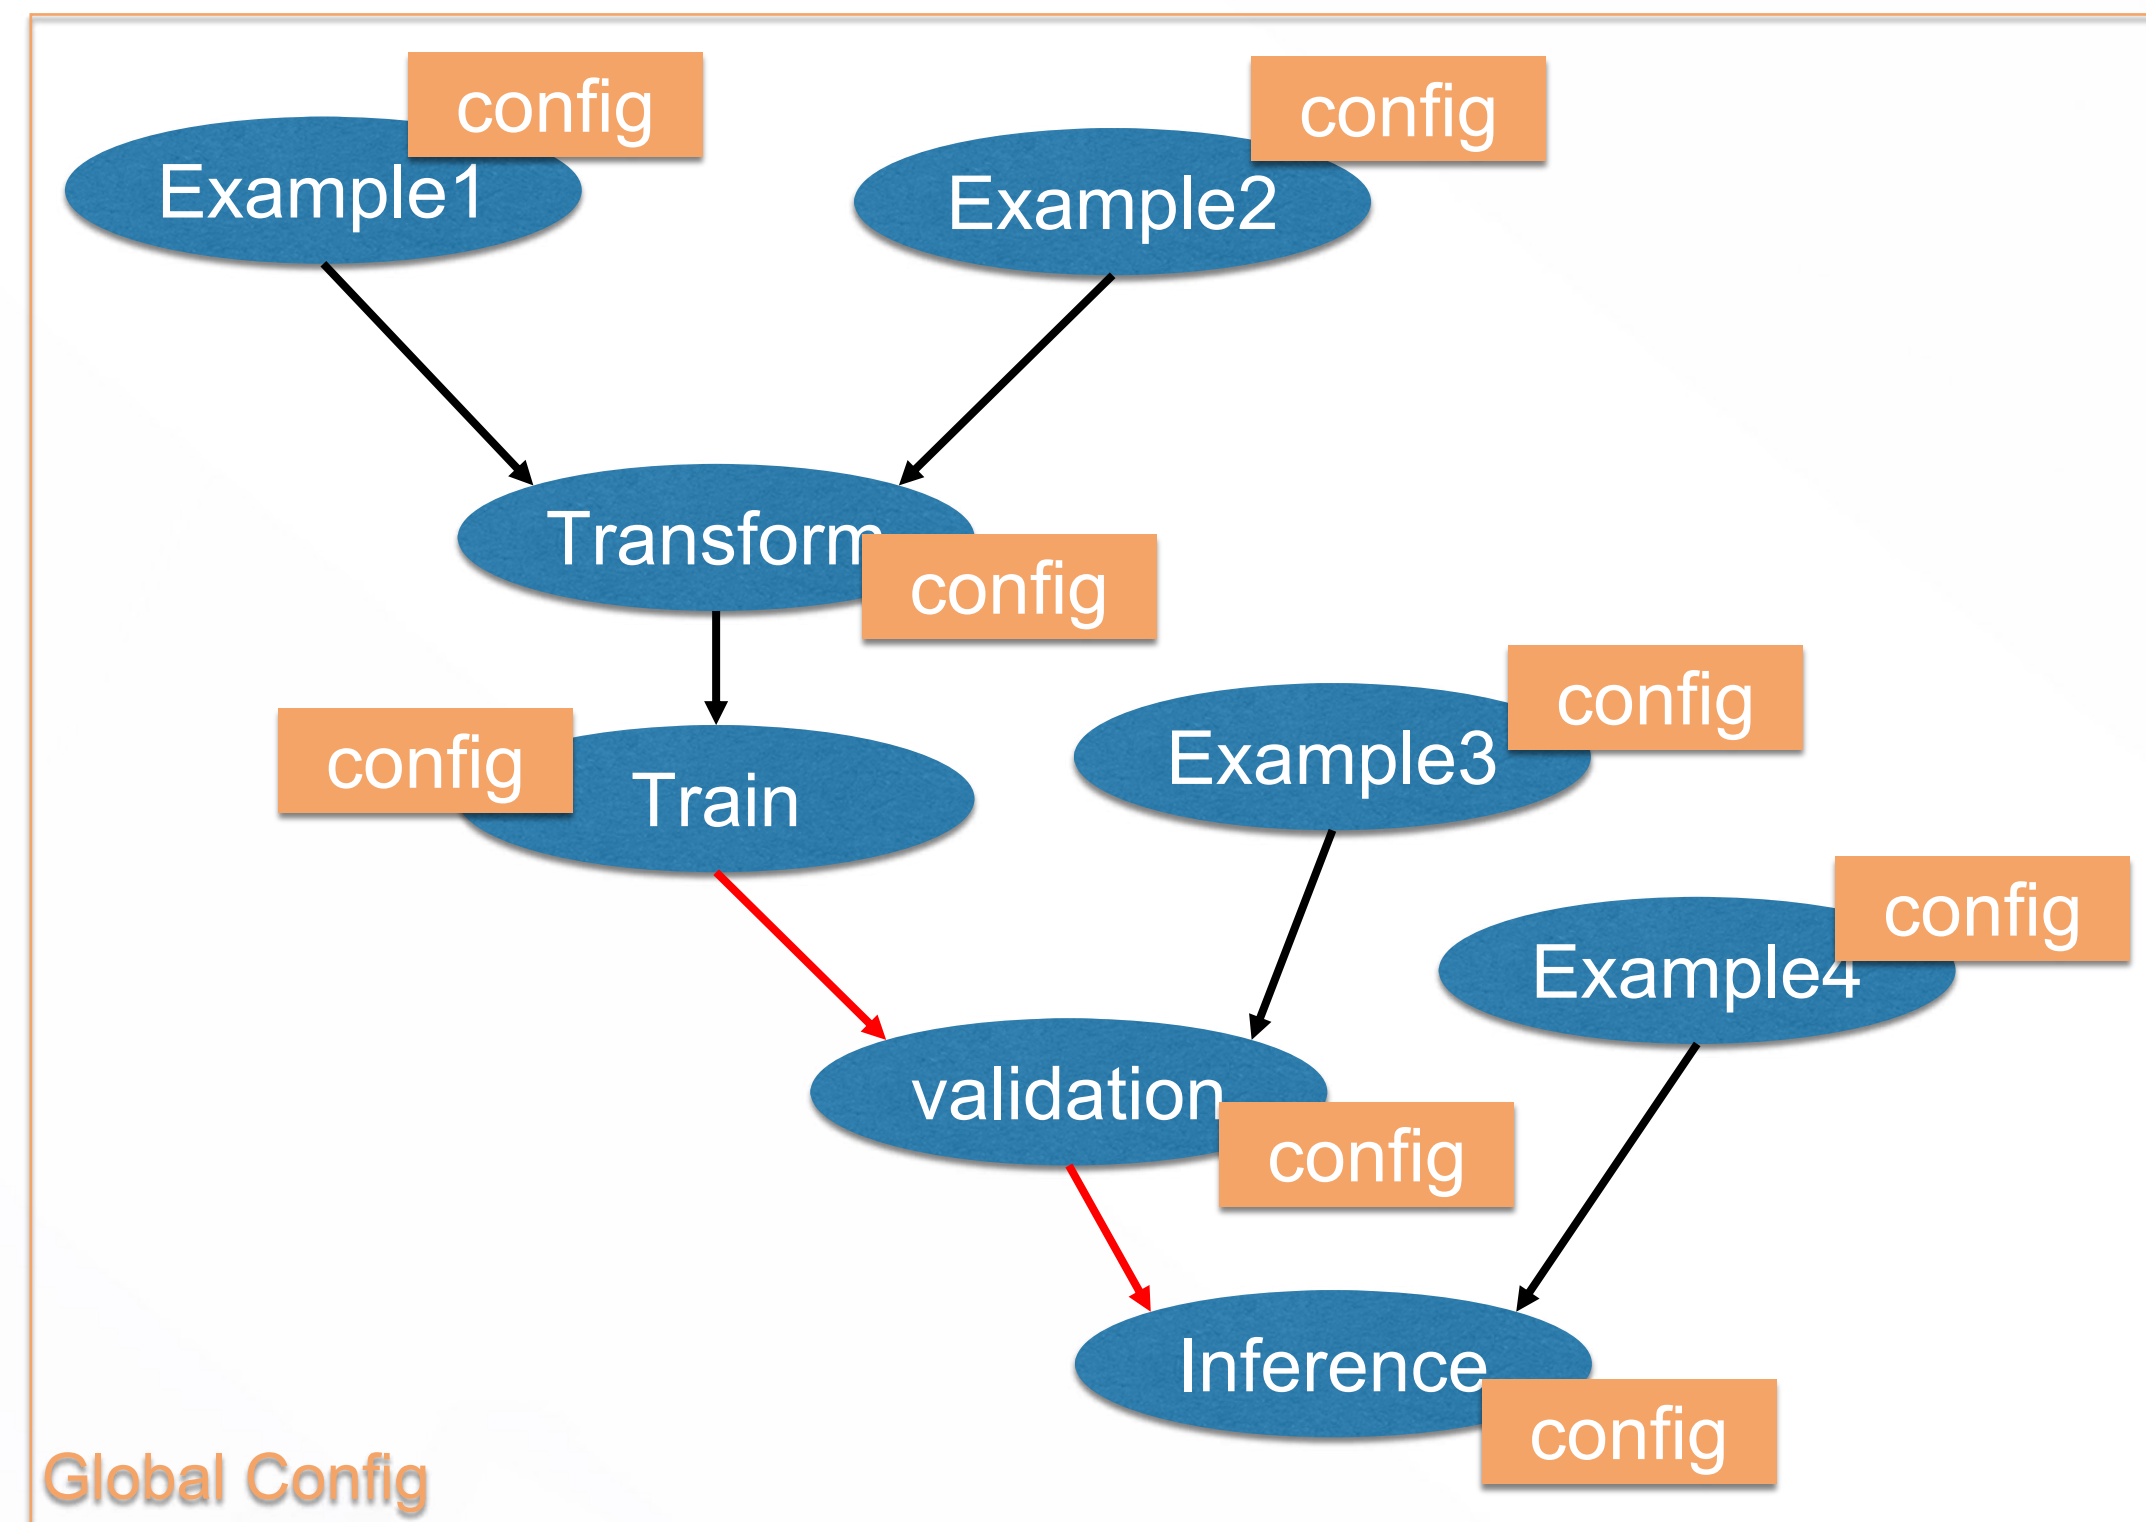
AI Graph → Workflow

- *Split*: AI Graph → Subgraphs
- *Translate*: Subgraph → Job

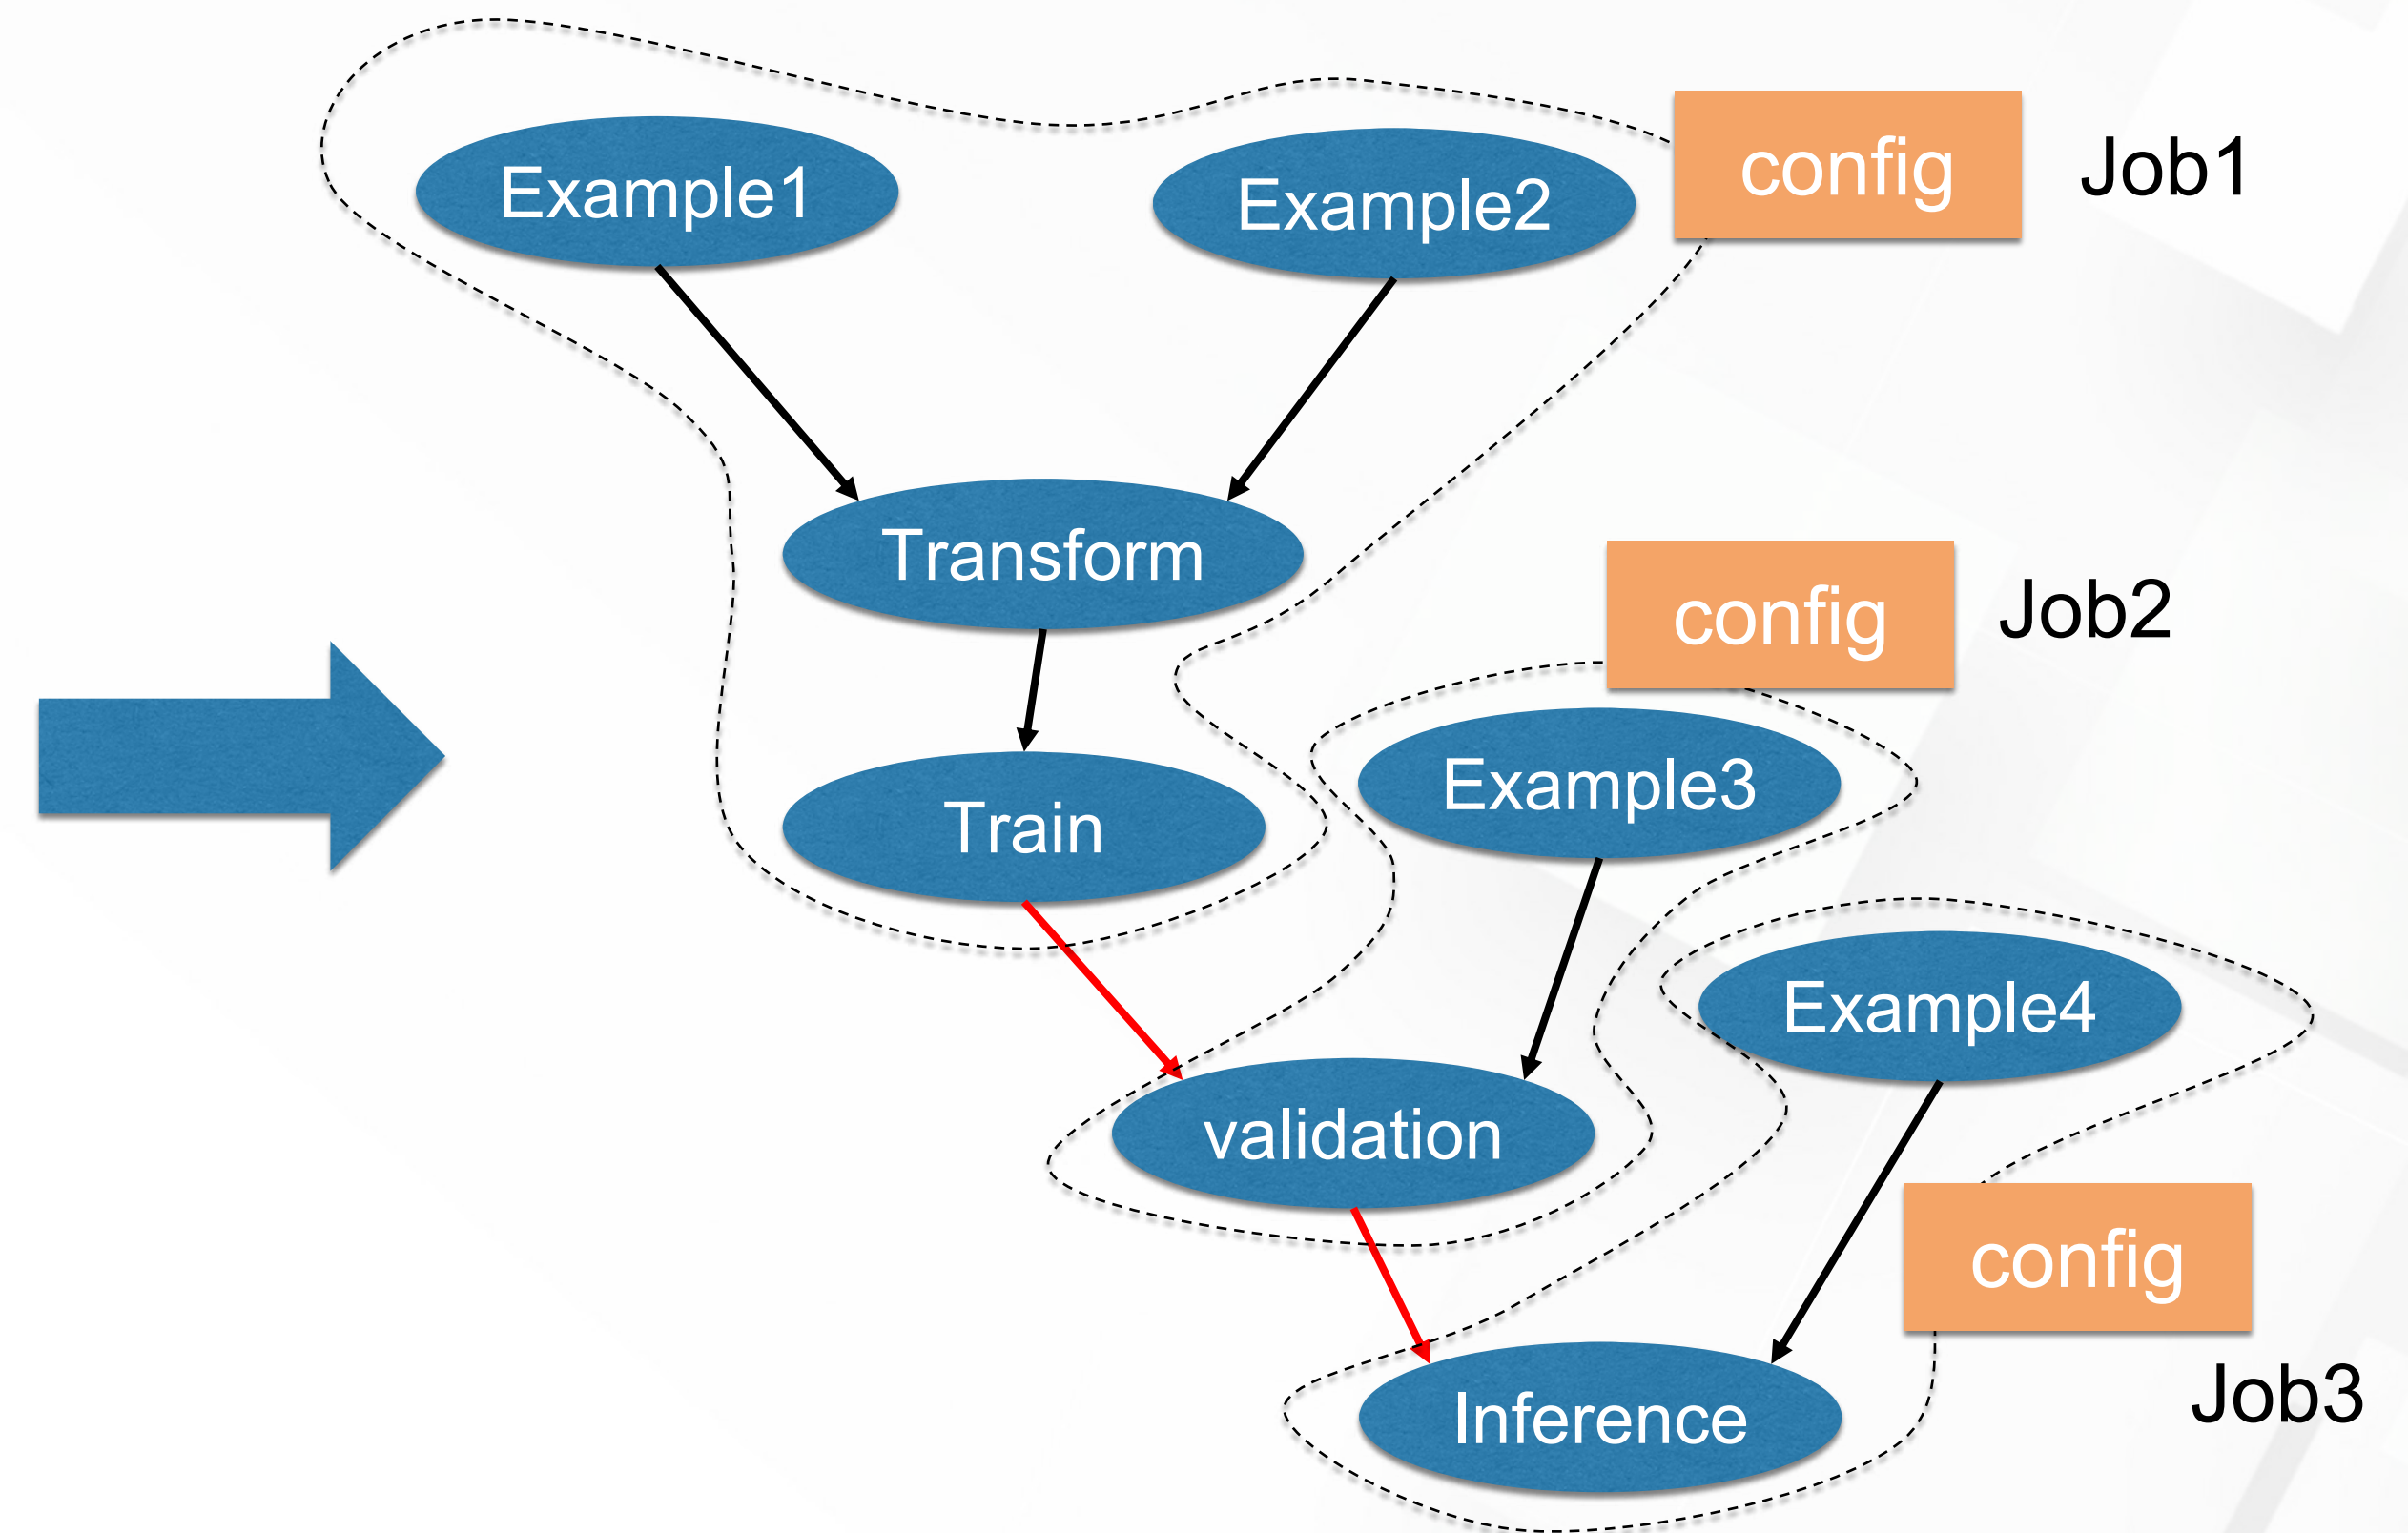


Workflow Compilation

AI Graph



Workflow



Supporting
Services

Notification Service

Metadata Service

Model Center

→
Data Edge

→
Control Edge

- - - - -
Message / Signal

Execution Blocks

- Specify computing engine and config with execution blocks

```
with af.global_config(af.BaseJobConfig(platform="K8S", engine="Flink", properties={"aa": "aa"})): Global Config
```

```
with af.job_config(config=config1):  
    .... } Flink Job
```

```
with af.engine('python'):  
    with af.job_config(config=config2):  
        input = af.read_example(input, ...)  
        processed = af.transform(input_data_list=[input],  
                                executor=PythonObjectExecutor(python_object=T1()))  
        write = af.write_example(input_data=processed, ...)
```

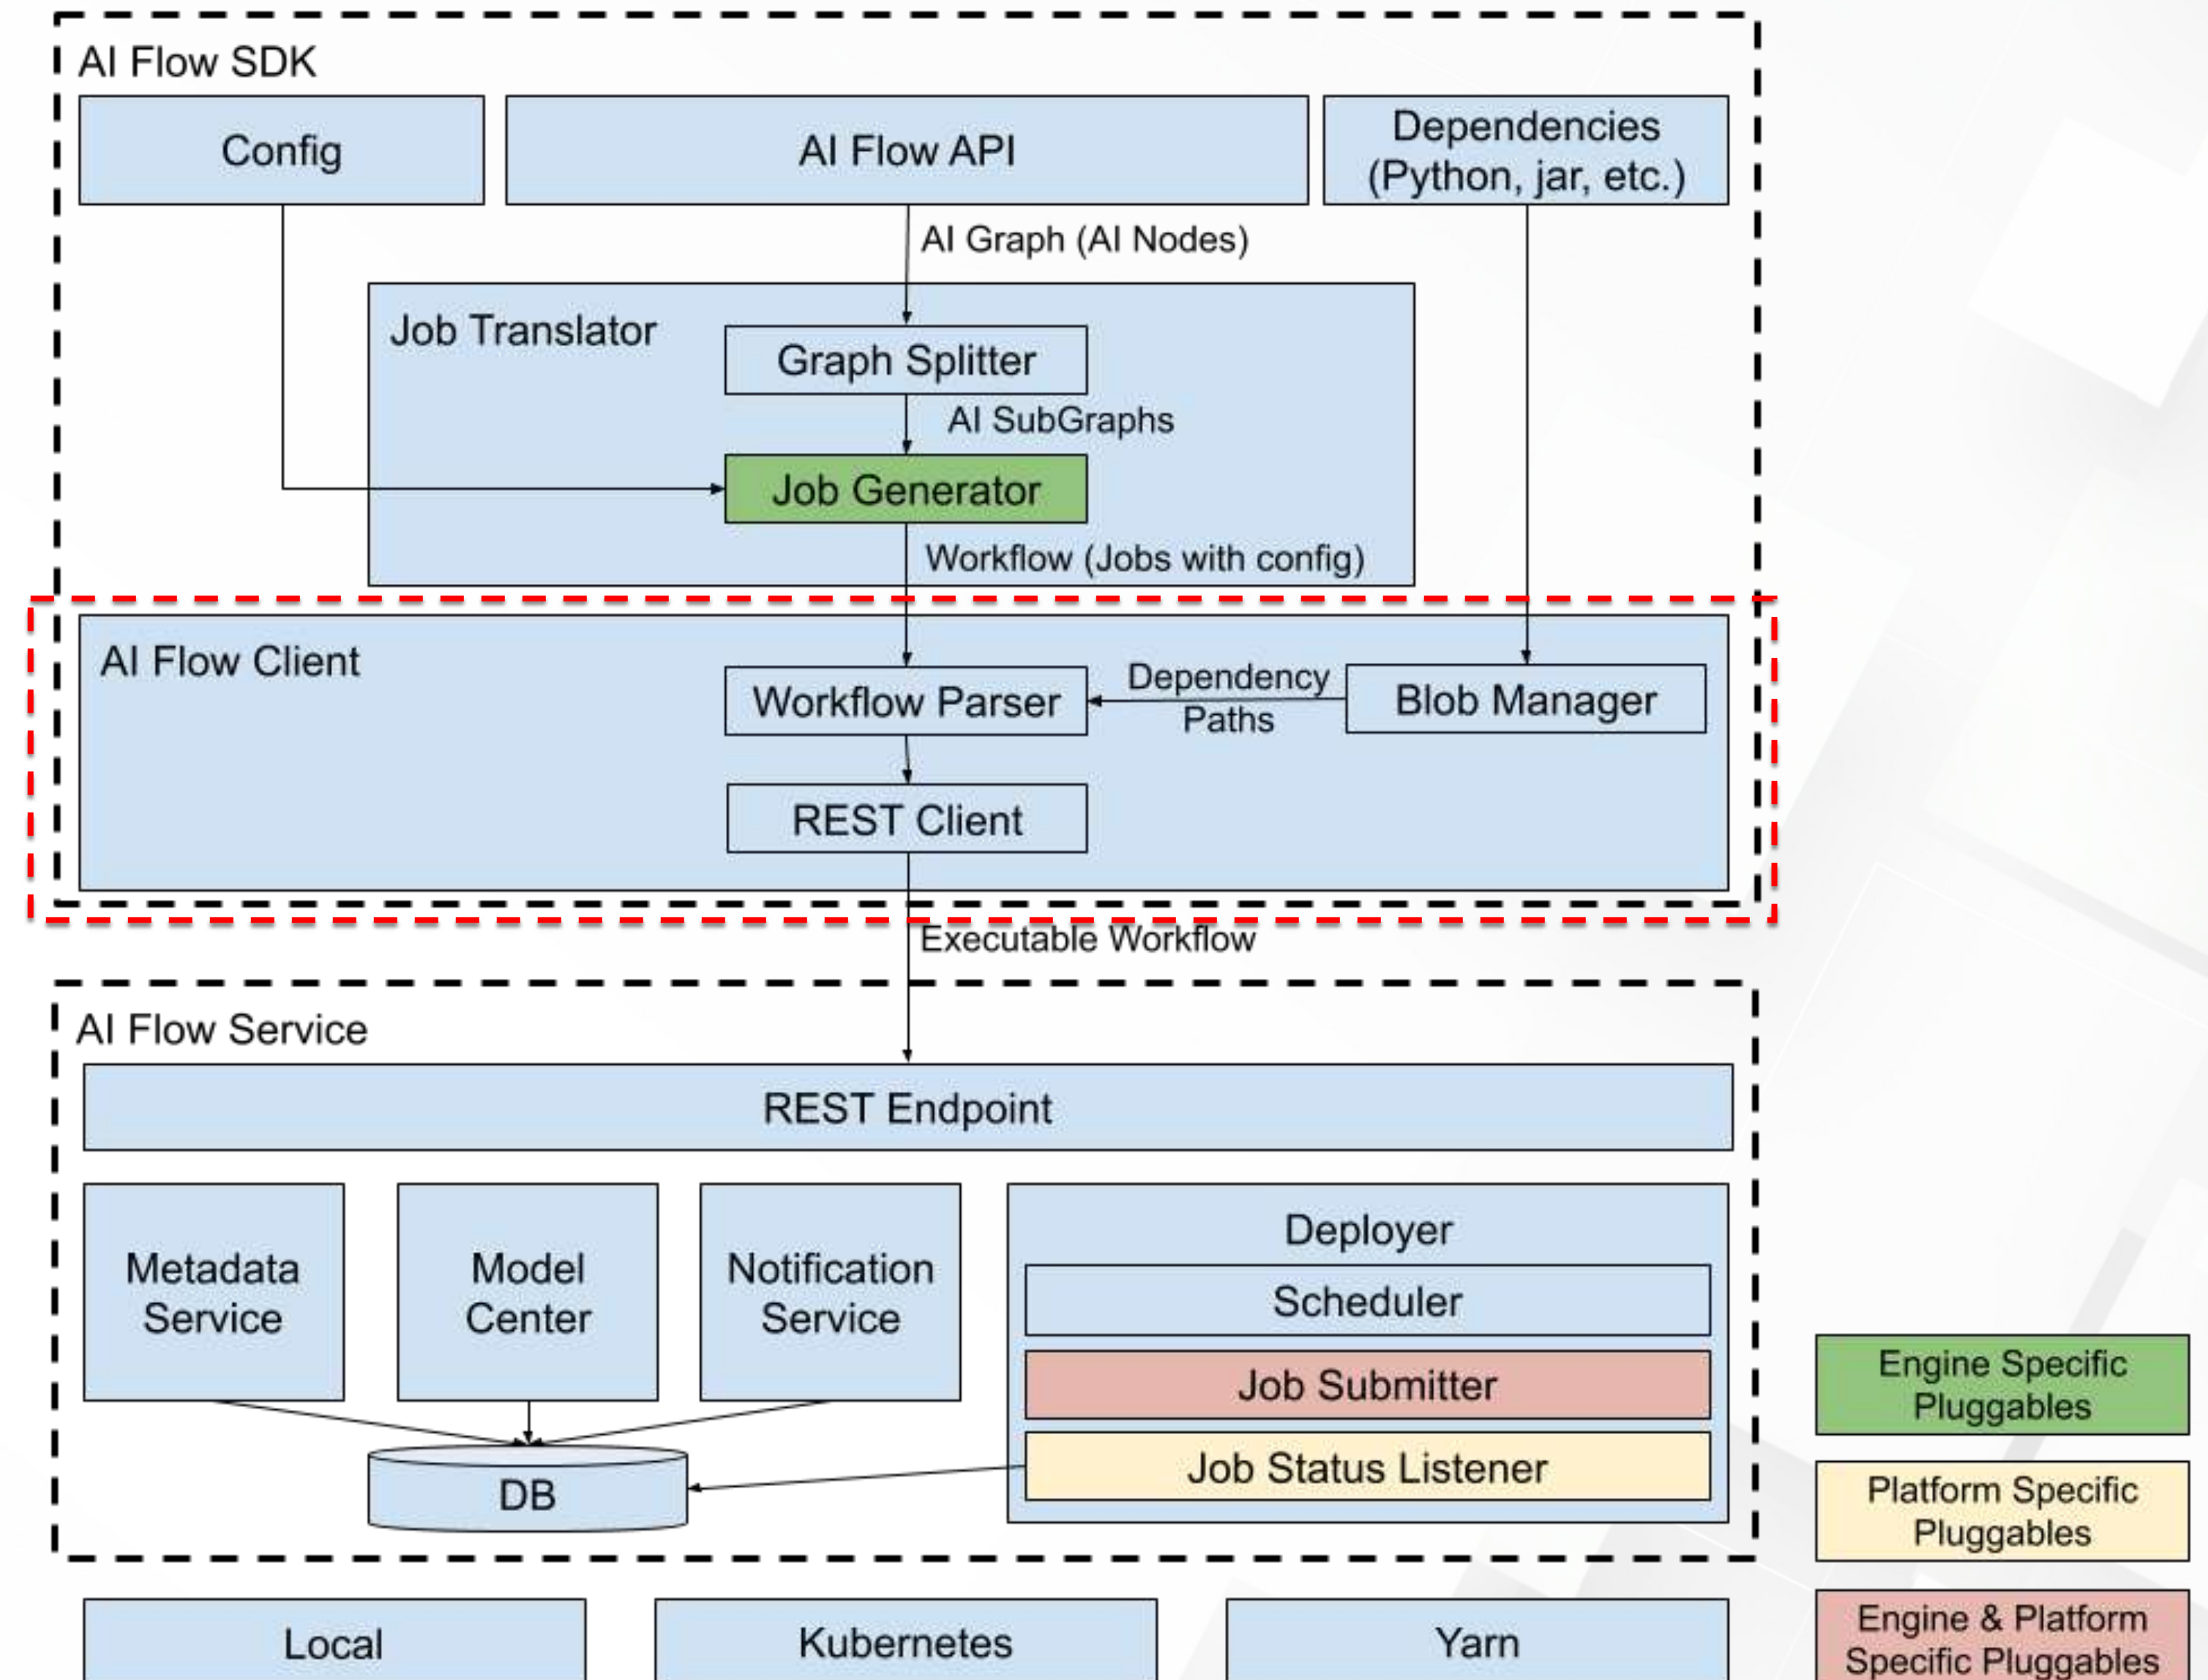
} Python Job

```
with af.job_config(config=config1):  
    .... } Flink Job
```

Construct Executable Workflows

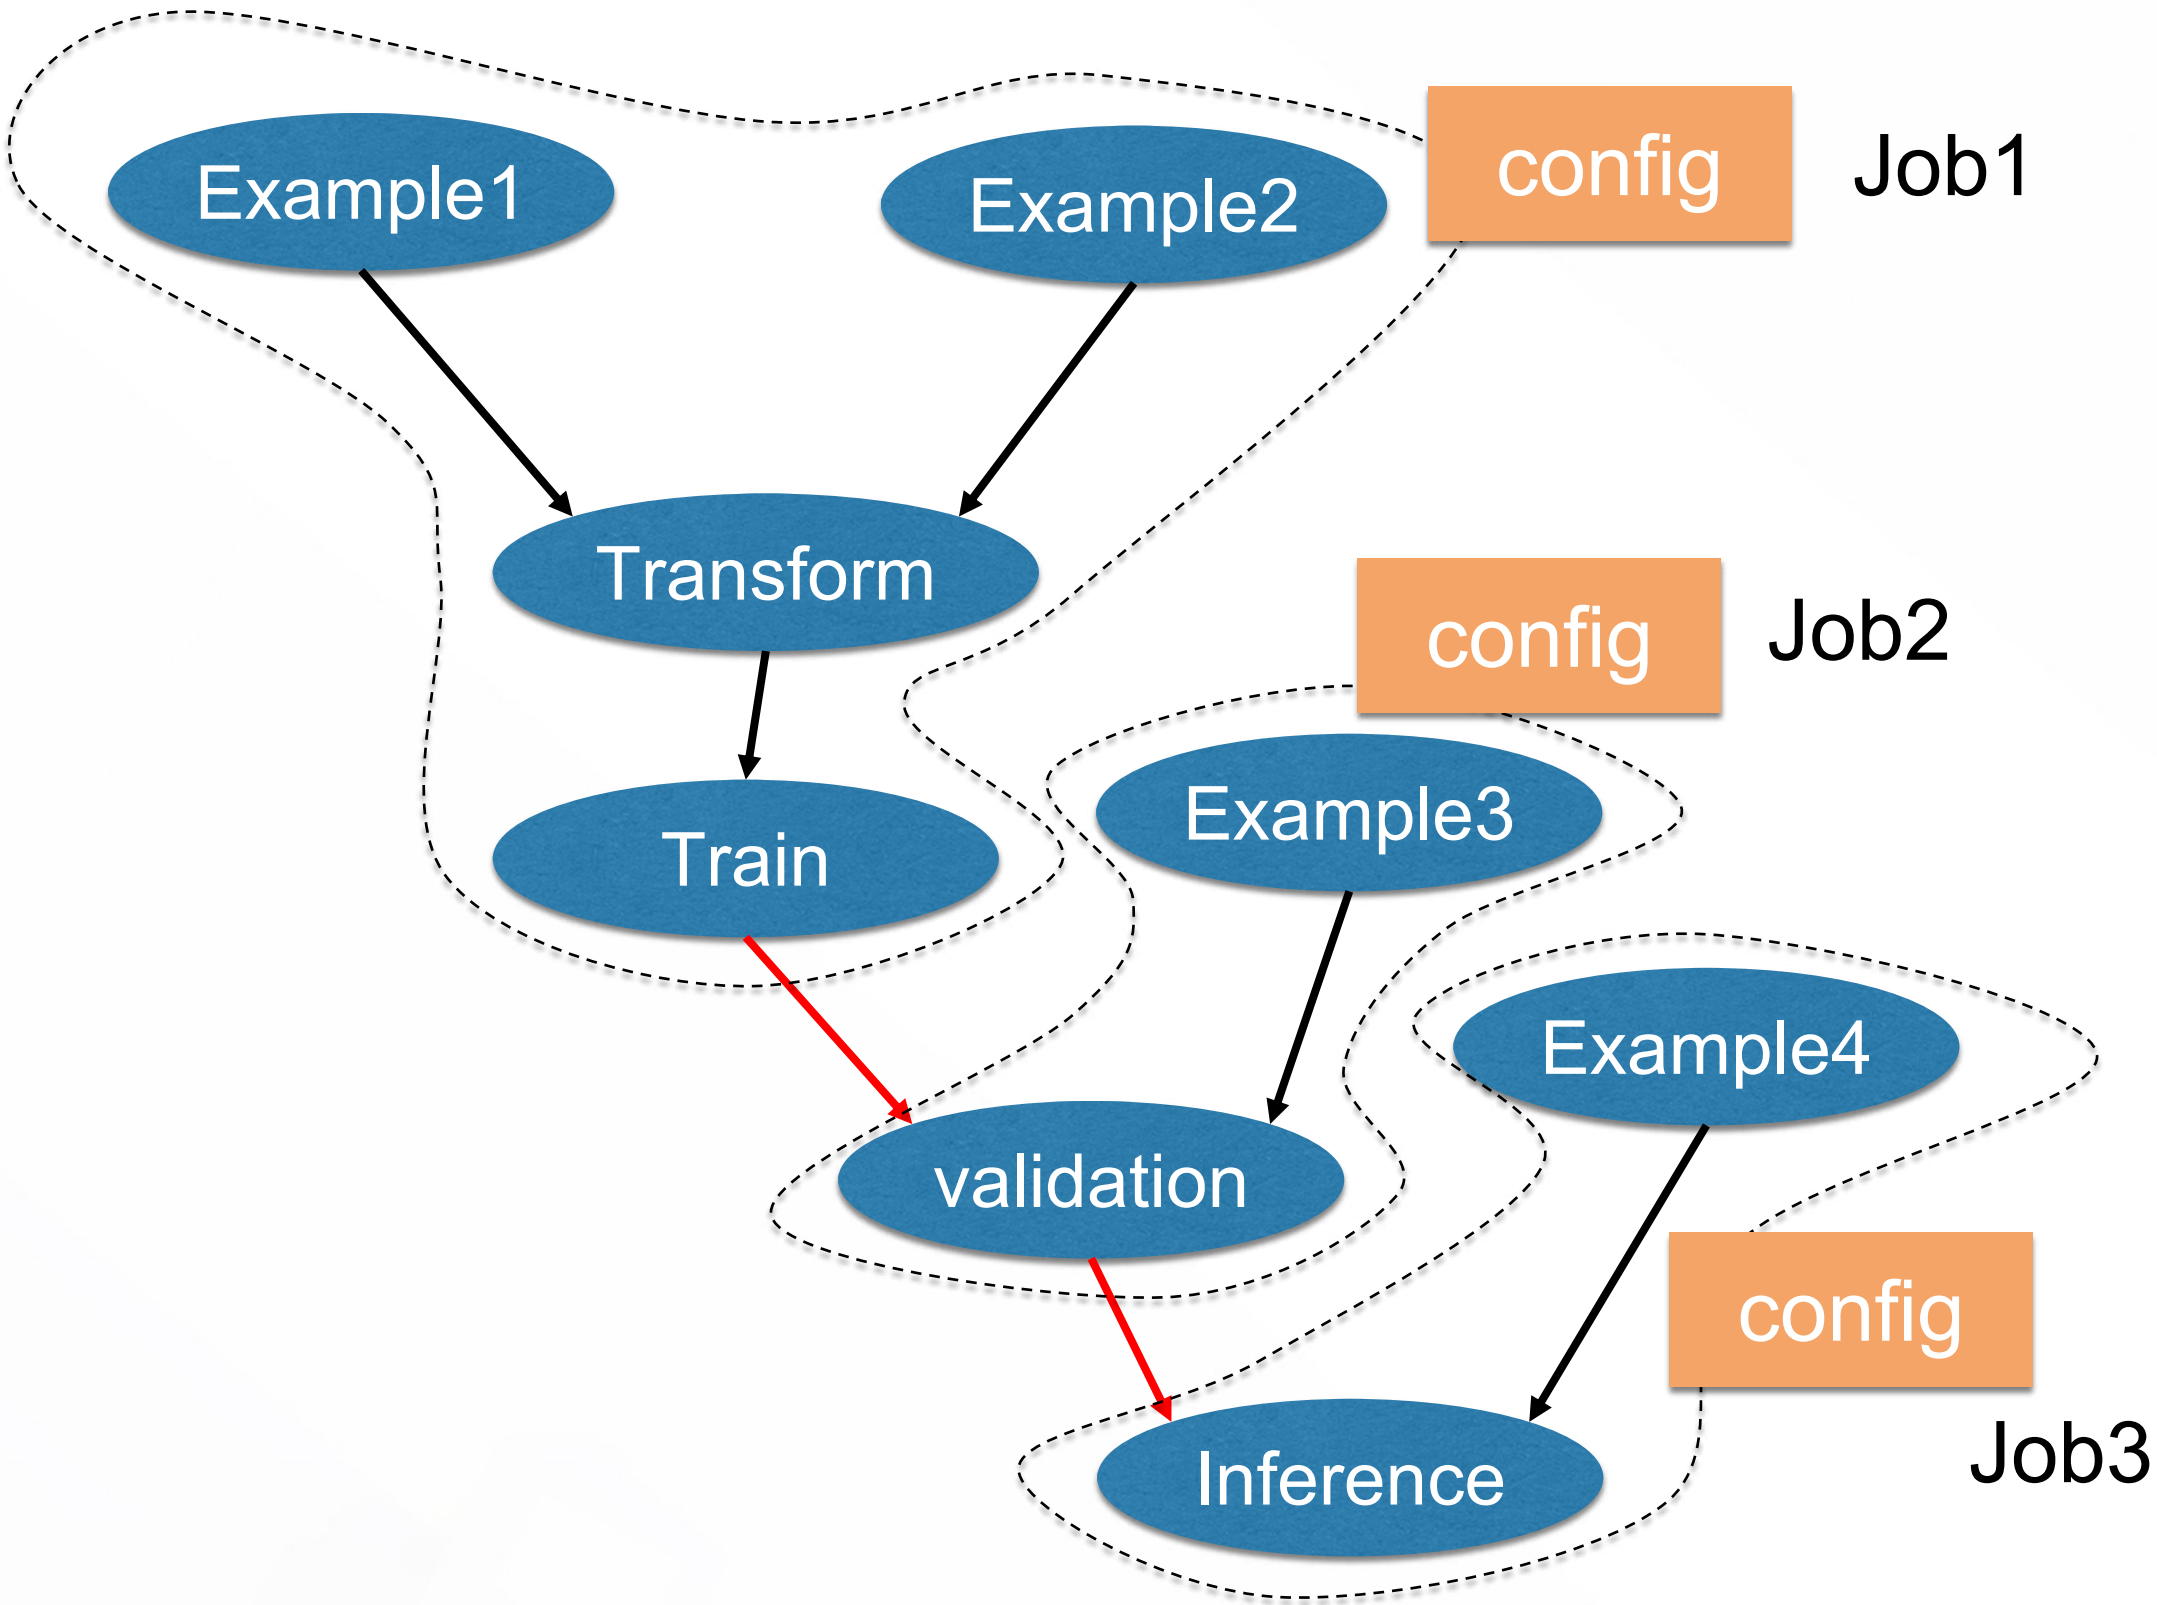
Workflow → Executable Workflow.

- Upload code & dependencies
- Update **Workflow** with URI.

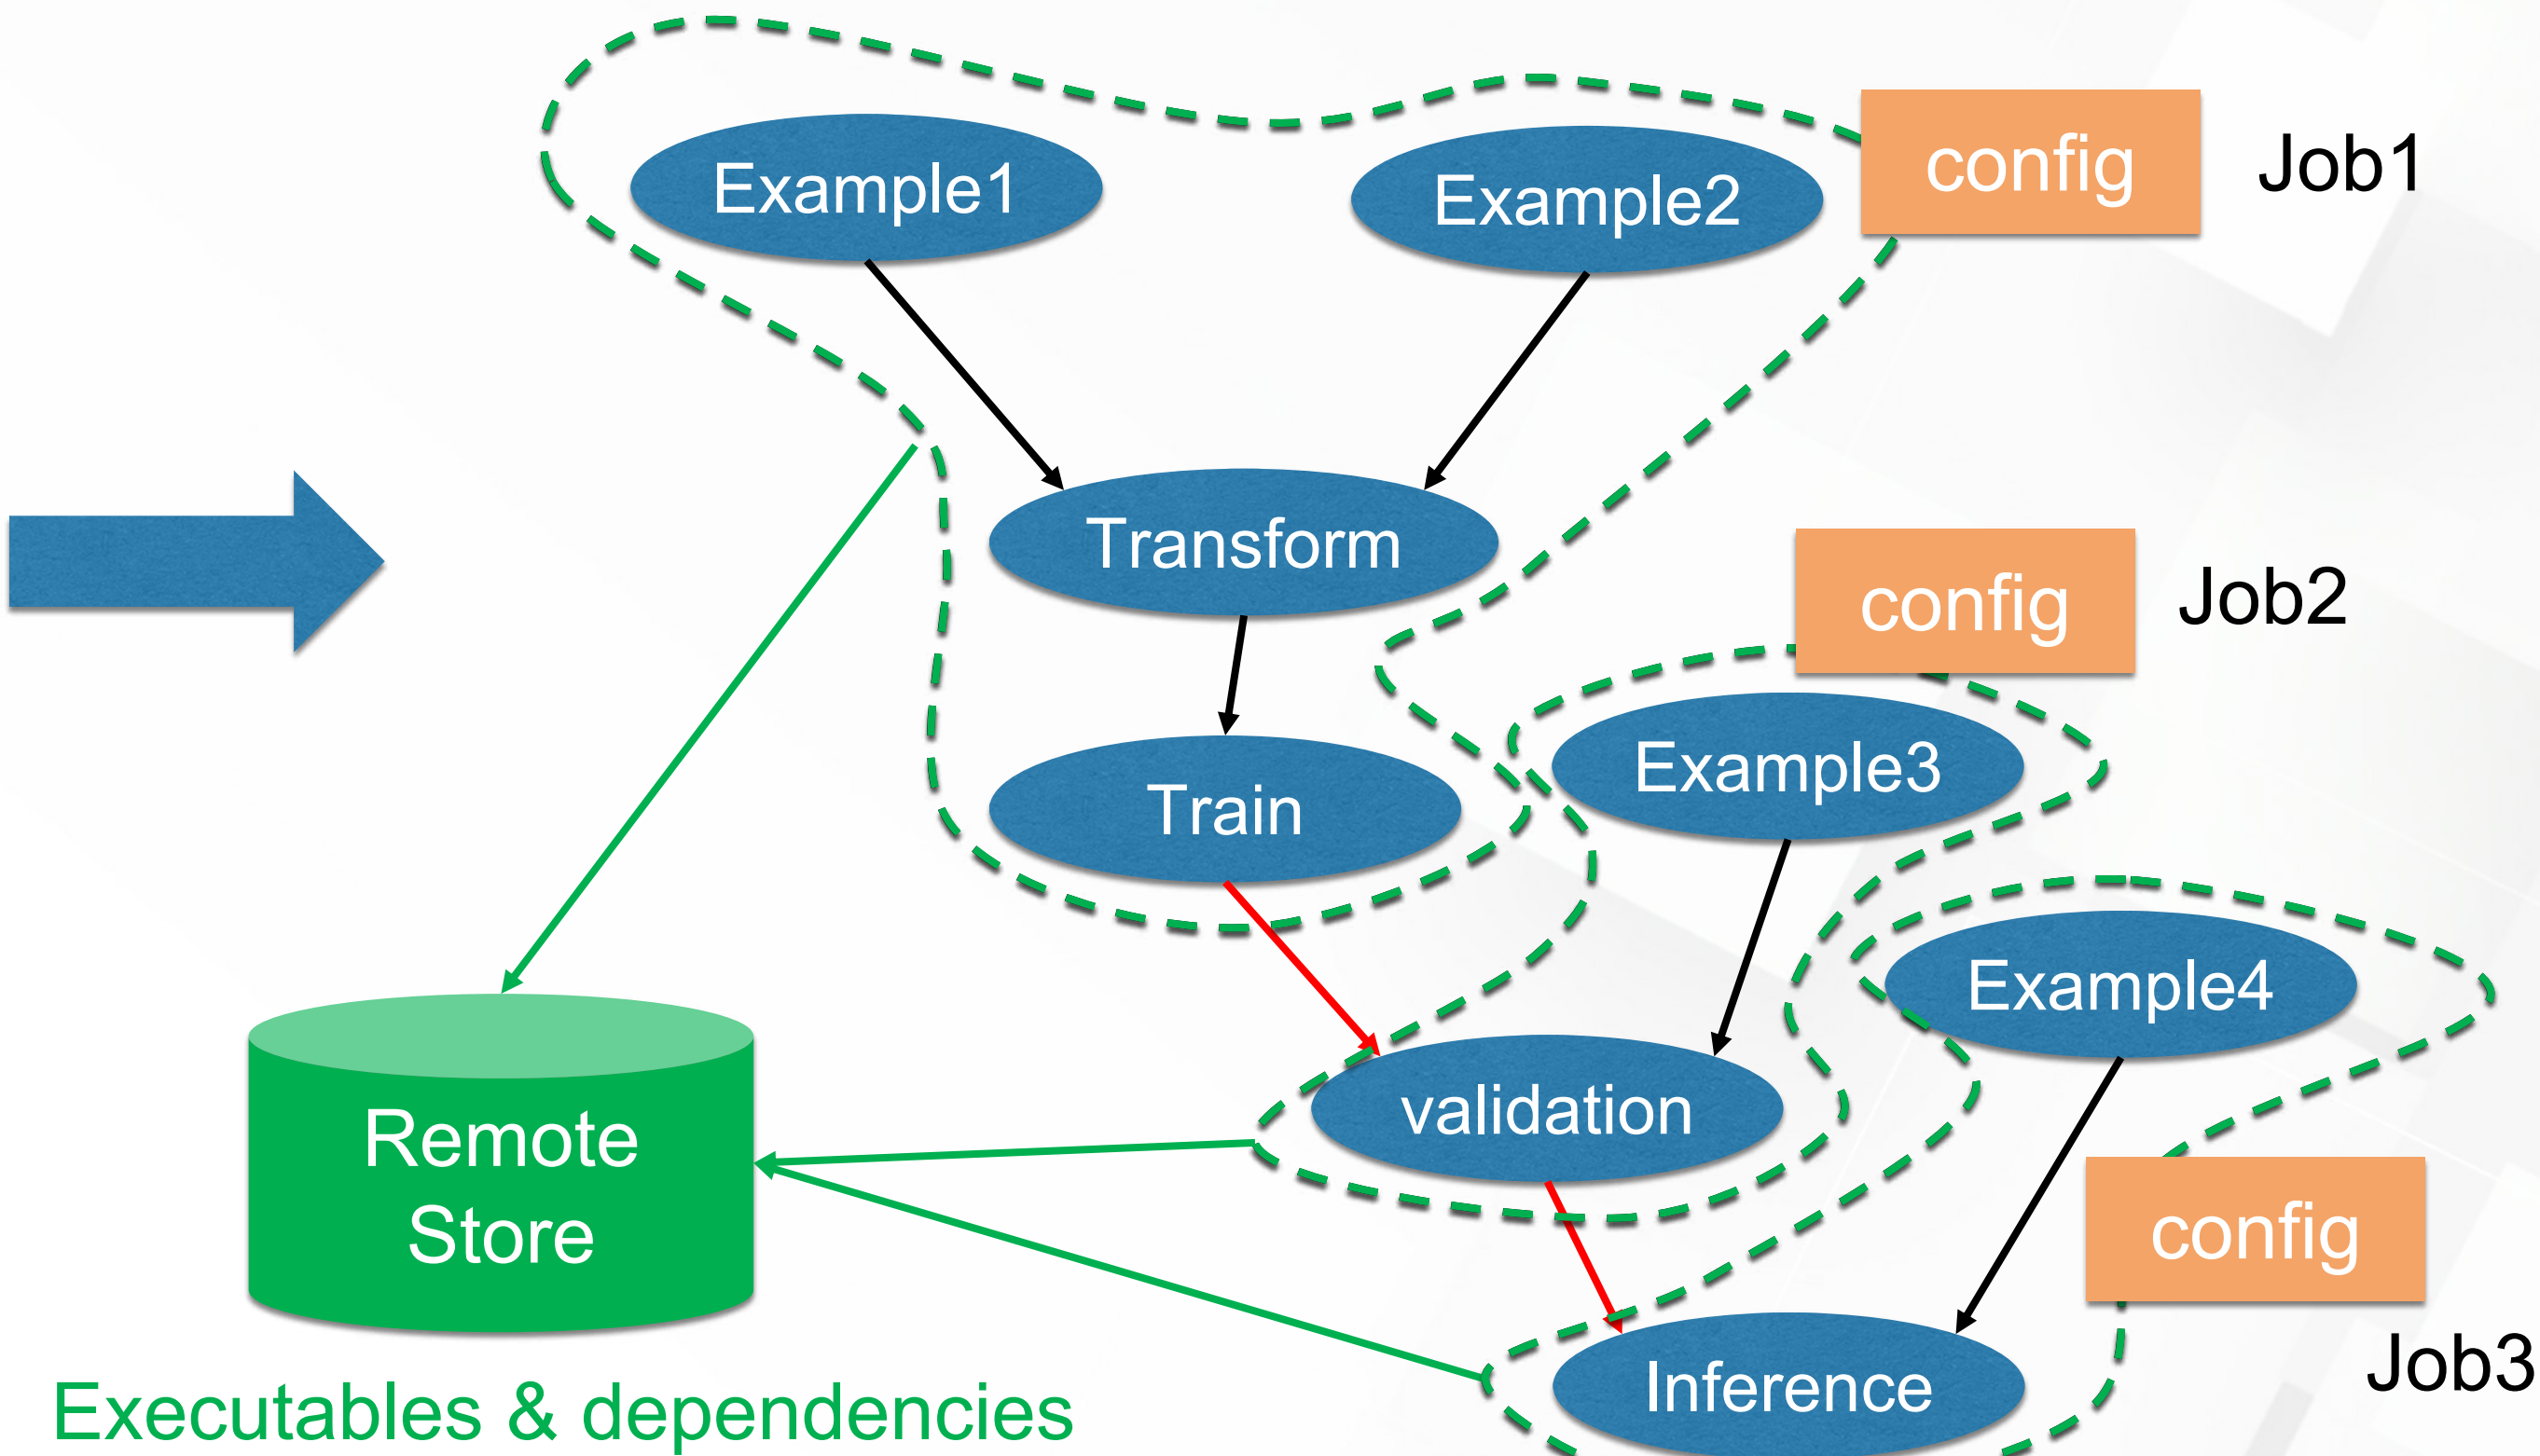


Executable Workflow Construction

Workflow



Executable Workflow



Supporting
Services

Notification Service

Metadata Service

Model Center

 Code and Dep.

 Data Edge

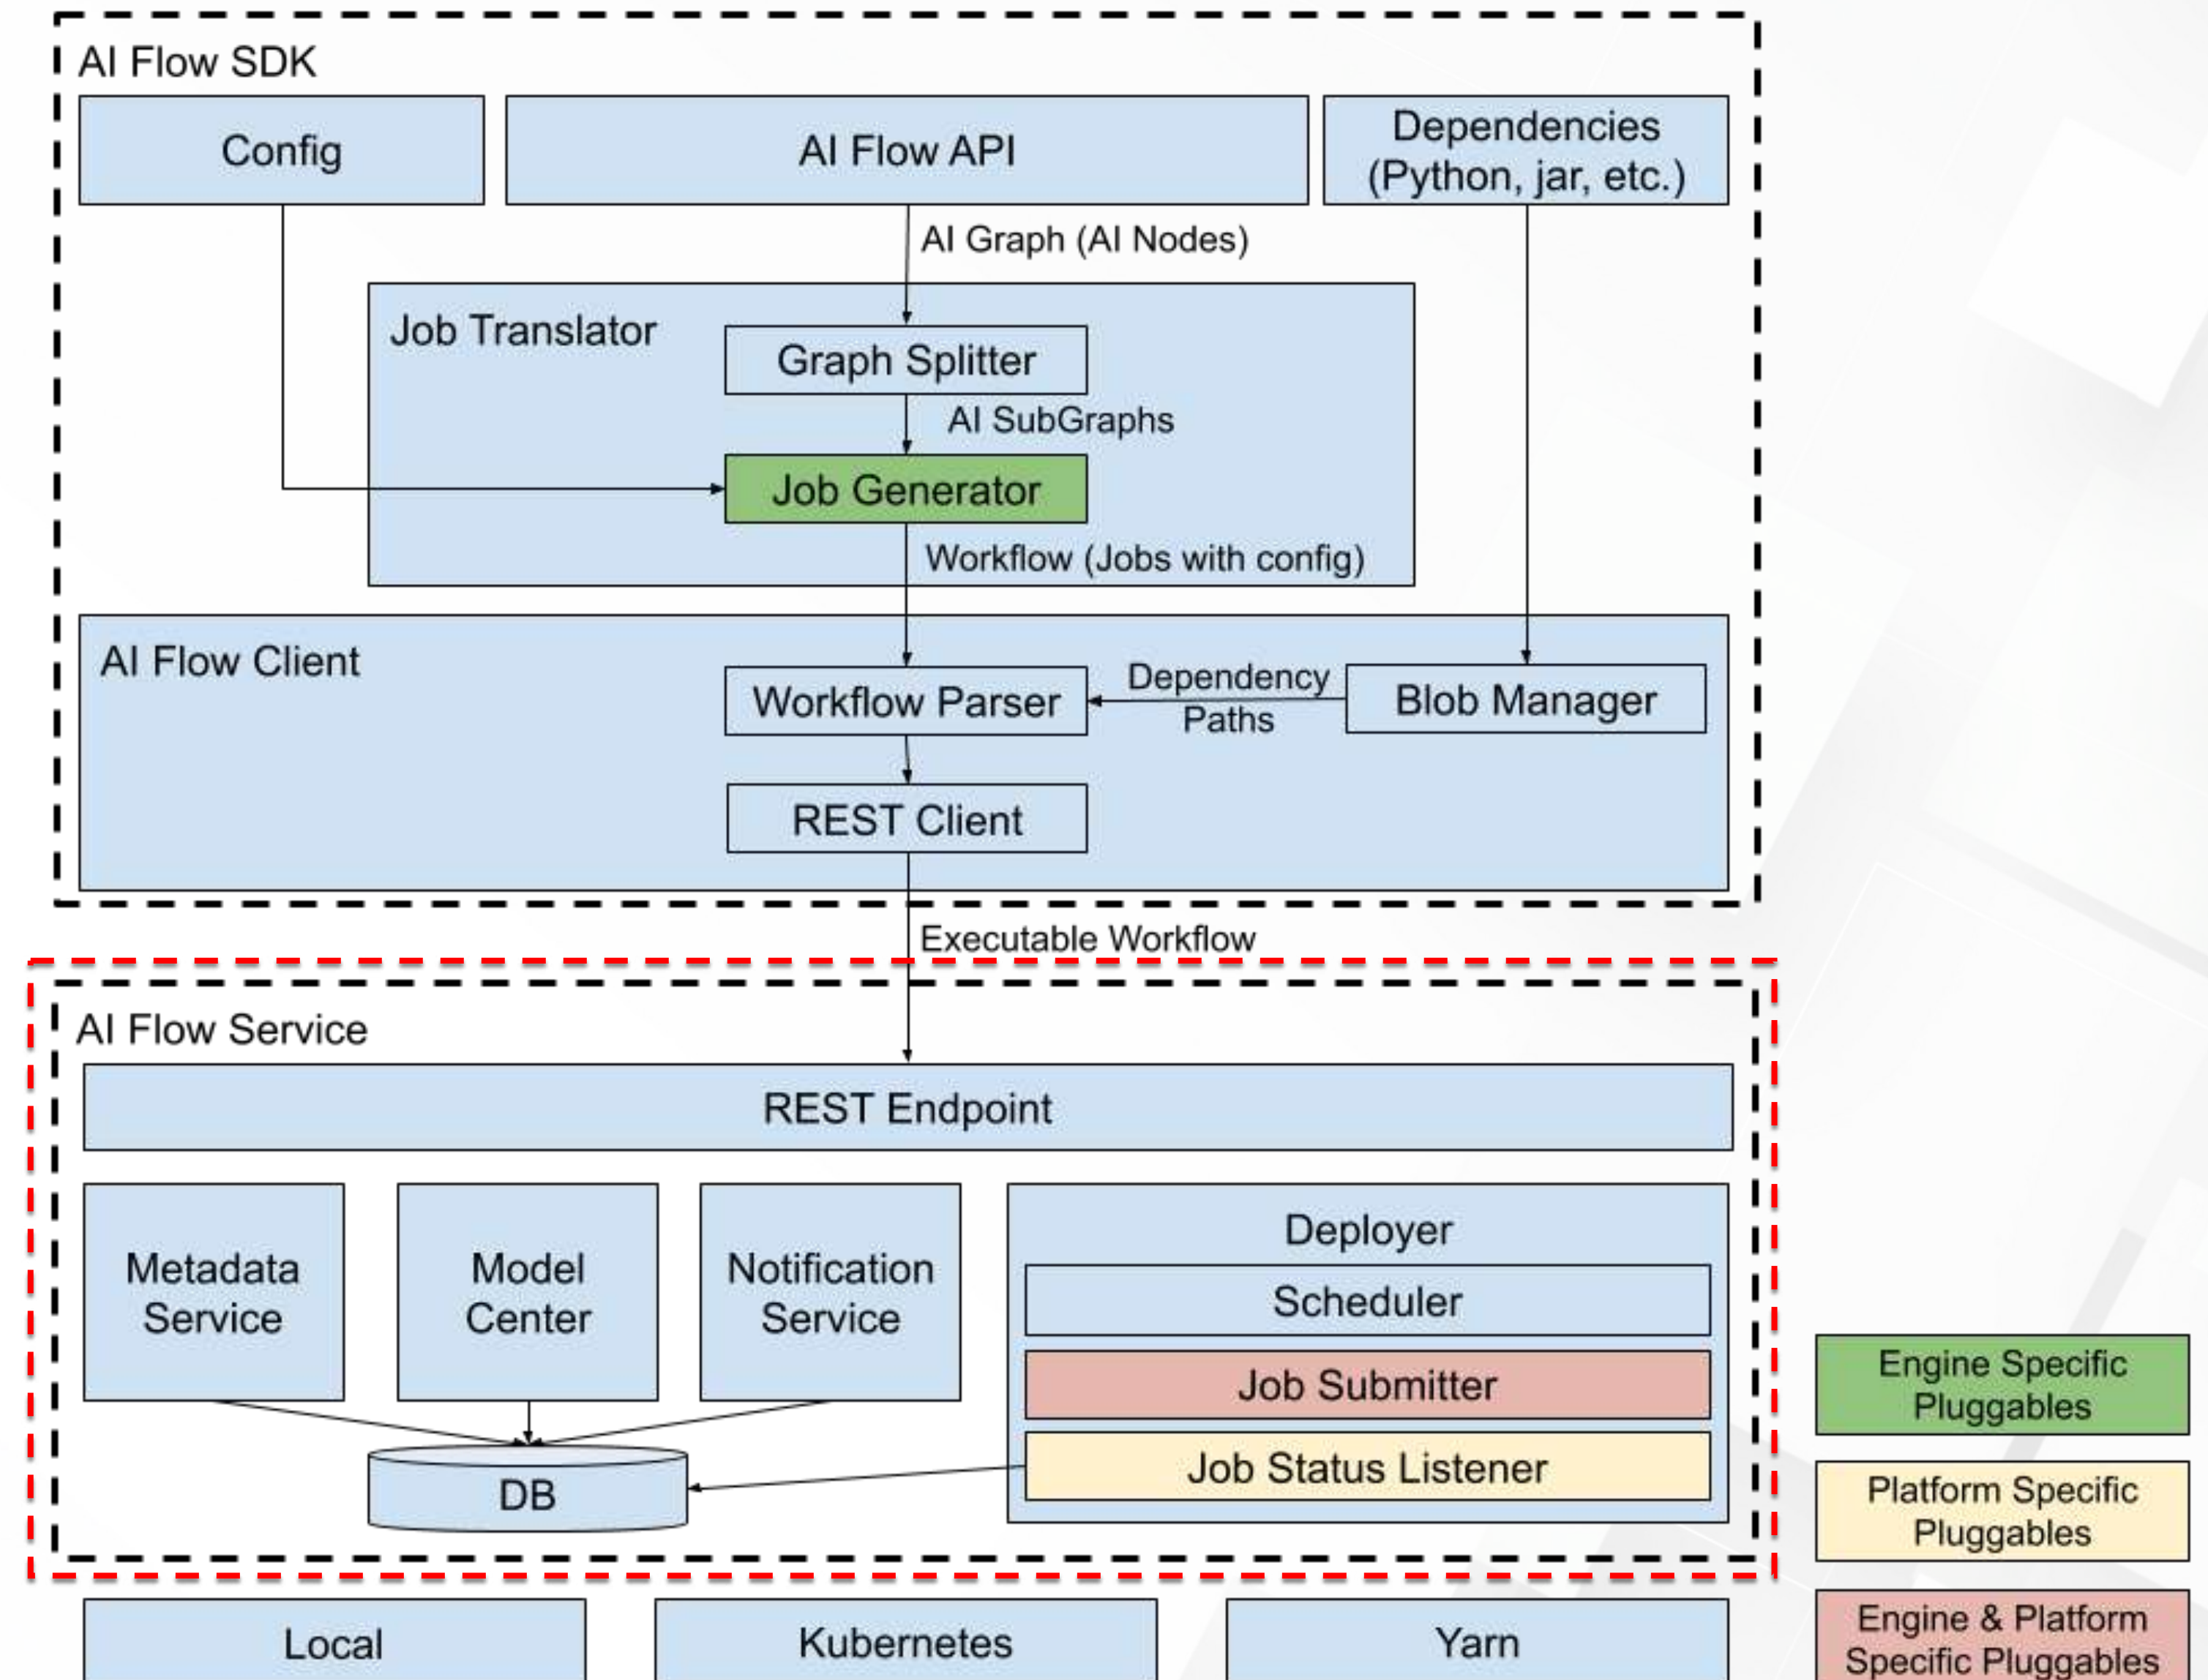
 Control Edge

 Message / Signal

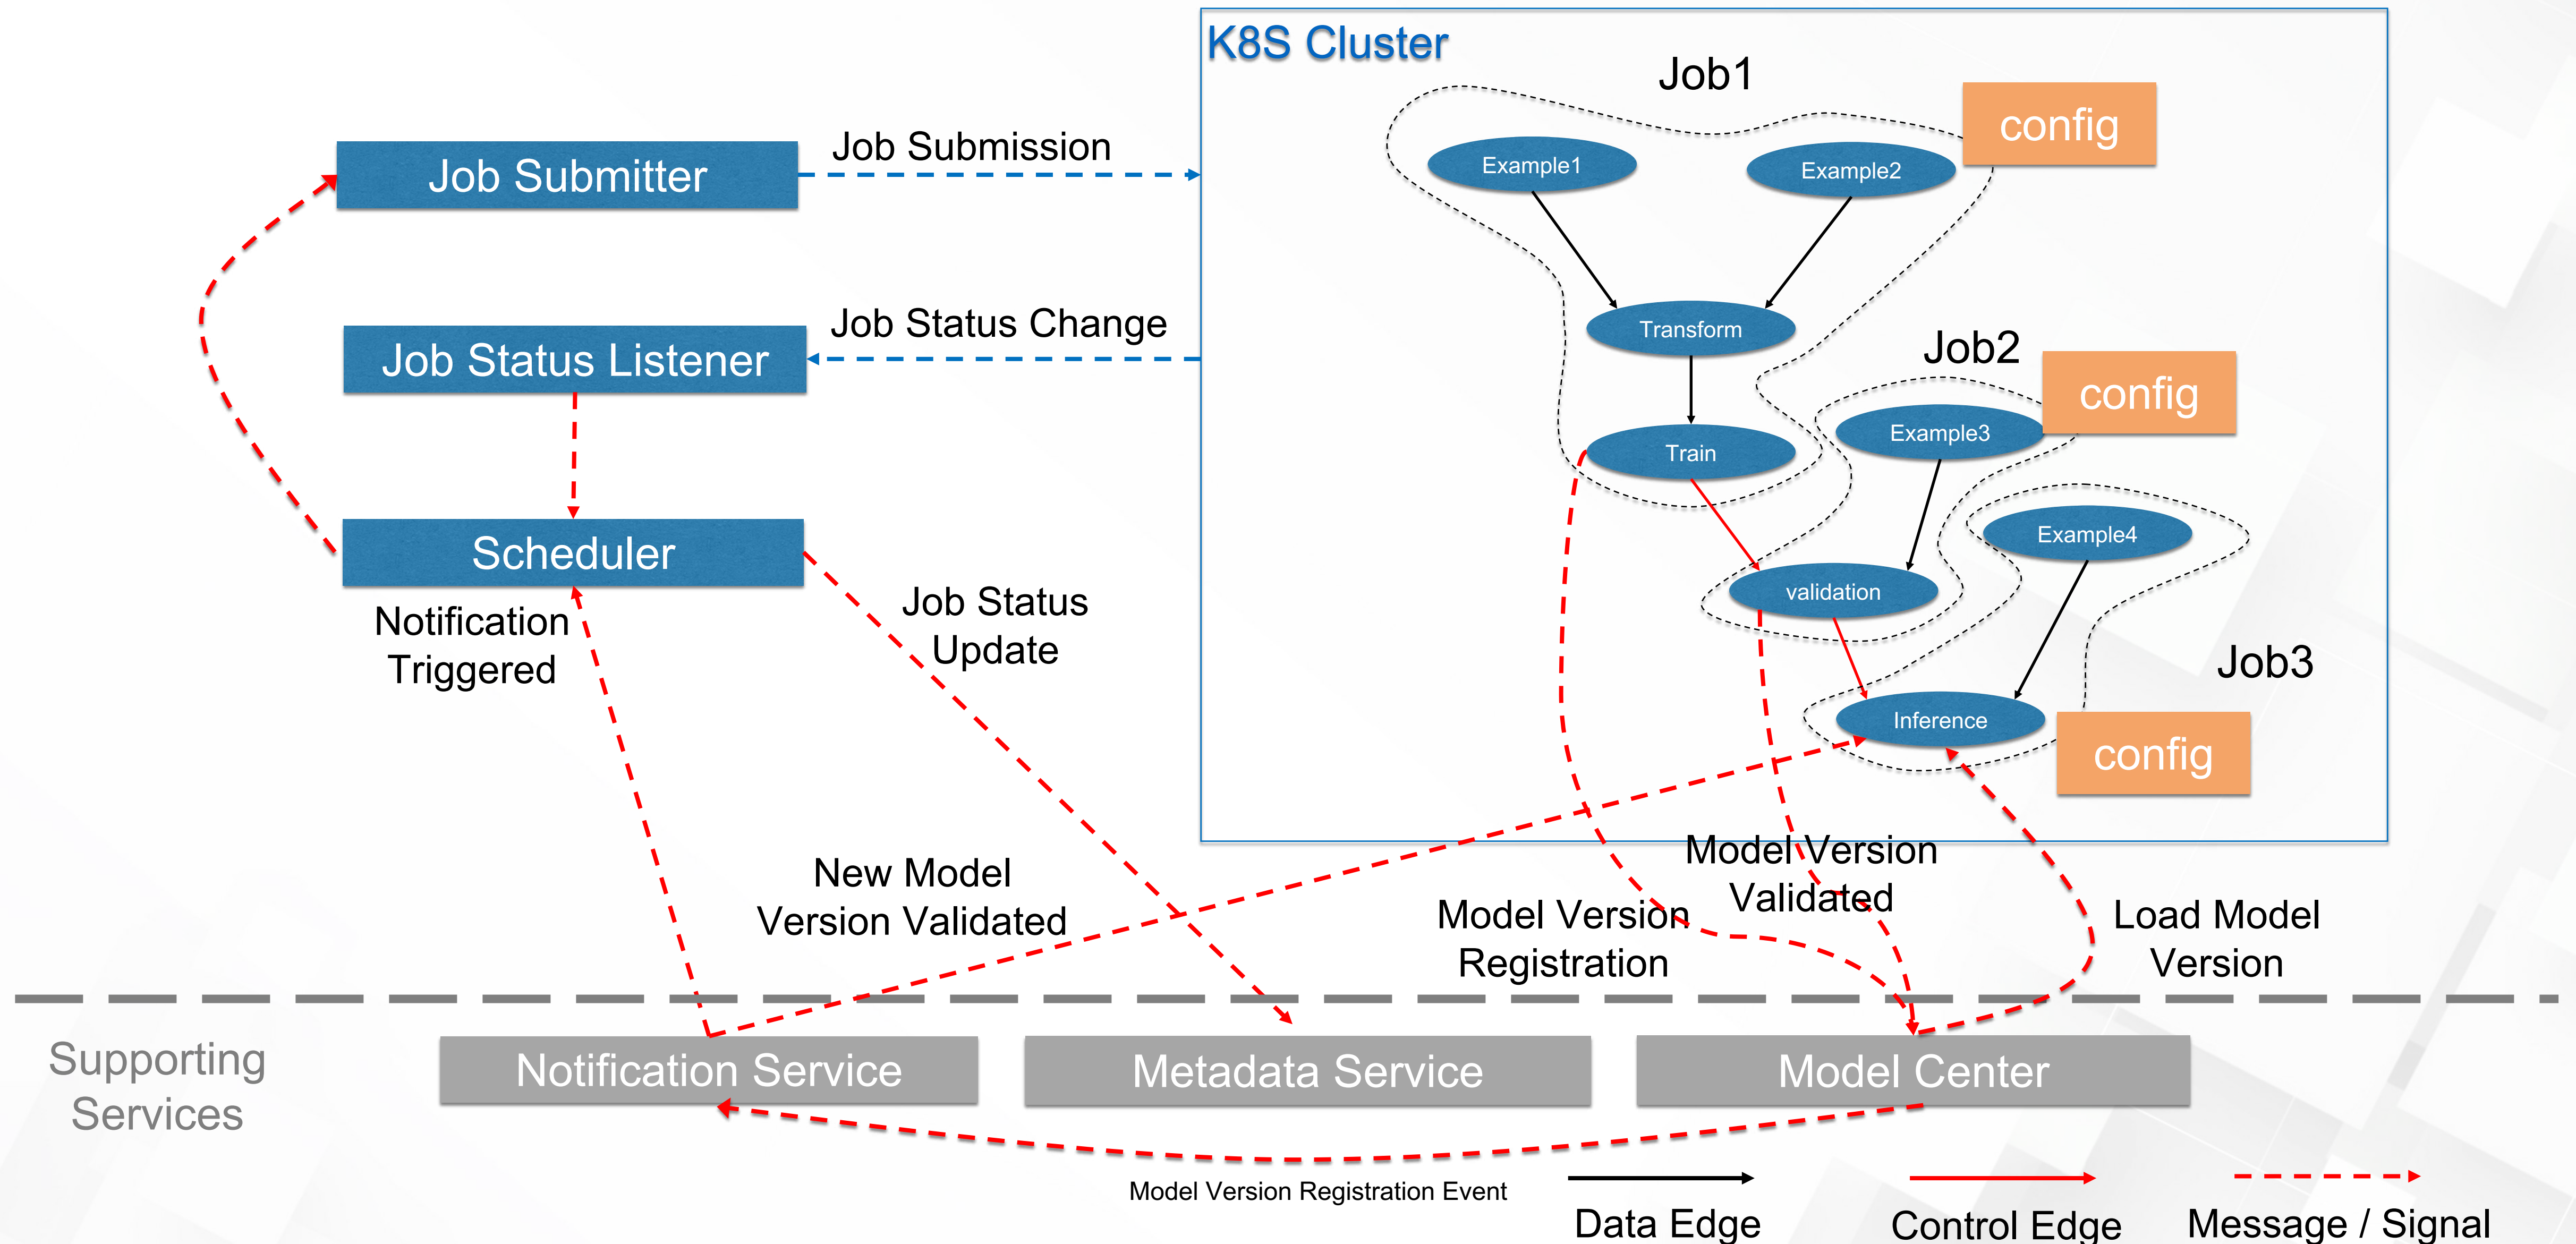
Workflow Execution

Run Executable Workflows

- Schedule jobs
- Submit jobs to different platforms



An Example



Flink AI Flow

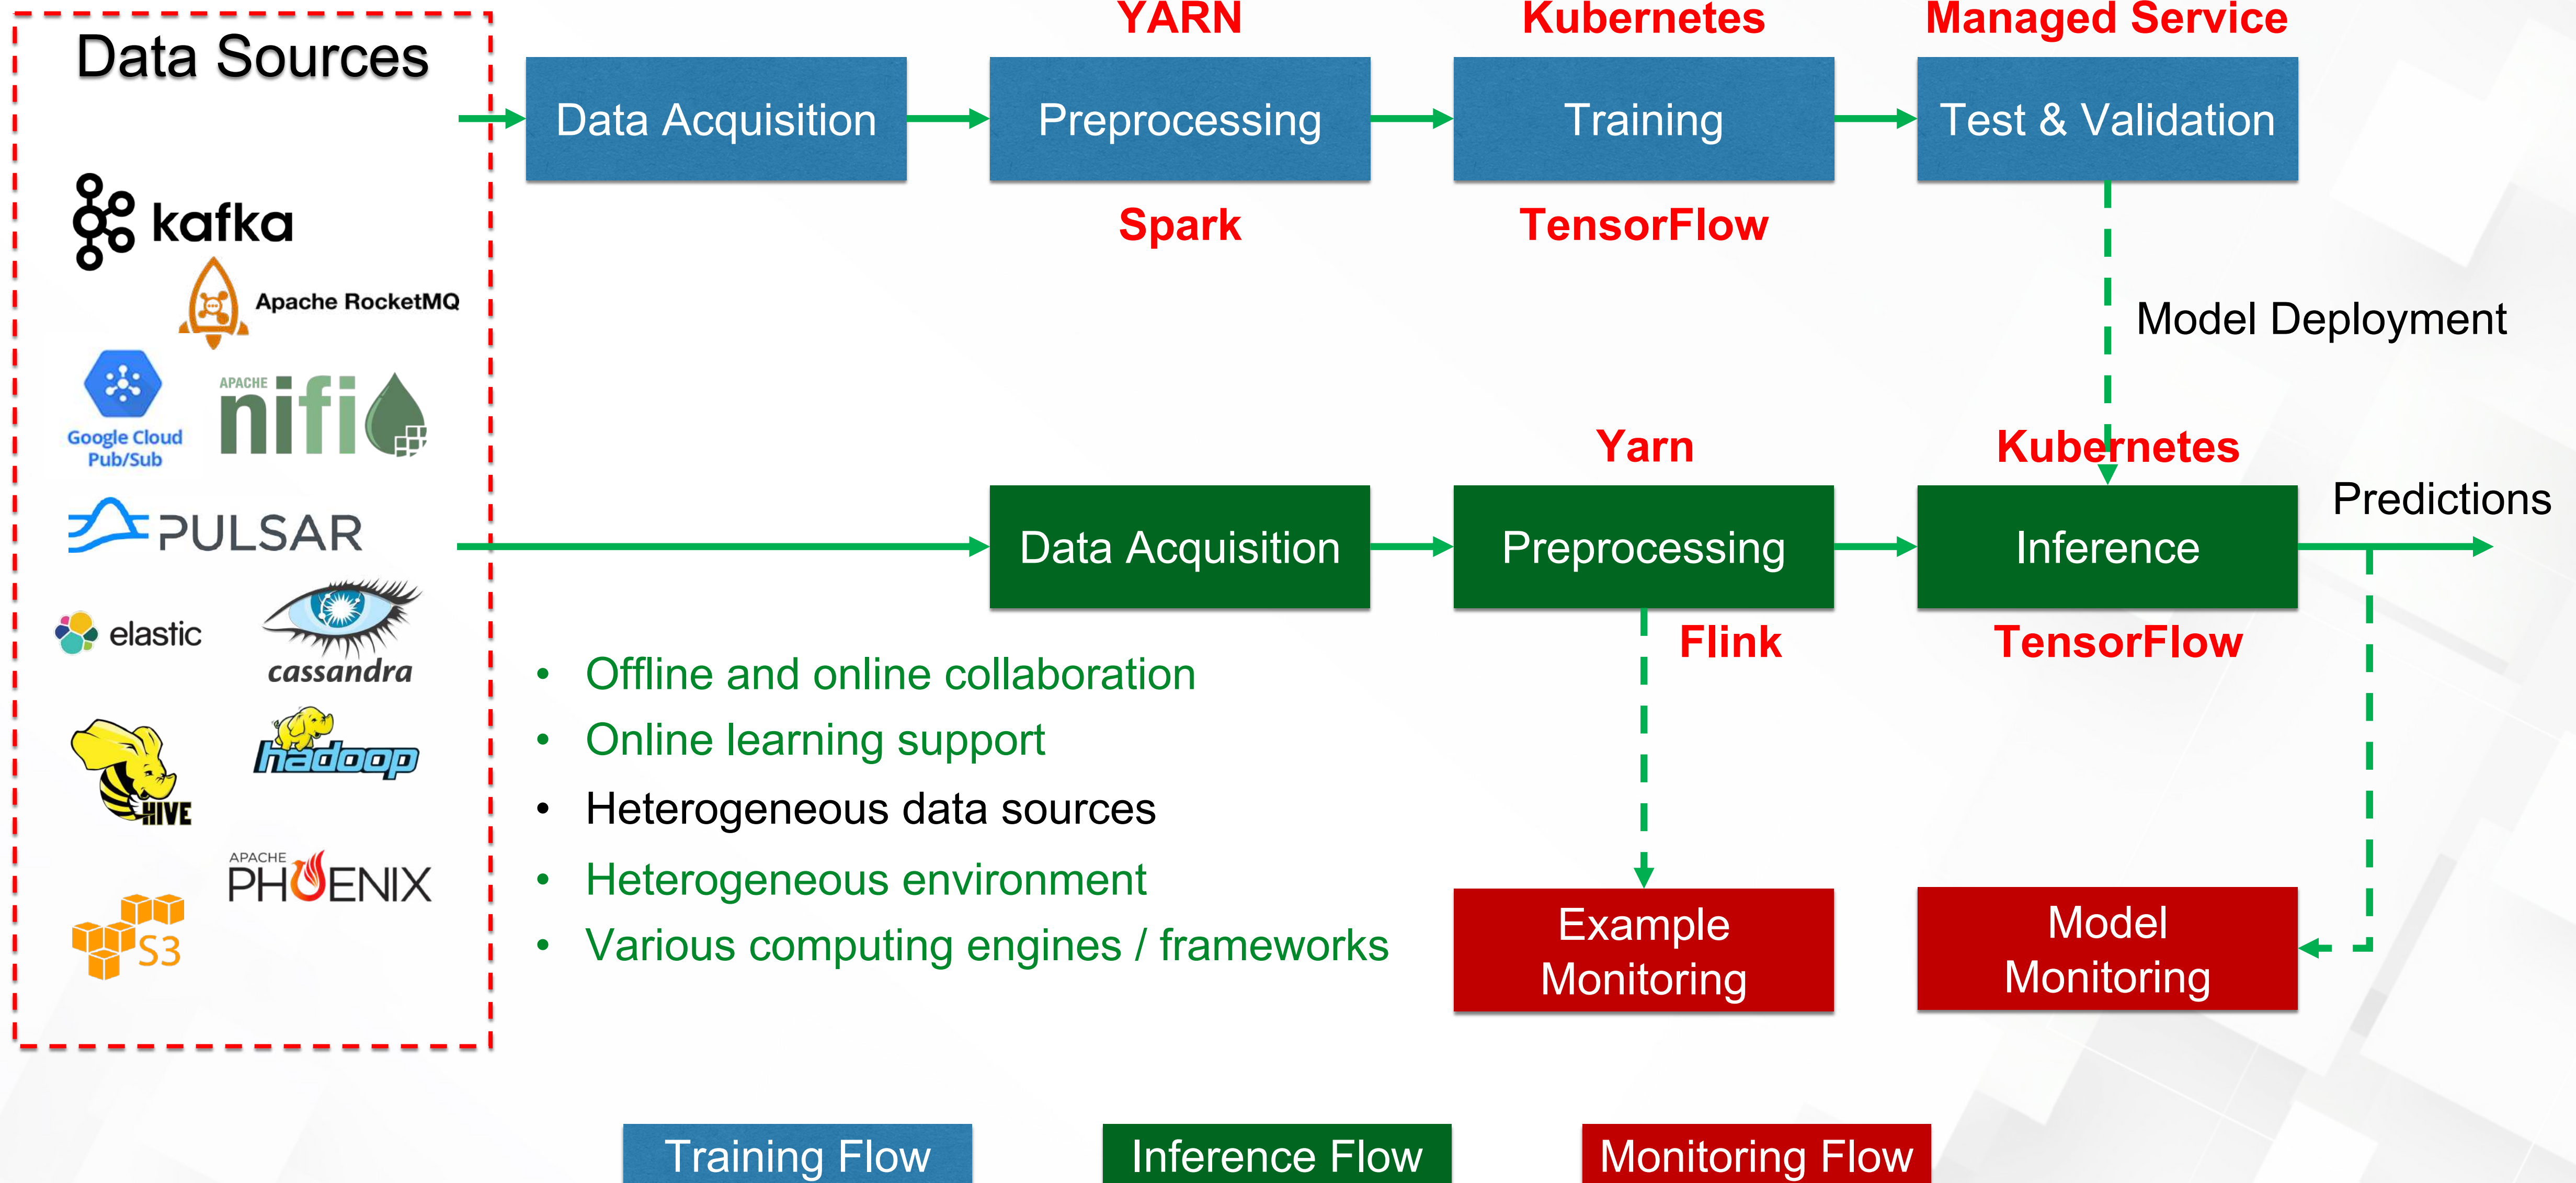
- Flink AI Flow
- Flink AI Flow & Flink ML Pipeline
- Flink AI Flow & PyFlink



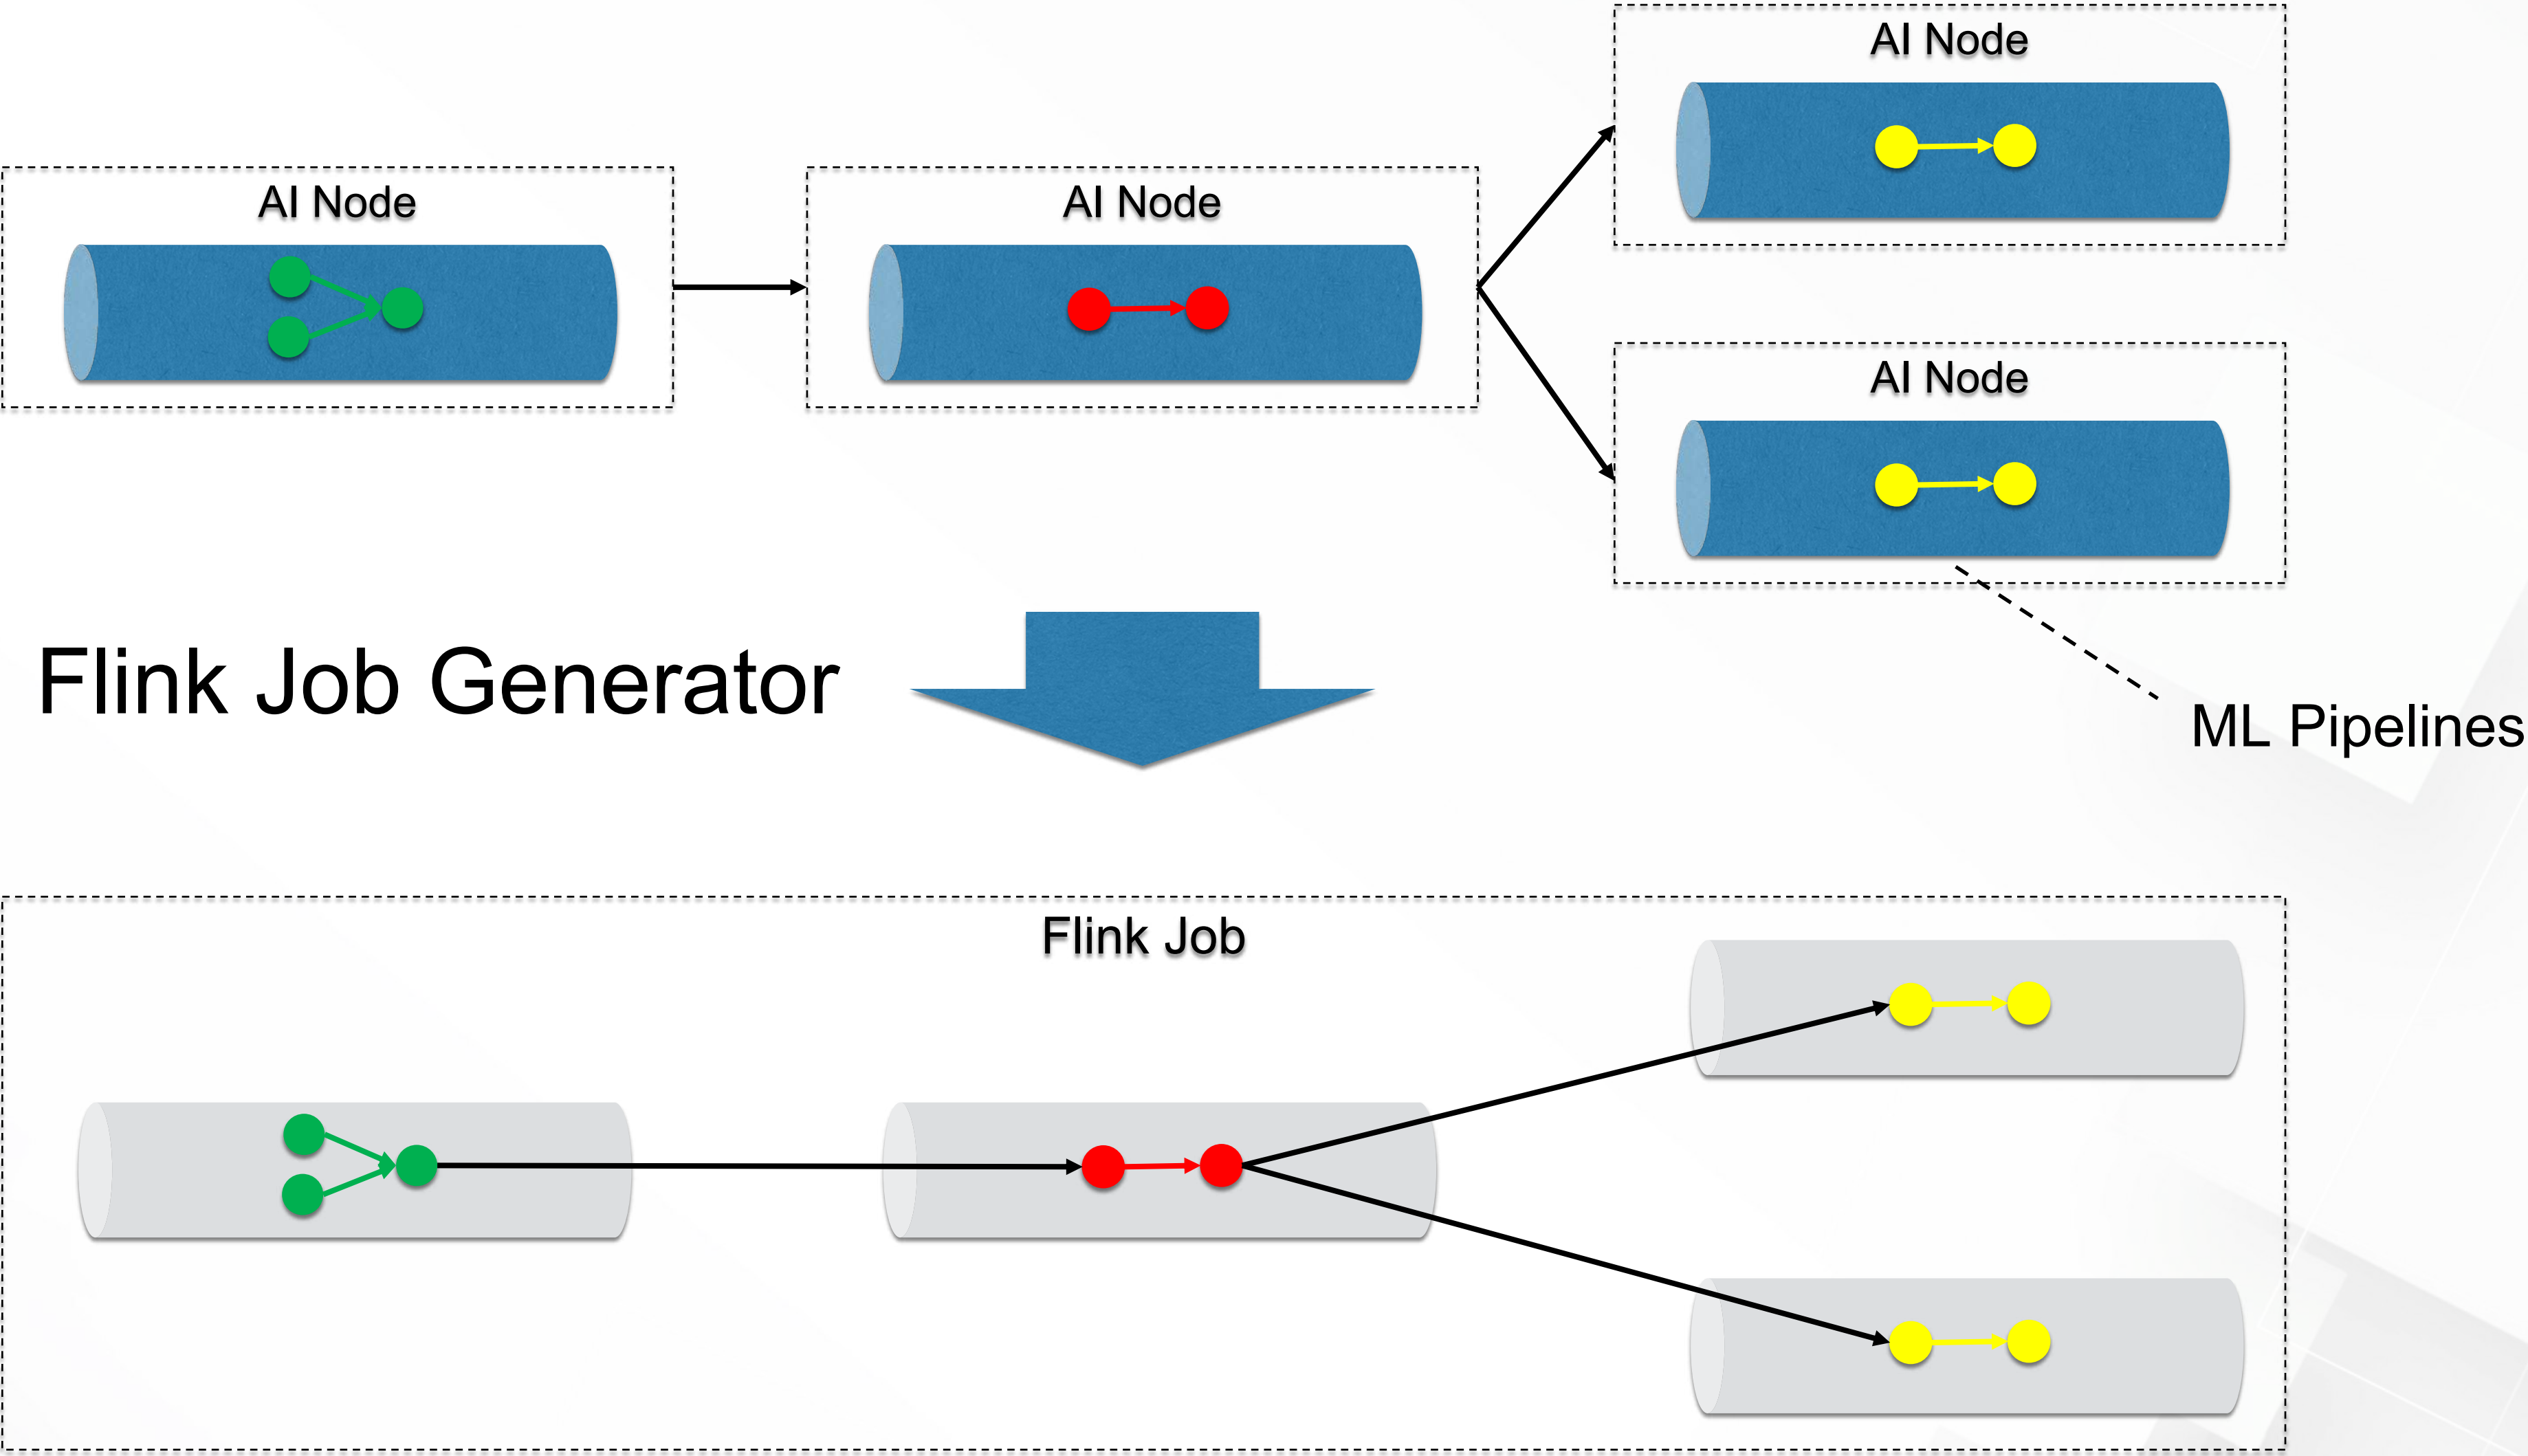
Flink AI Flow

- Implementation of AI Flow semantics
 - Flink Job Generator
 - Flink Job Submitter (Local / K8S)
 - ~1000 lines of code (easy to integrate)
- Work seamlessly with Flink Ecosystem
 - PyFlink
 - Flink ML Pipeline
 - TF on Flink


Flink AI Flow



Flink AI Flow & ML Pipeline




Flink AI Flow & PyFlink




engine = Python

As an ordinary Python function
(Users deal with IO)



engine = Flink

As chainable PyFlink Processing Functions
(Flink AI Flow deals with IO)



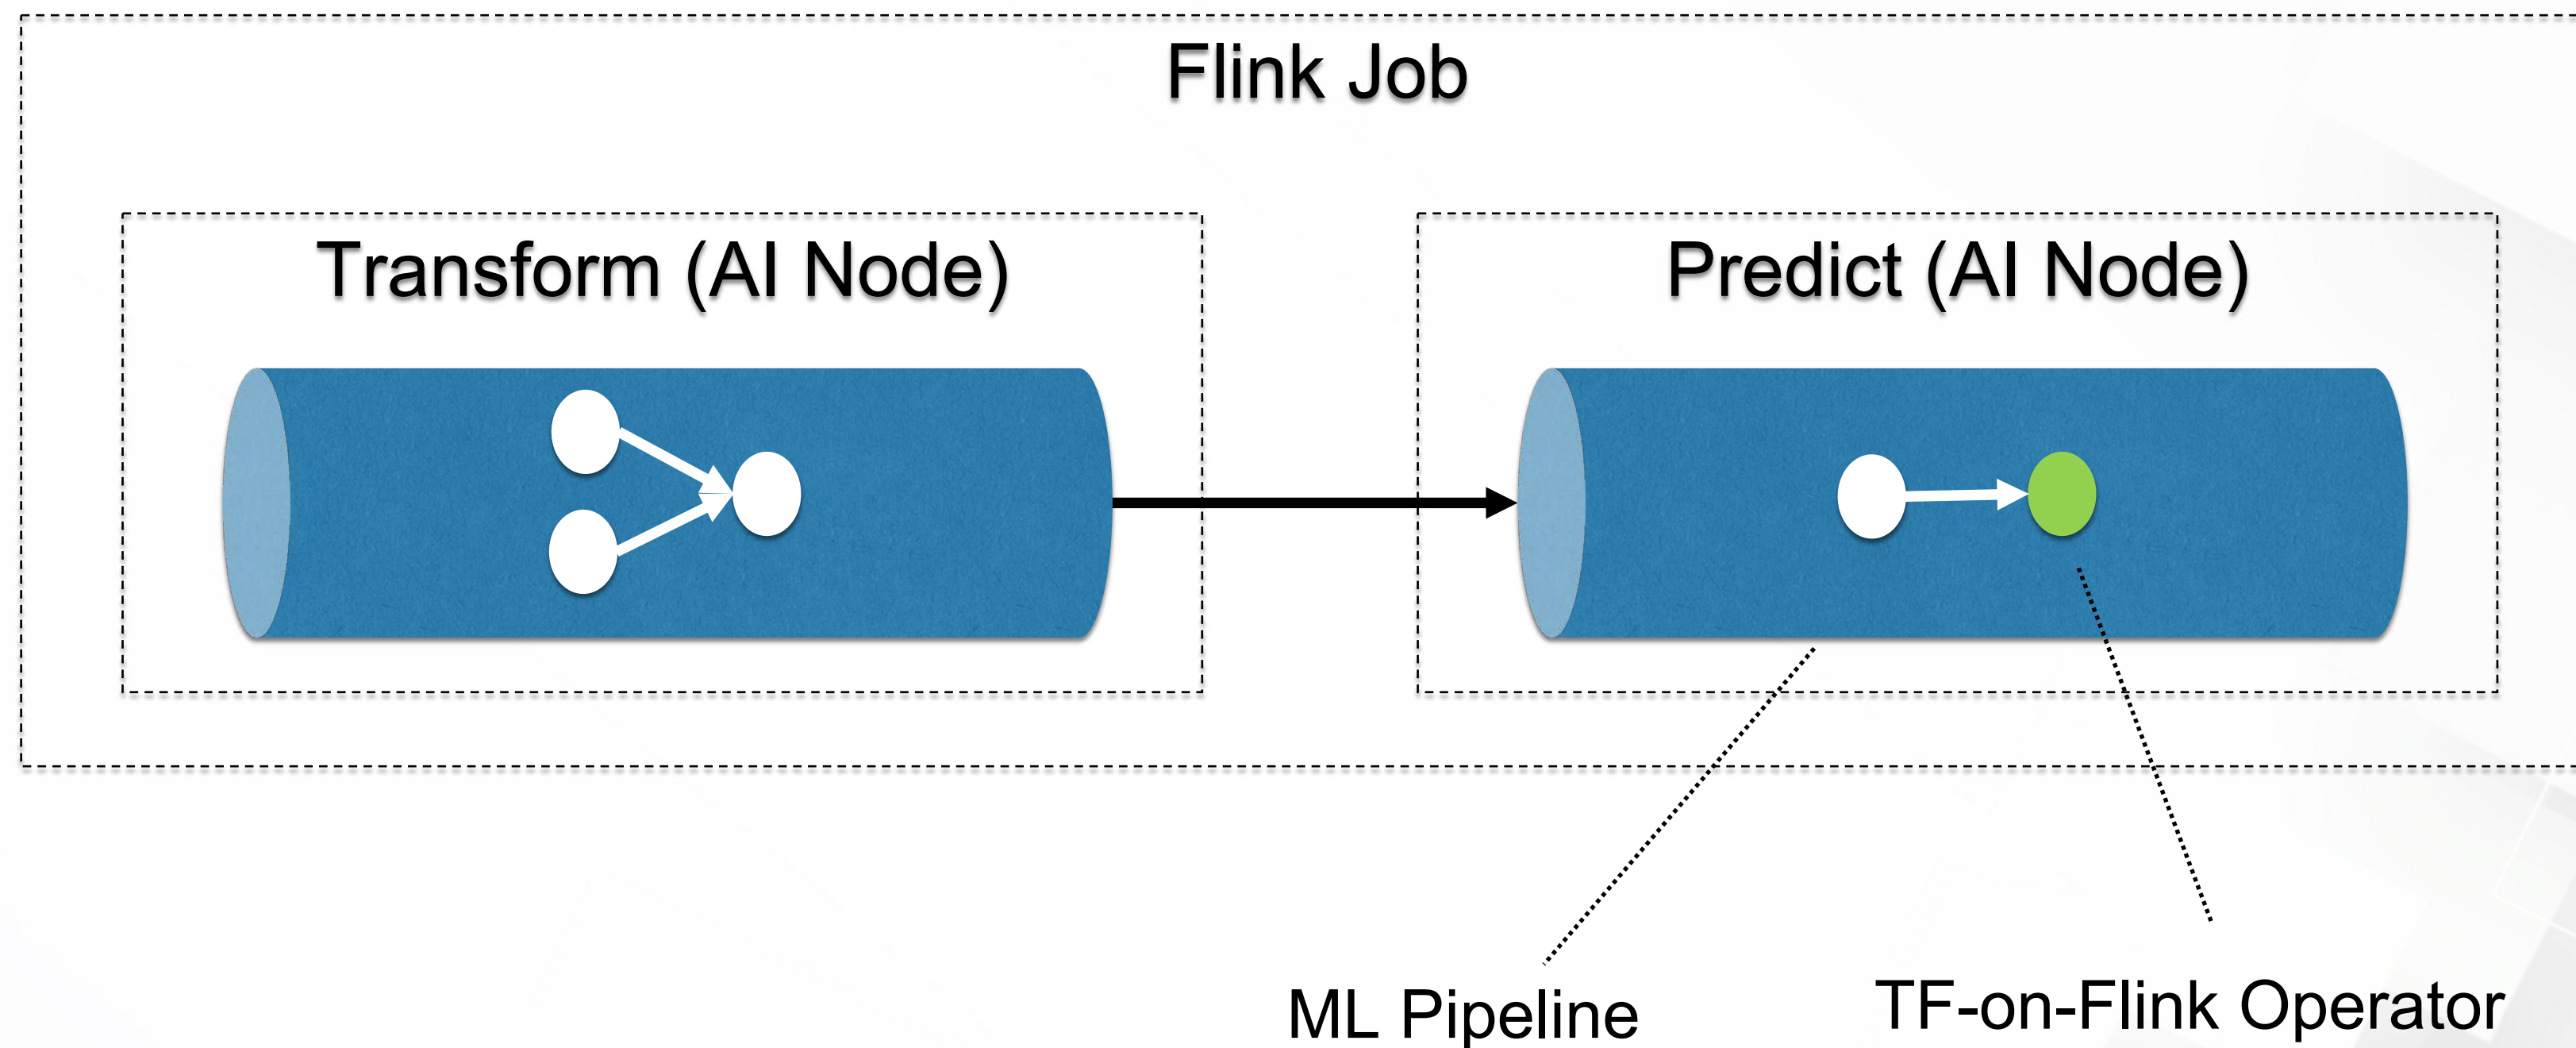
engine = Flink

As PyFlink ML Pipeline Stages
(Flink AI Flow deals with IO)

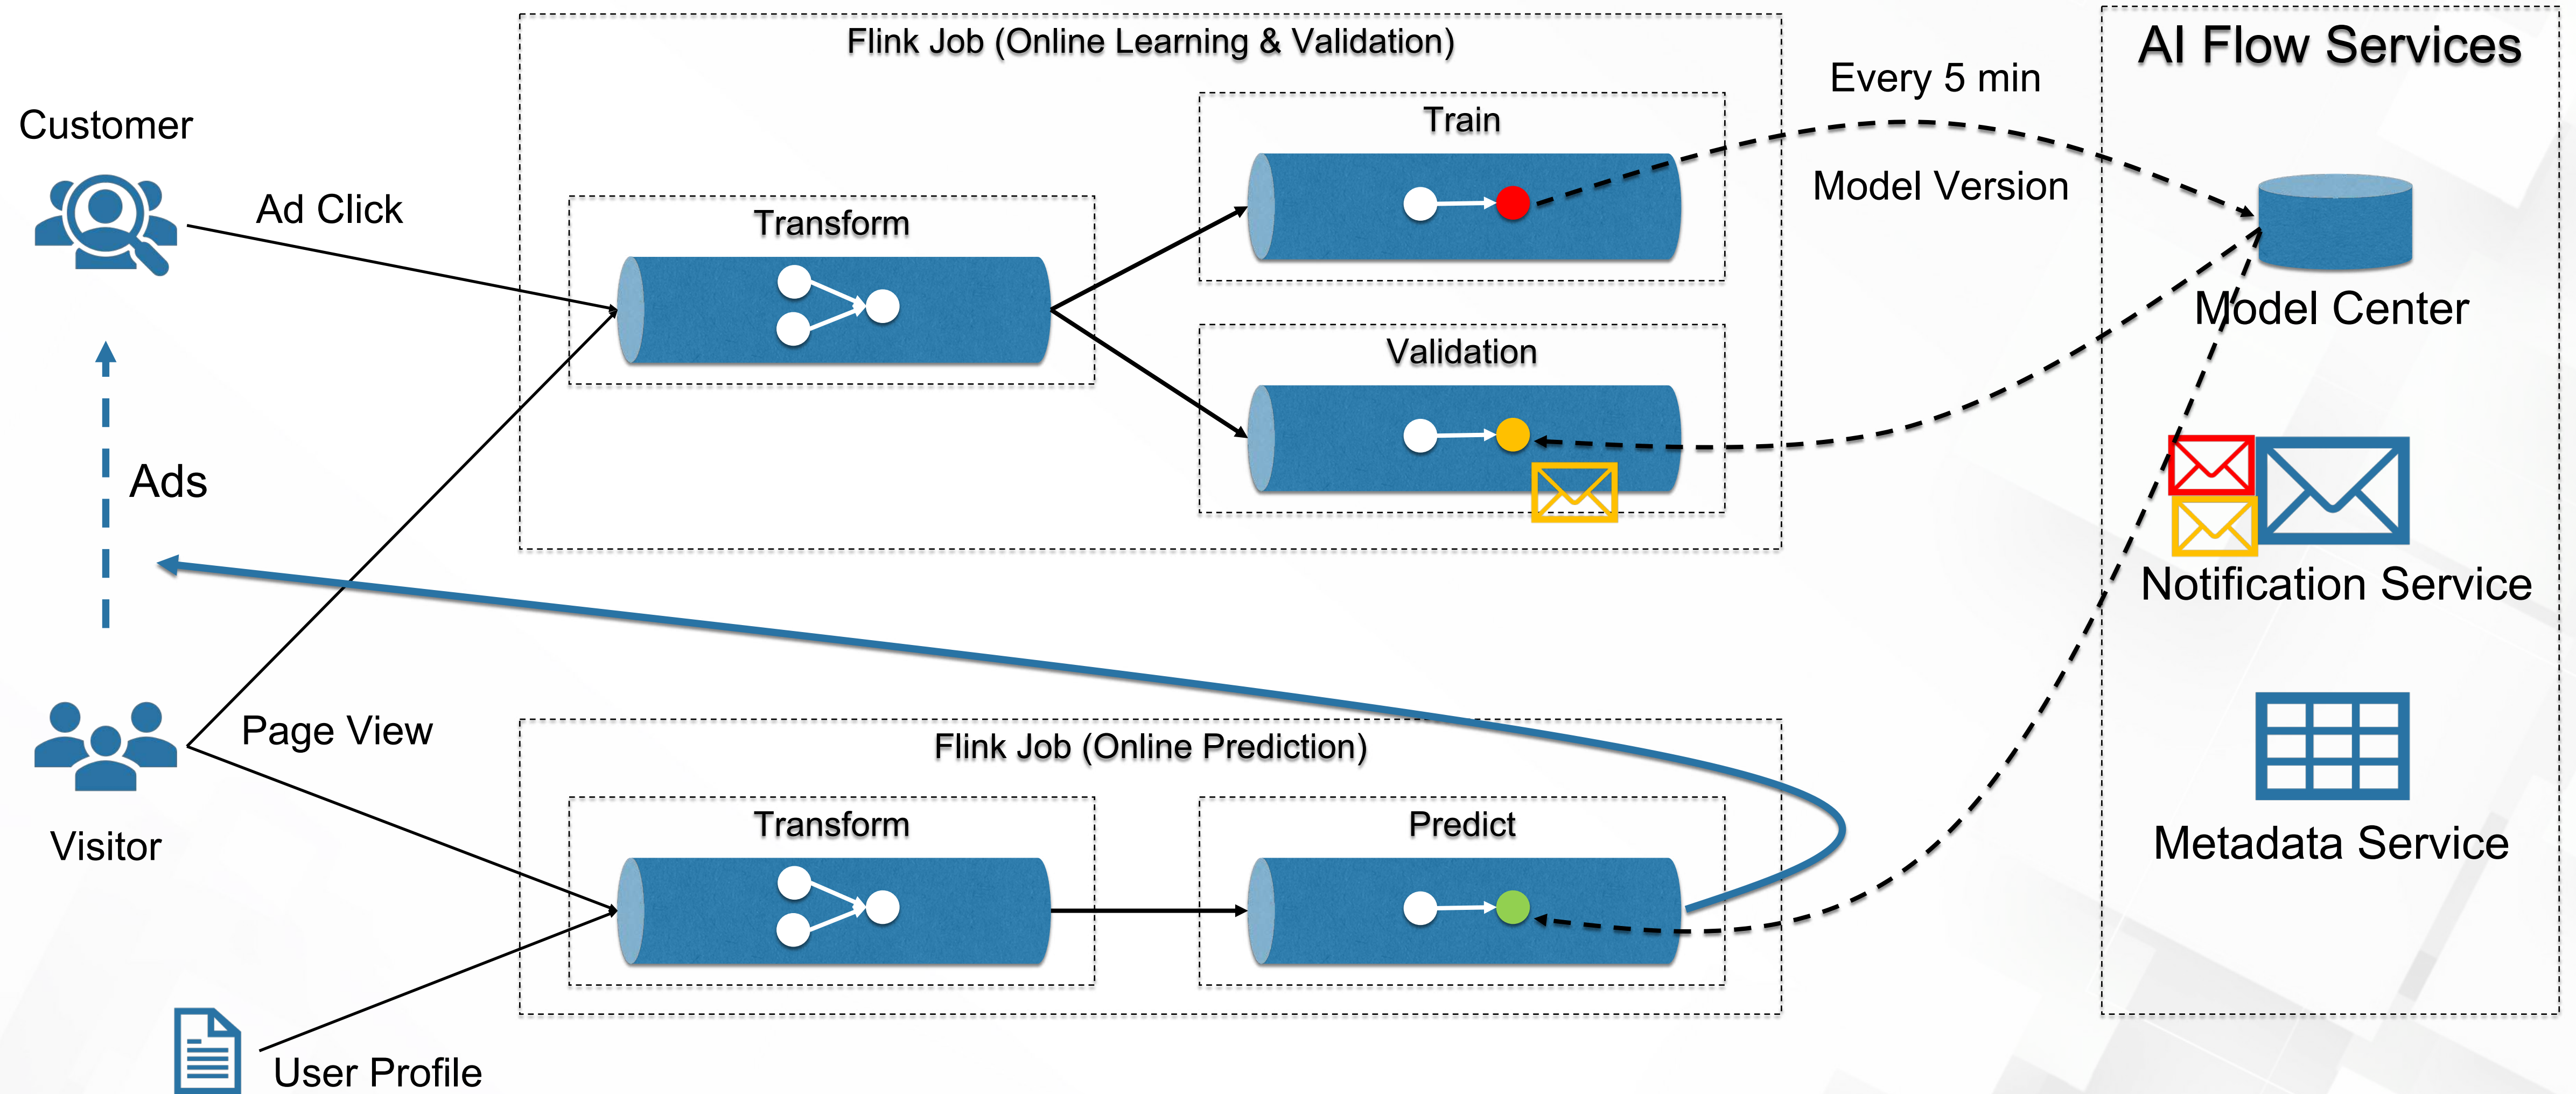
Available after Flink 1.11

Flink AI Flow & TF-on-Flink

- Recap for TF-on-Flink
 - TensorFlow as a Flink Operator
 - No intermediate storage between Flink and TF



An Example



Summary



Summary – AI Flow

- A **workflow** with **streaming job** support
- **Supporting services** to facilitate AI use cases
- Works perfectly well with Apache Flink
 - Also adaptive to other engines

AI Flow + Apache Flink
AI made easy!

Open Source Soon!

WE ARE HIRING!!!



Summer Interns: Scan QR Code

Full-Time: Send resume to jiangjie.qj@alibaba-inc.com

