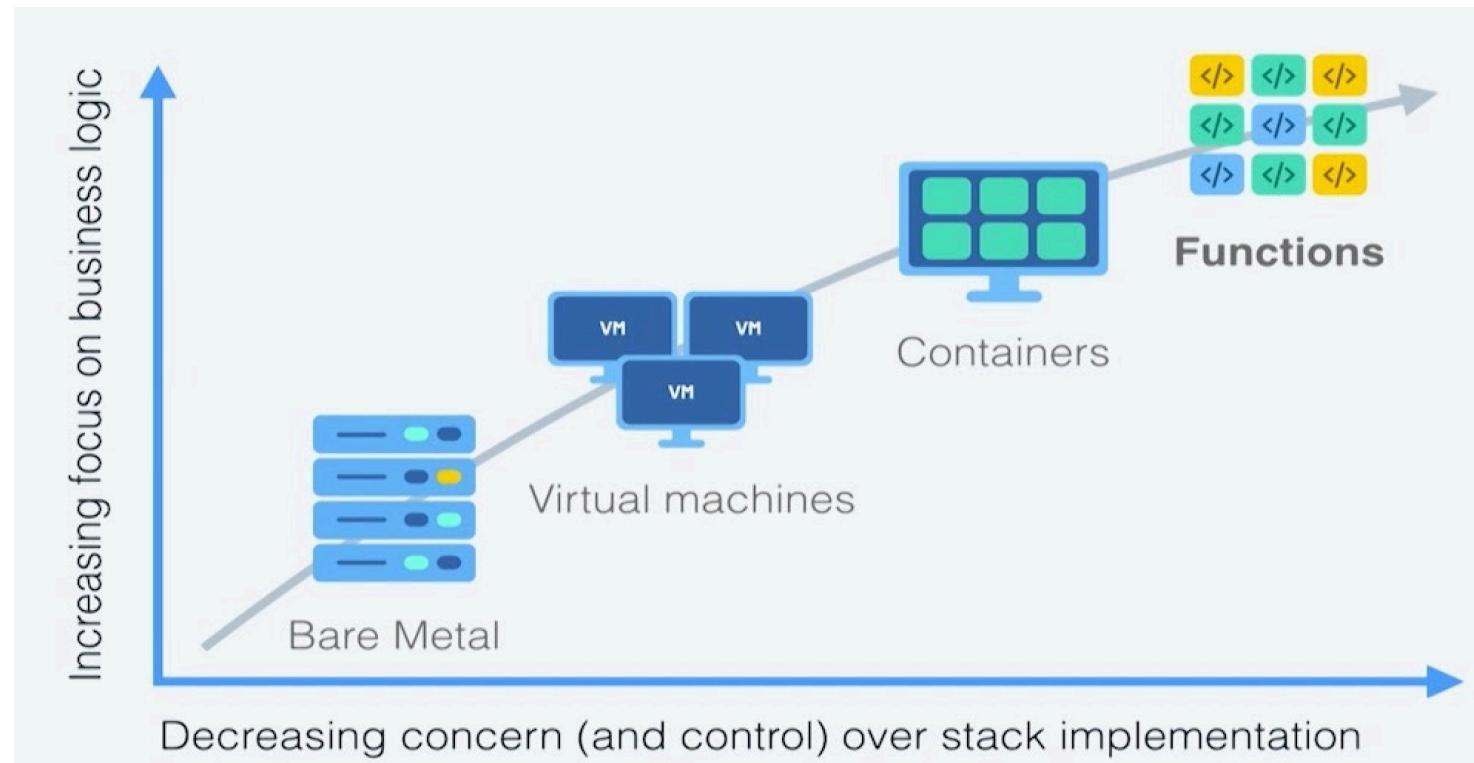


Stateful Functions 2.0

Stream Processing meets Serverless Apps

Serverless



Data Center

- Hardware based scaling
- Abstracts the physical hosting environment

IaaS

- Operating system based scaling
- Abstracts the hardware

PaaS

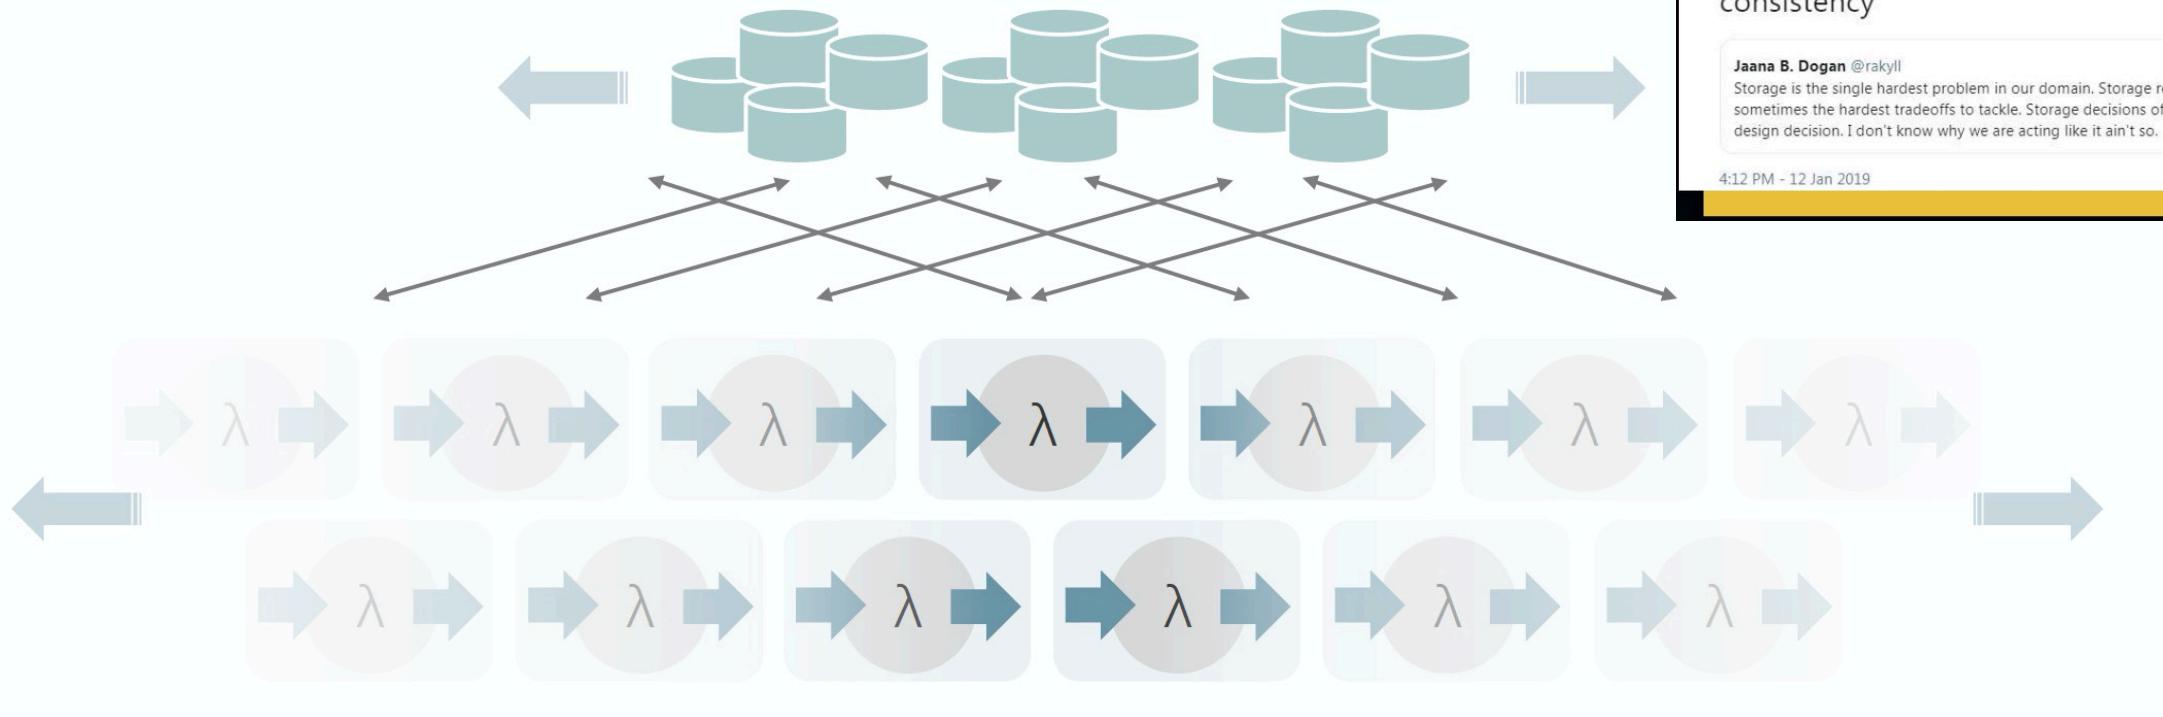
- Application based scaling
- Abstracts the operating system

Serverless

- Function based scaling
- Abstracts the language runtime



Problem of data handling



tim gross
@0x74696d

Follow

OMG yes! So much energy being poured into orchestrating stateless applications. That isn't *totally* trivial but it's pretty damn close relative to state and storage. And application devs too often pretend selecting a RDBMS means they don't have to worry about state consistency

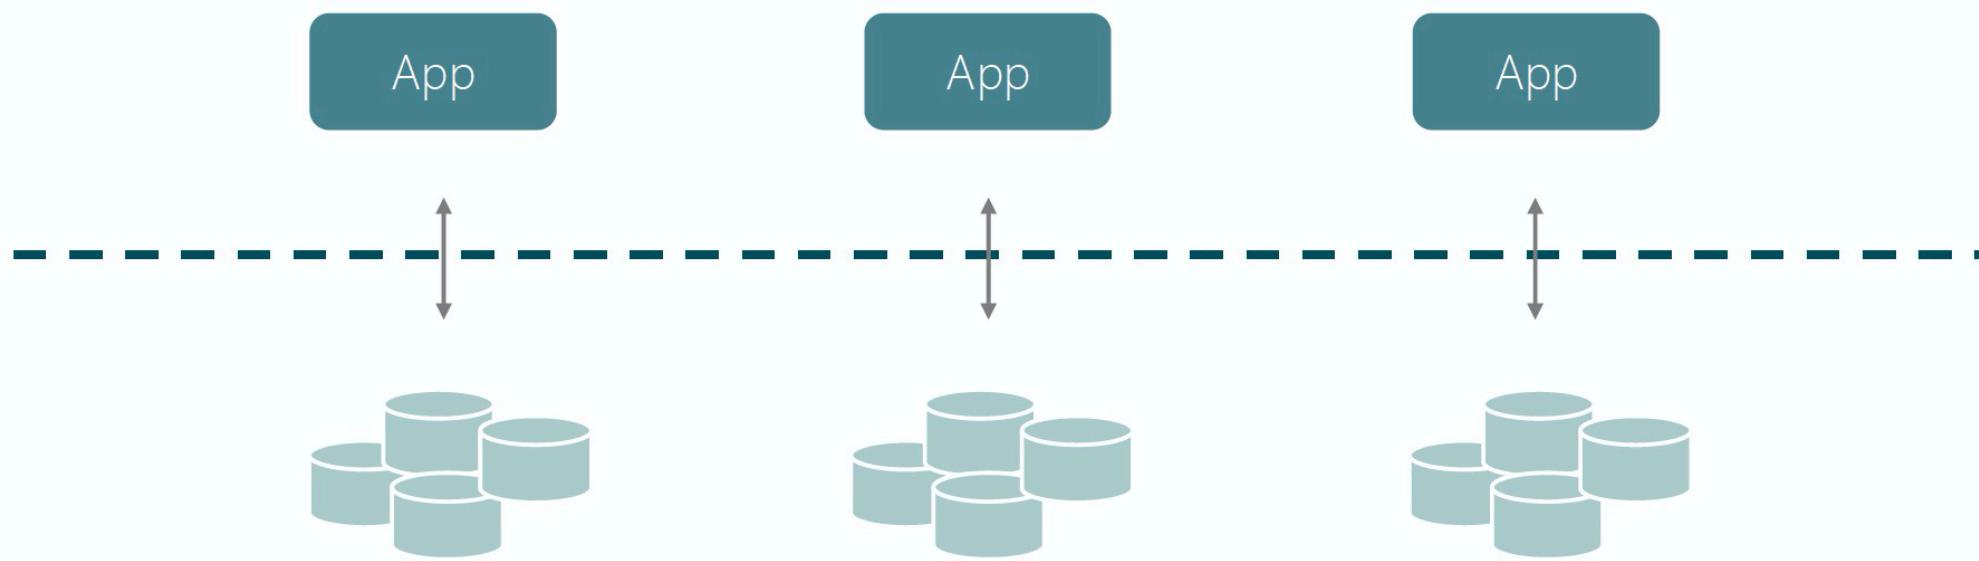
Jaana B. Dogan @rakyll

Storage is the single hardest problem in our domain. Storage related tradeoffs are sometimes the hardest tradeoffs to tackle. Storage decisions often impact every other design decision. I don't know why we are acting like it ain't so.

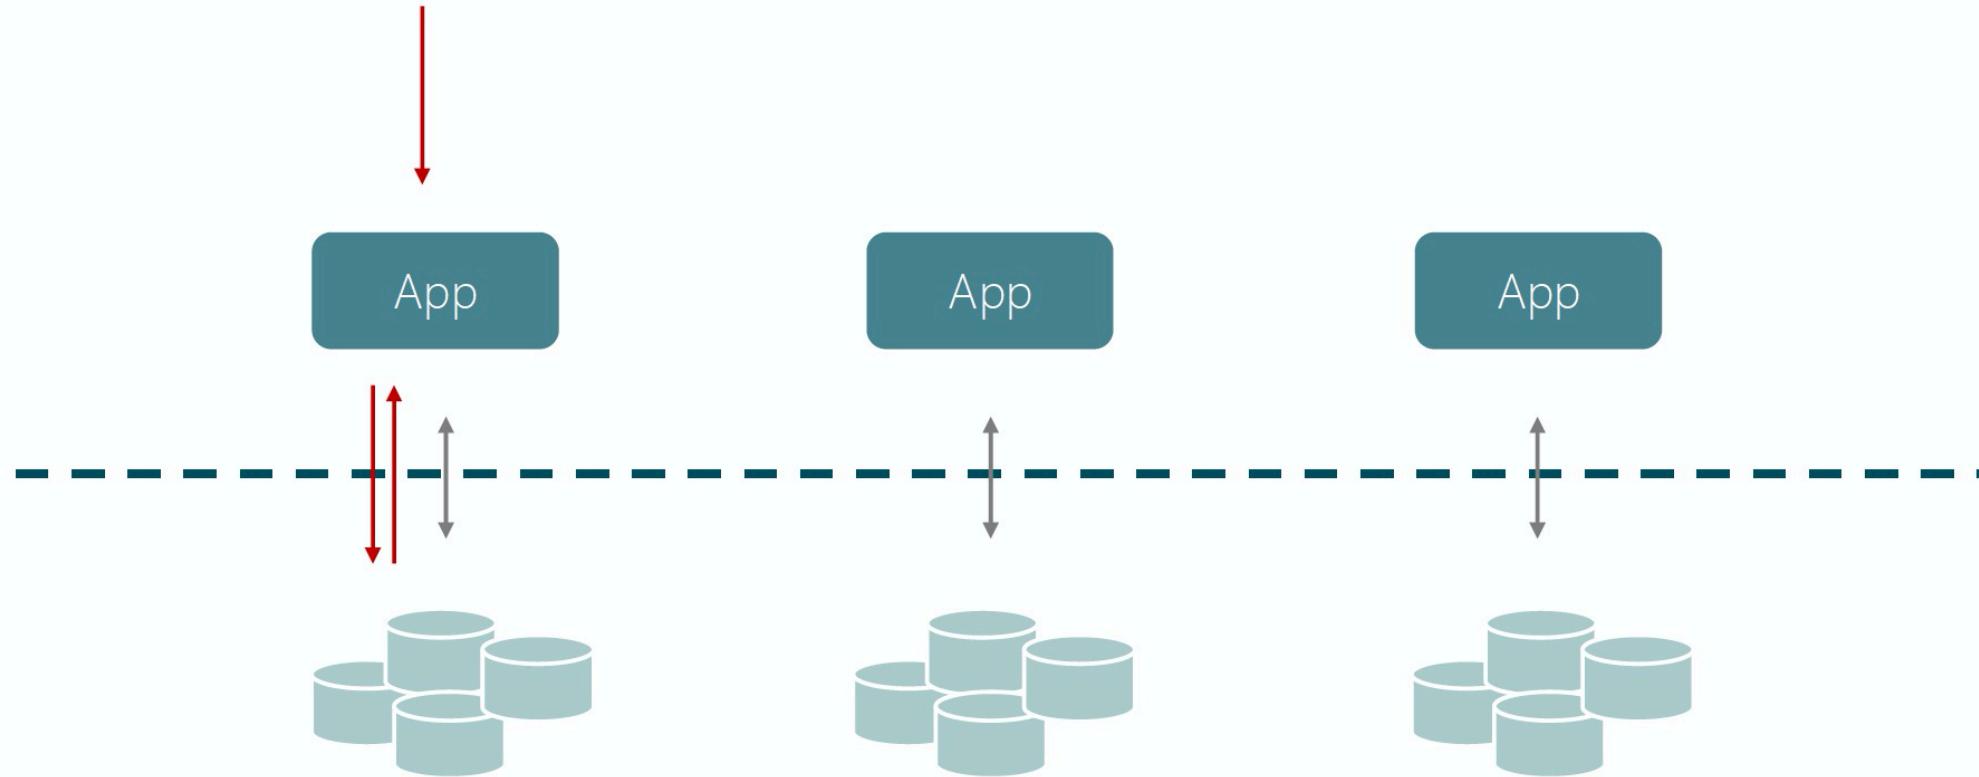
4:12 PM - 12 Jan 2019



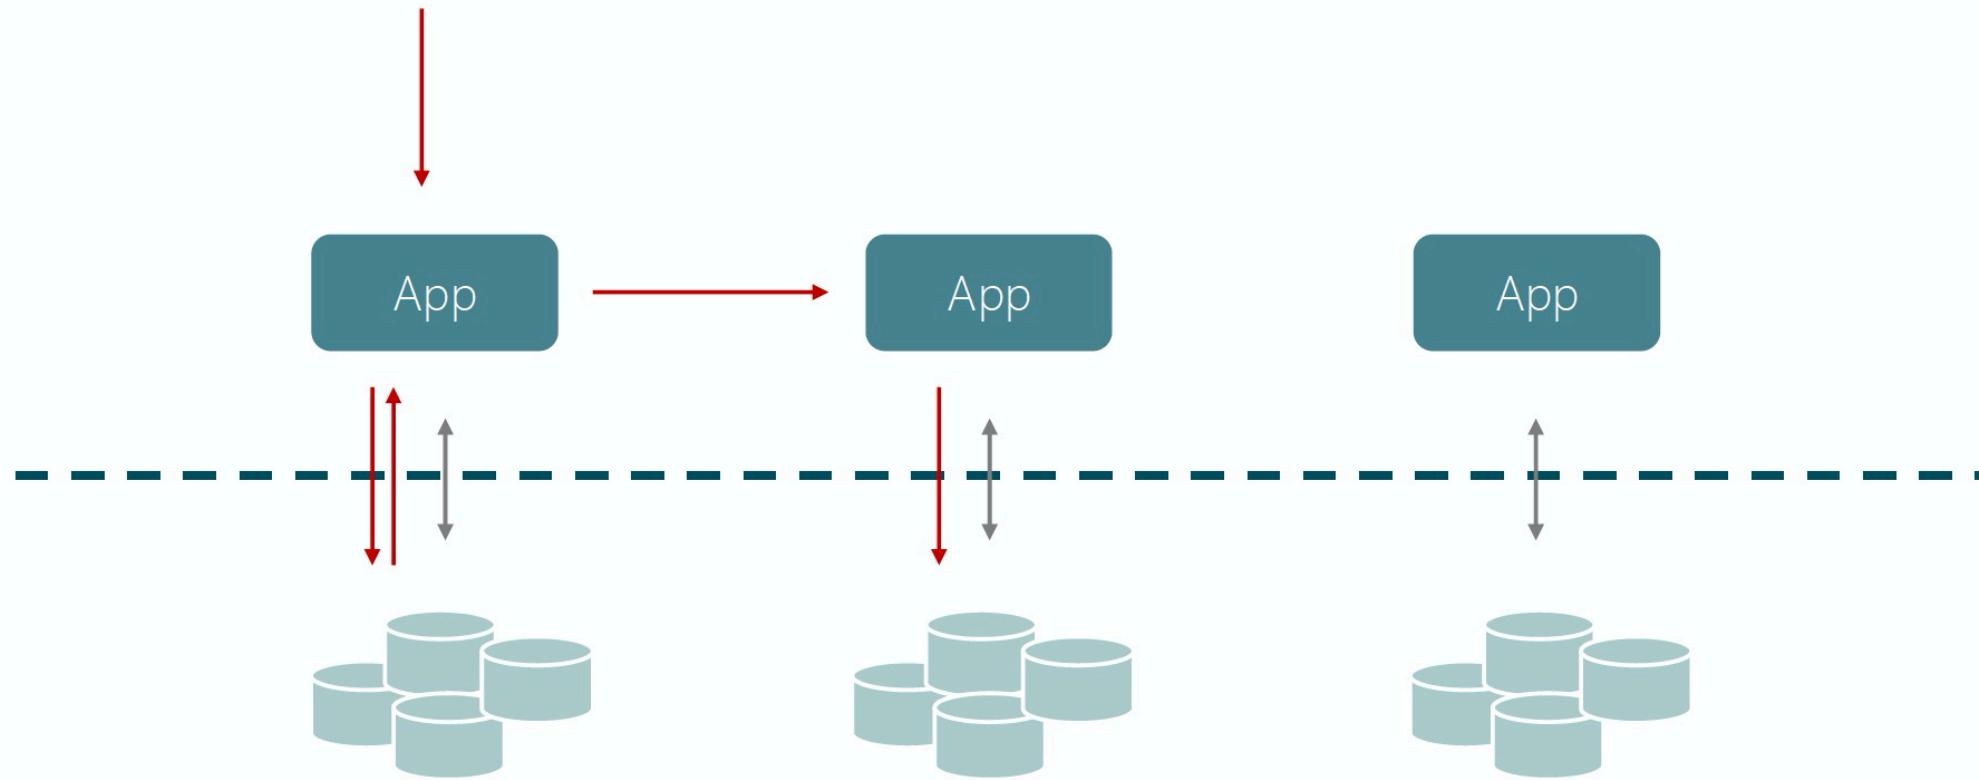
The Applications handles messaging (and consistency)



The Applications handles messaging (and consistency)

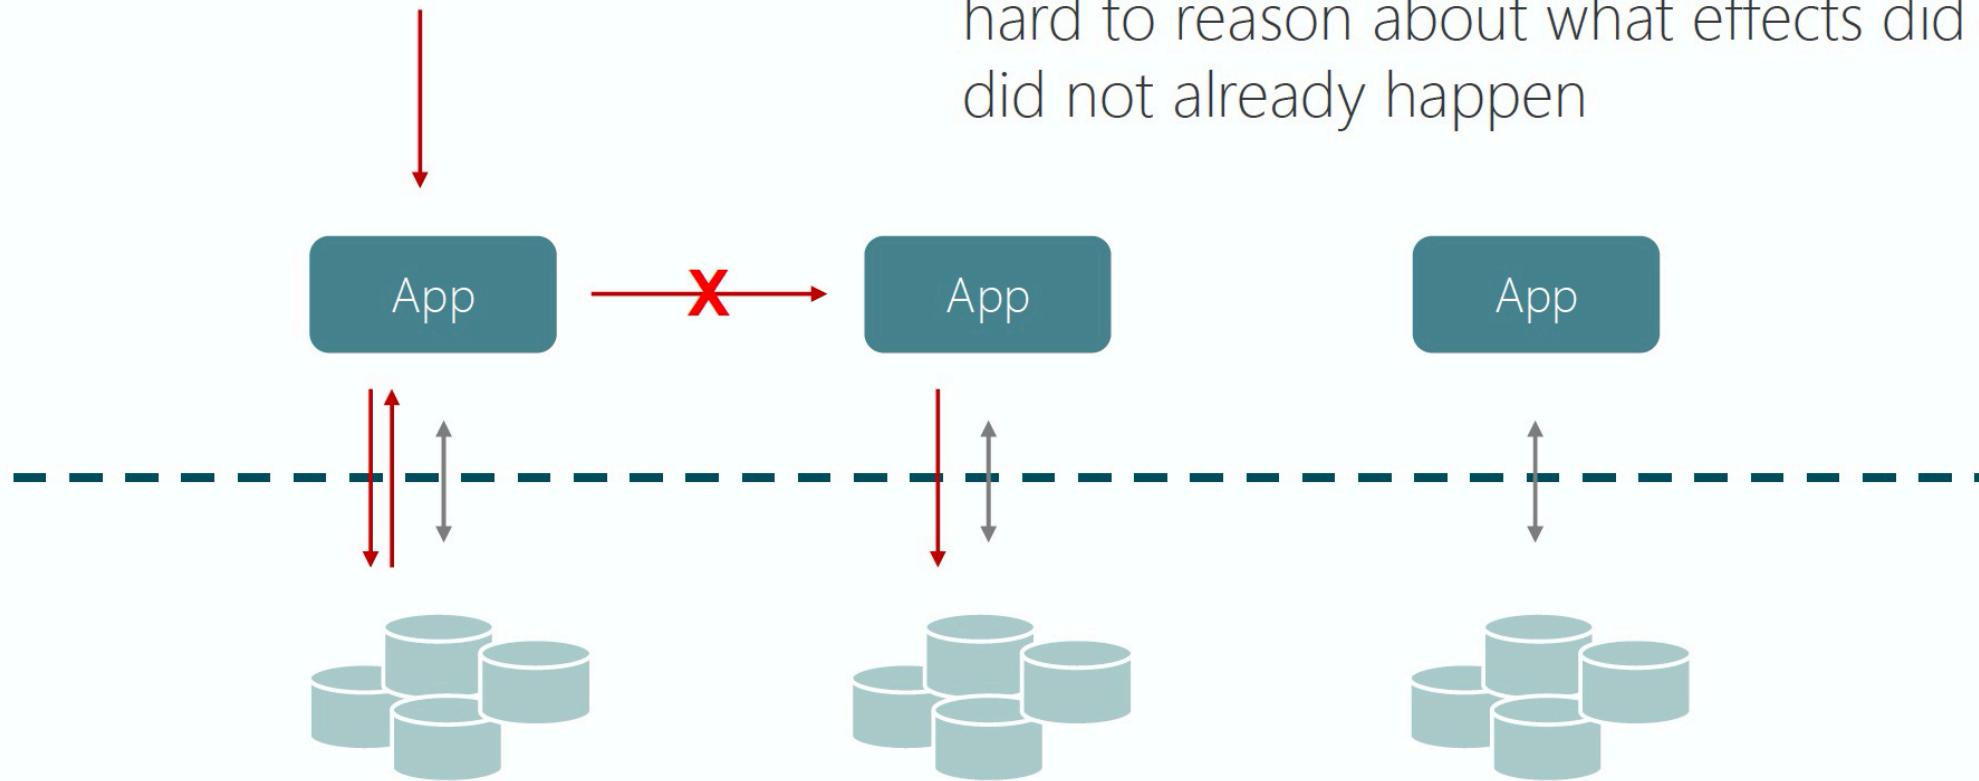


The Applications handles messaging (and consistency)

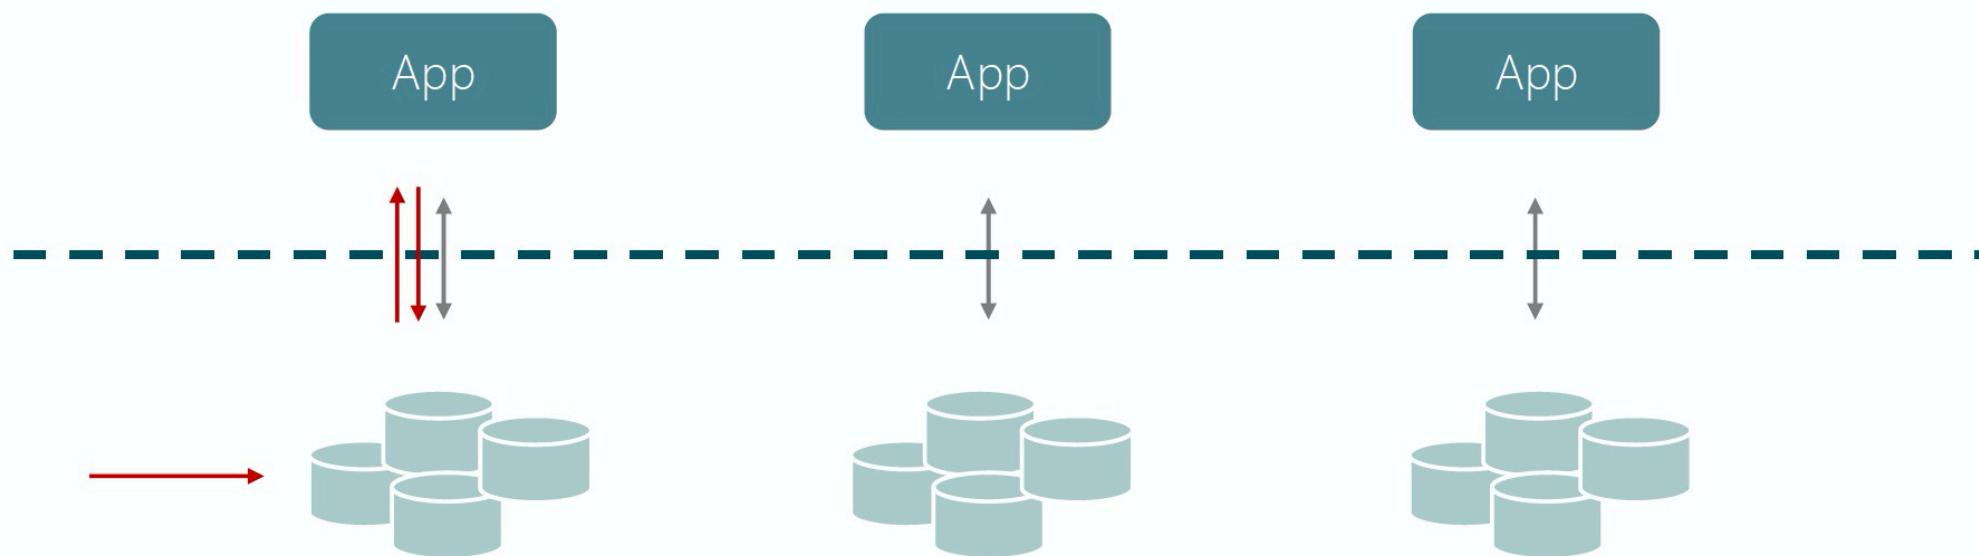


The Applications handles messaging (and consistency)

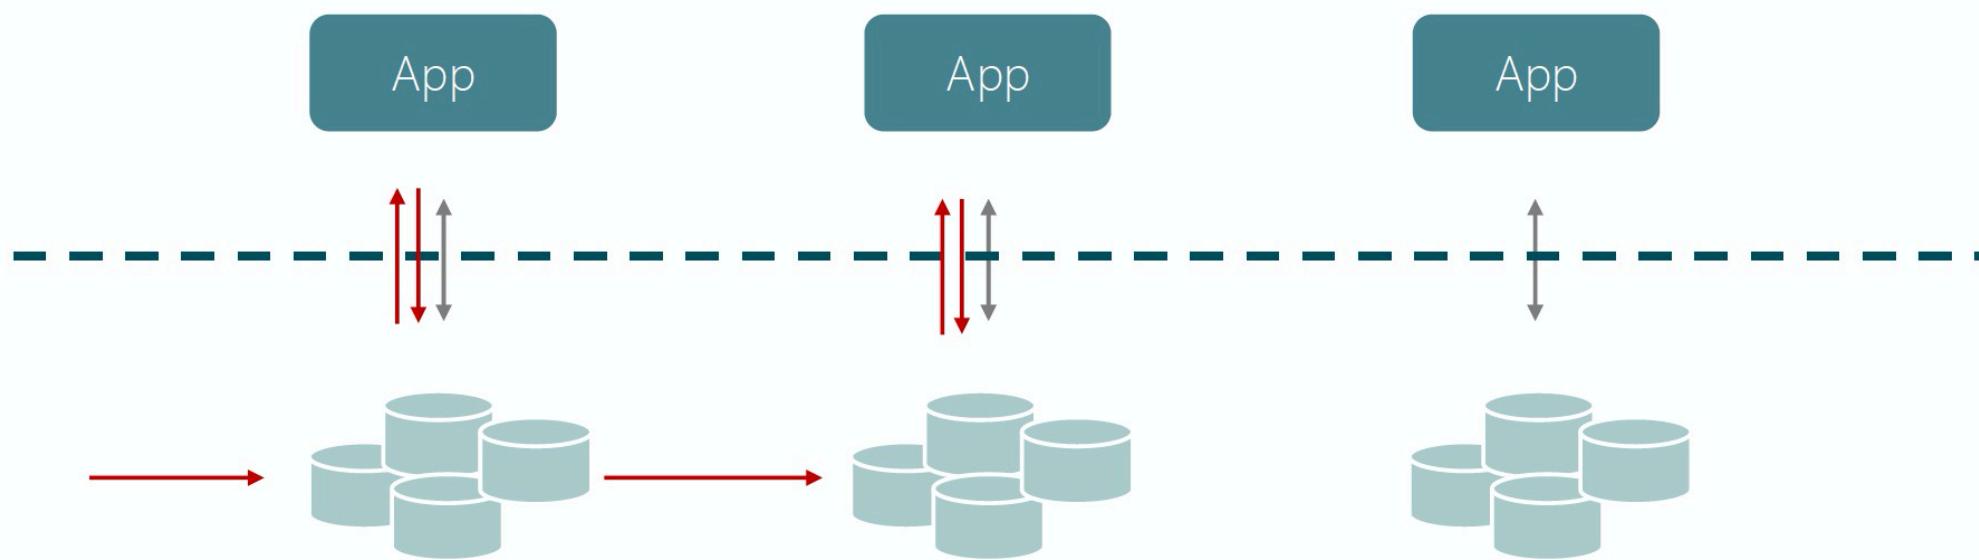
For any failure in any call, it becomes hard to reason about what effects did or did not already happen



The State Layer handles messaging (and consistency)

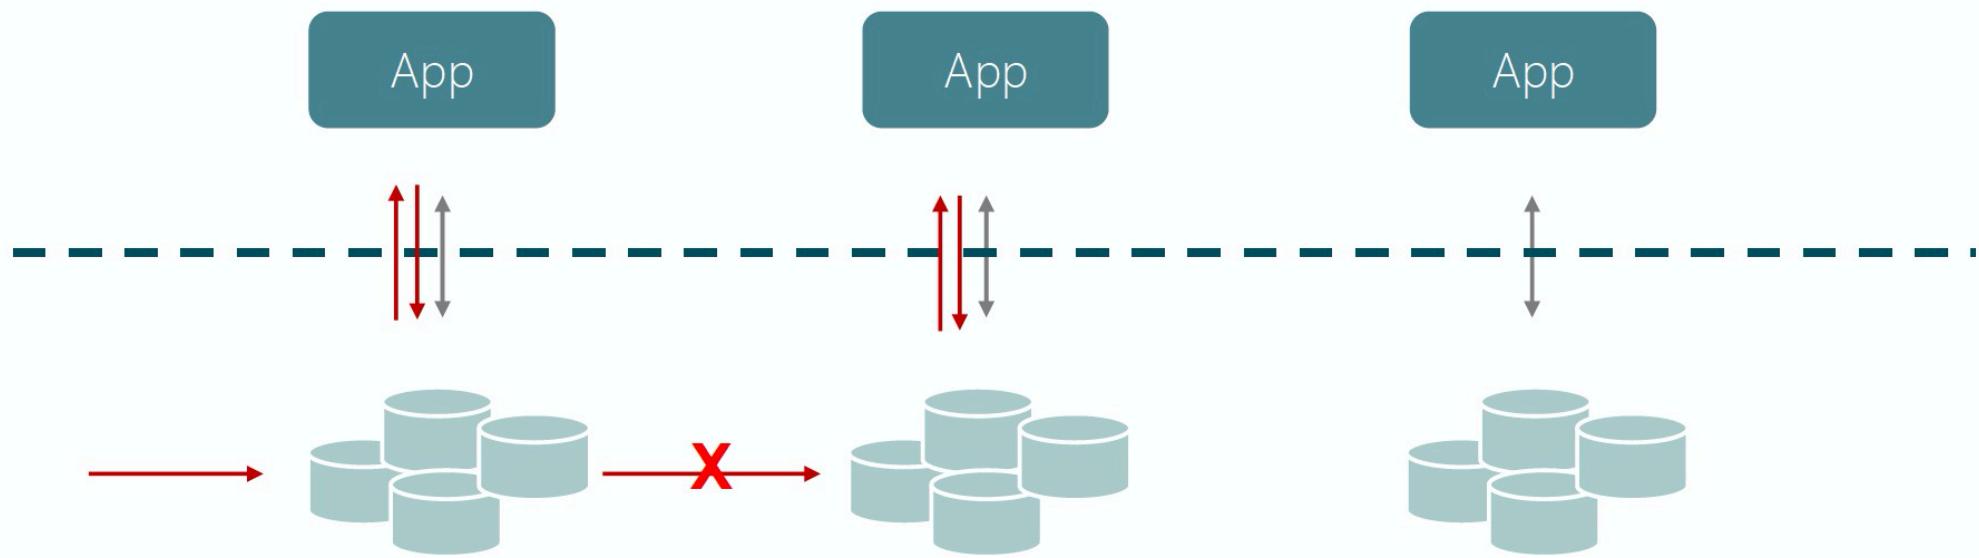


The State Layer handles messaging (and consistency)



The State Layer handles messaging (and consistency)

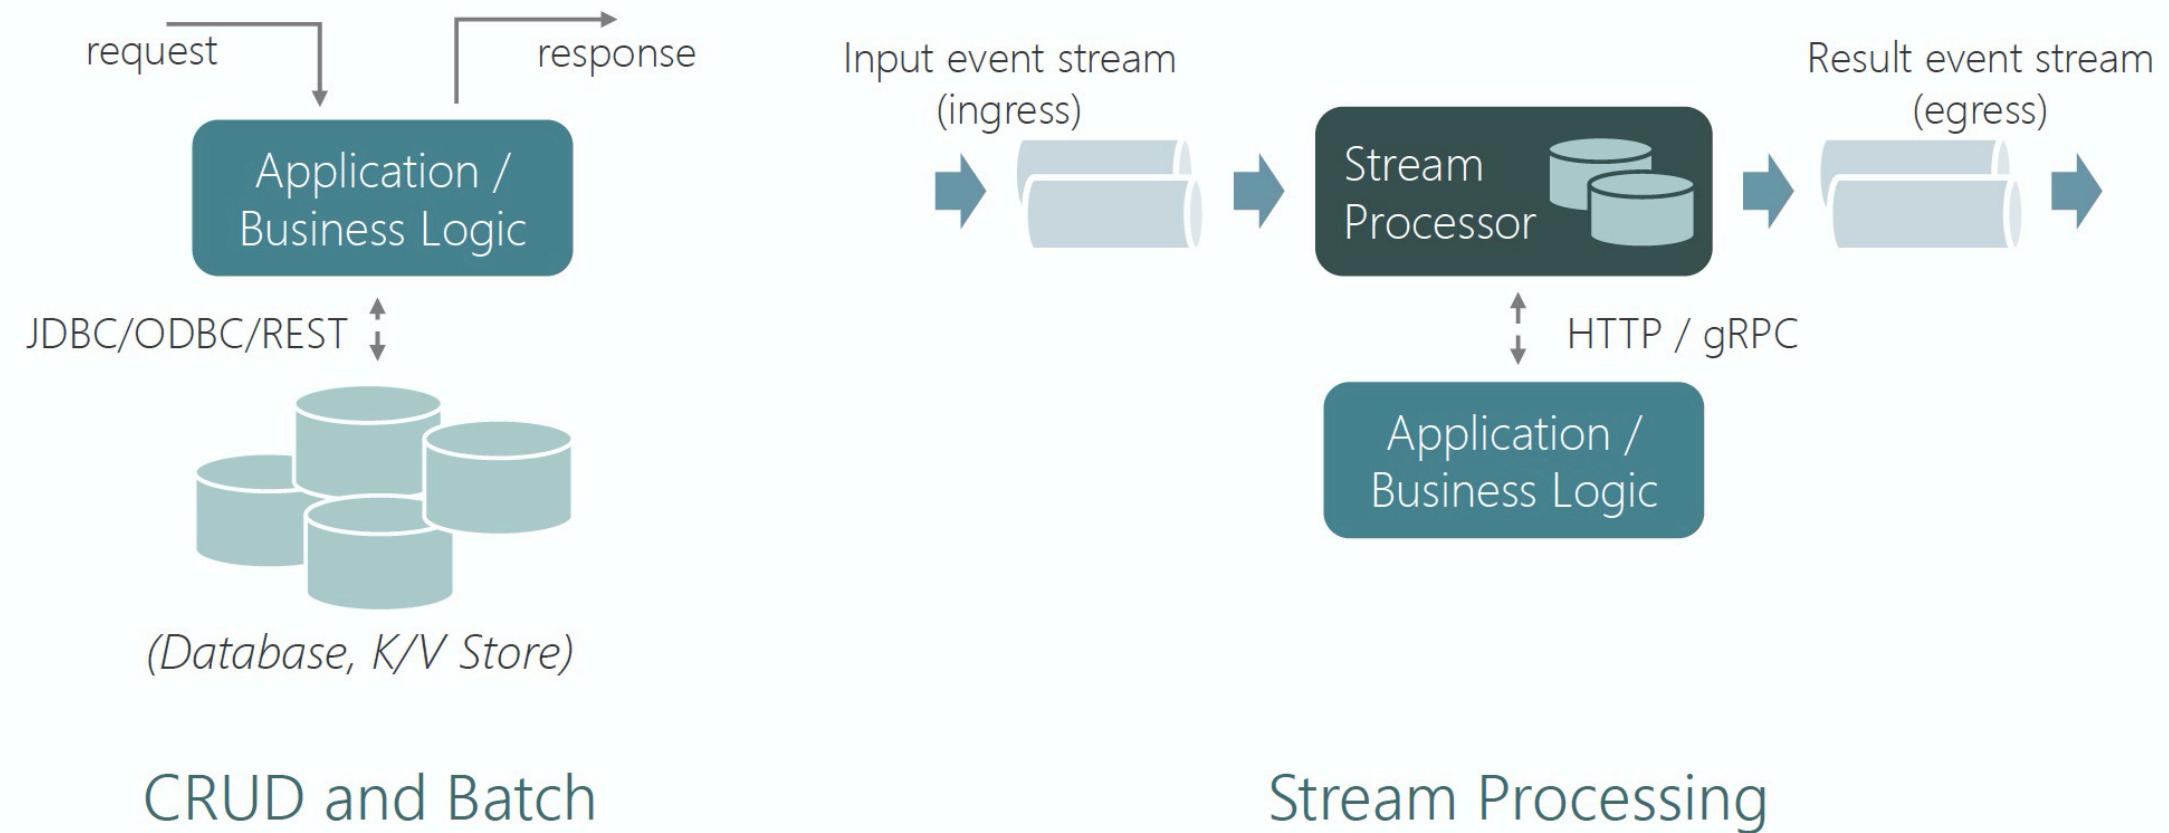
application logic is pure/stateless,
so re-tries are always idempotent



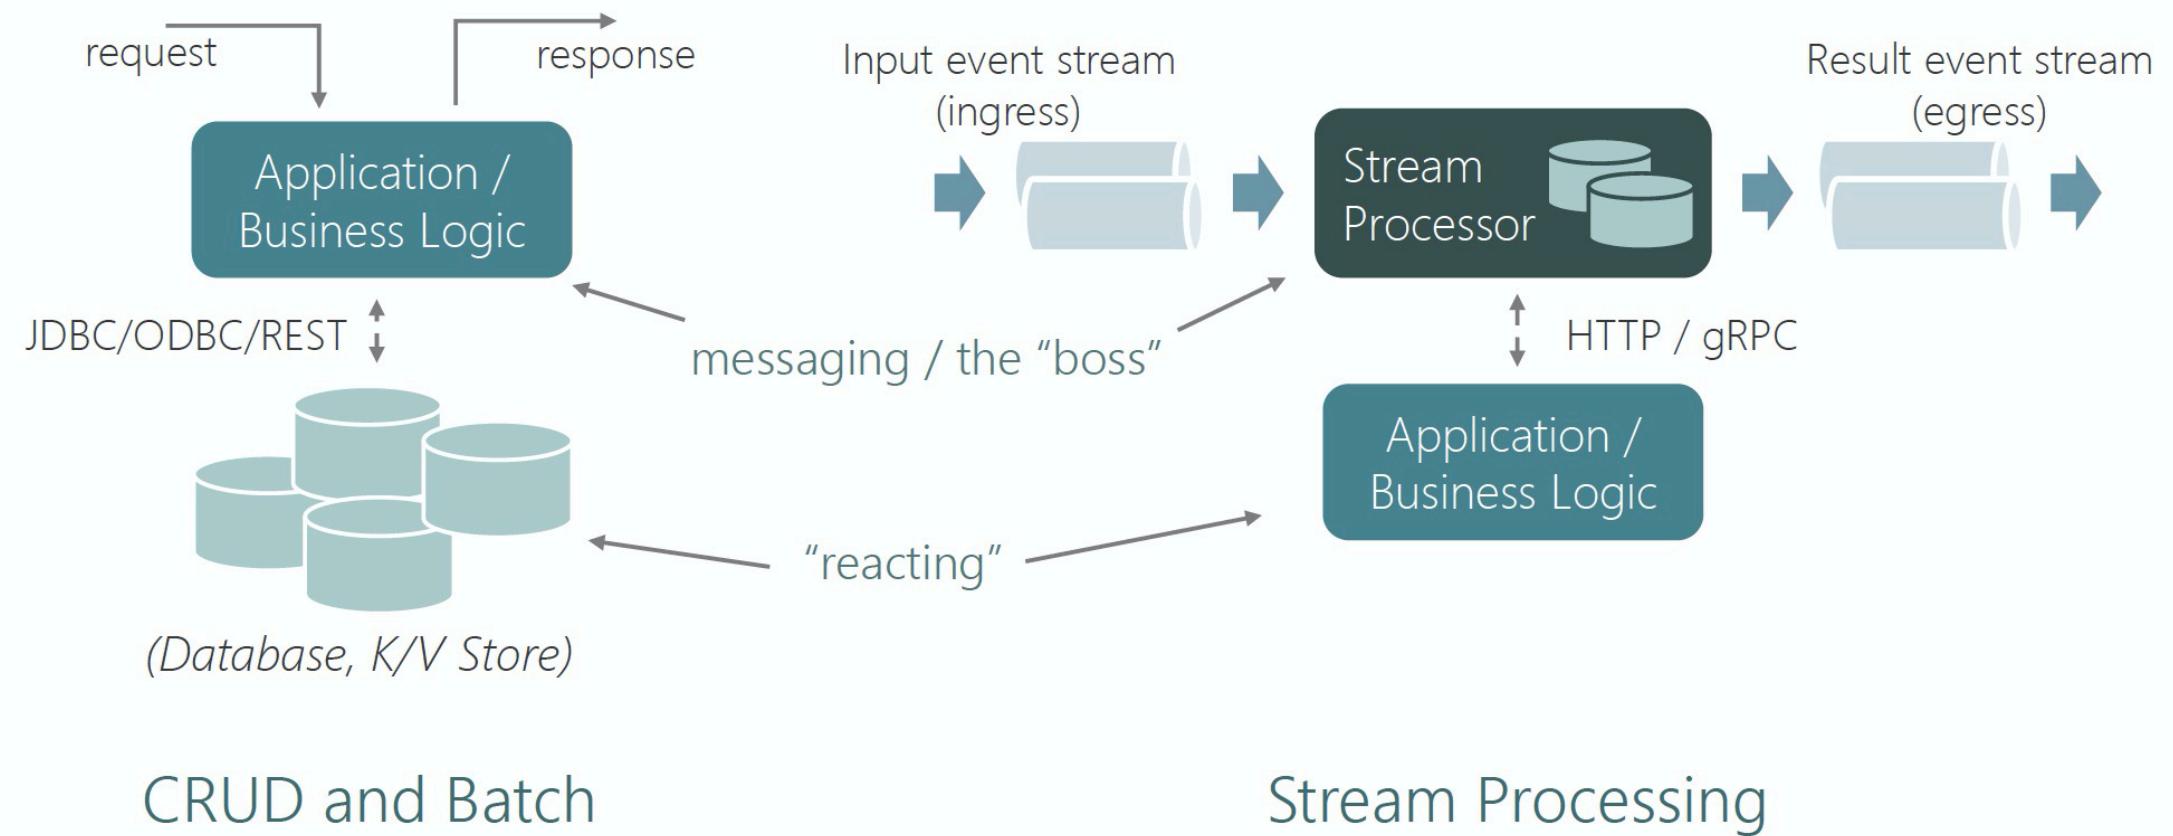
state updates and messaging are
easily made “atomic”



Compute / State Separation

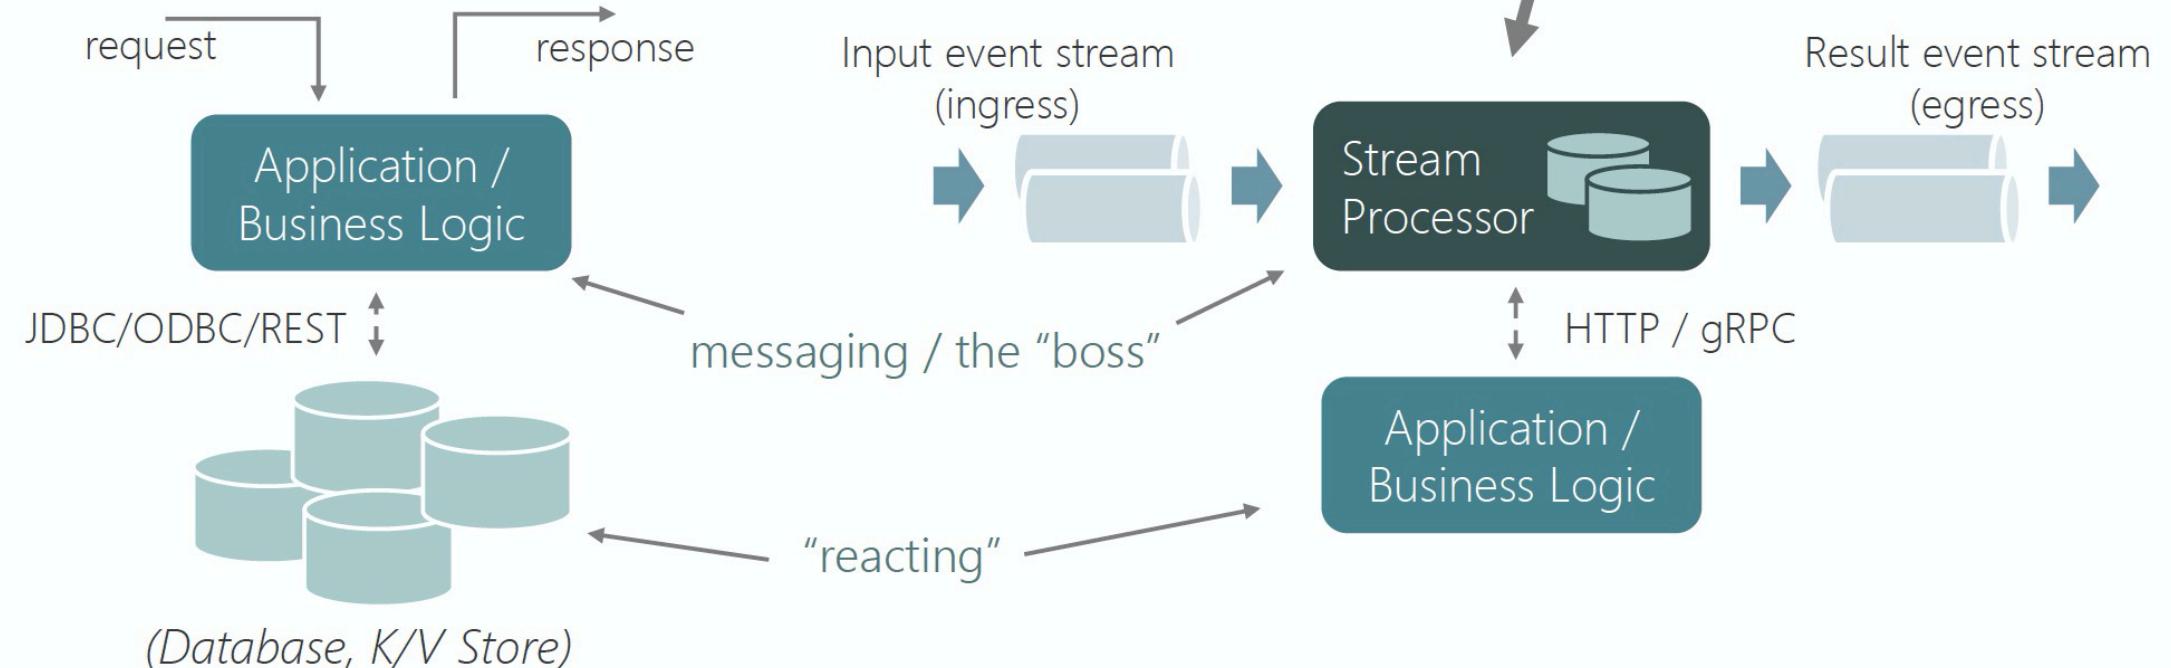


Compute / State Separation



Compute / State Separation

Stream Proc. + StateFun = Event-driven DB?

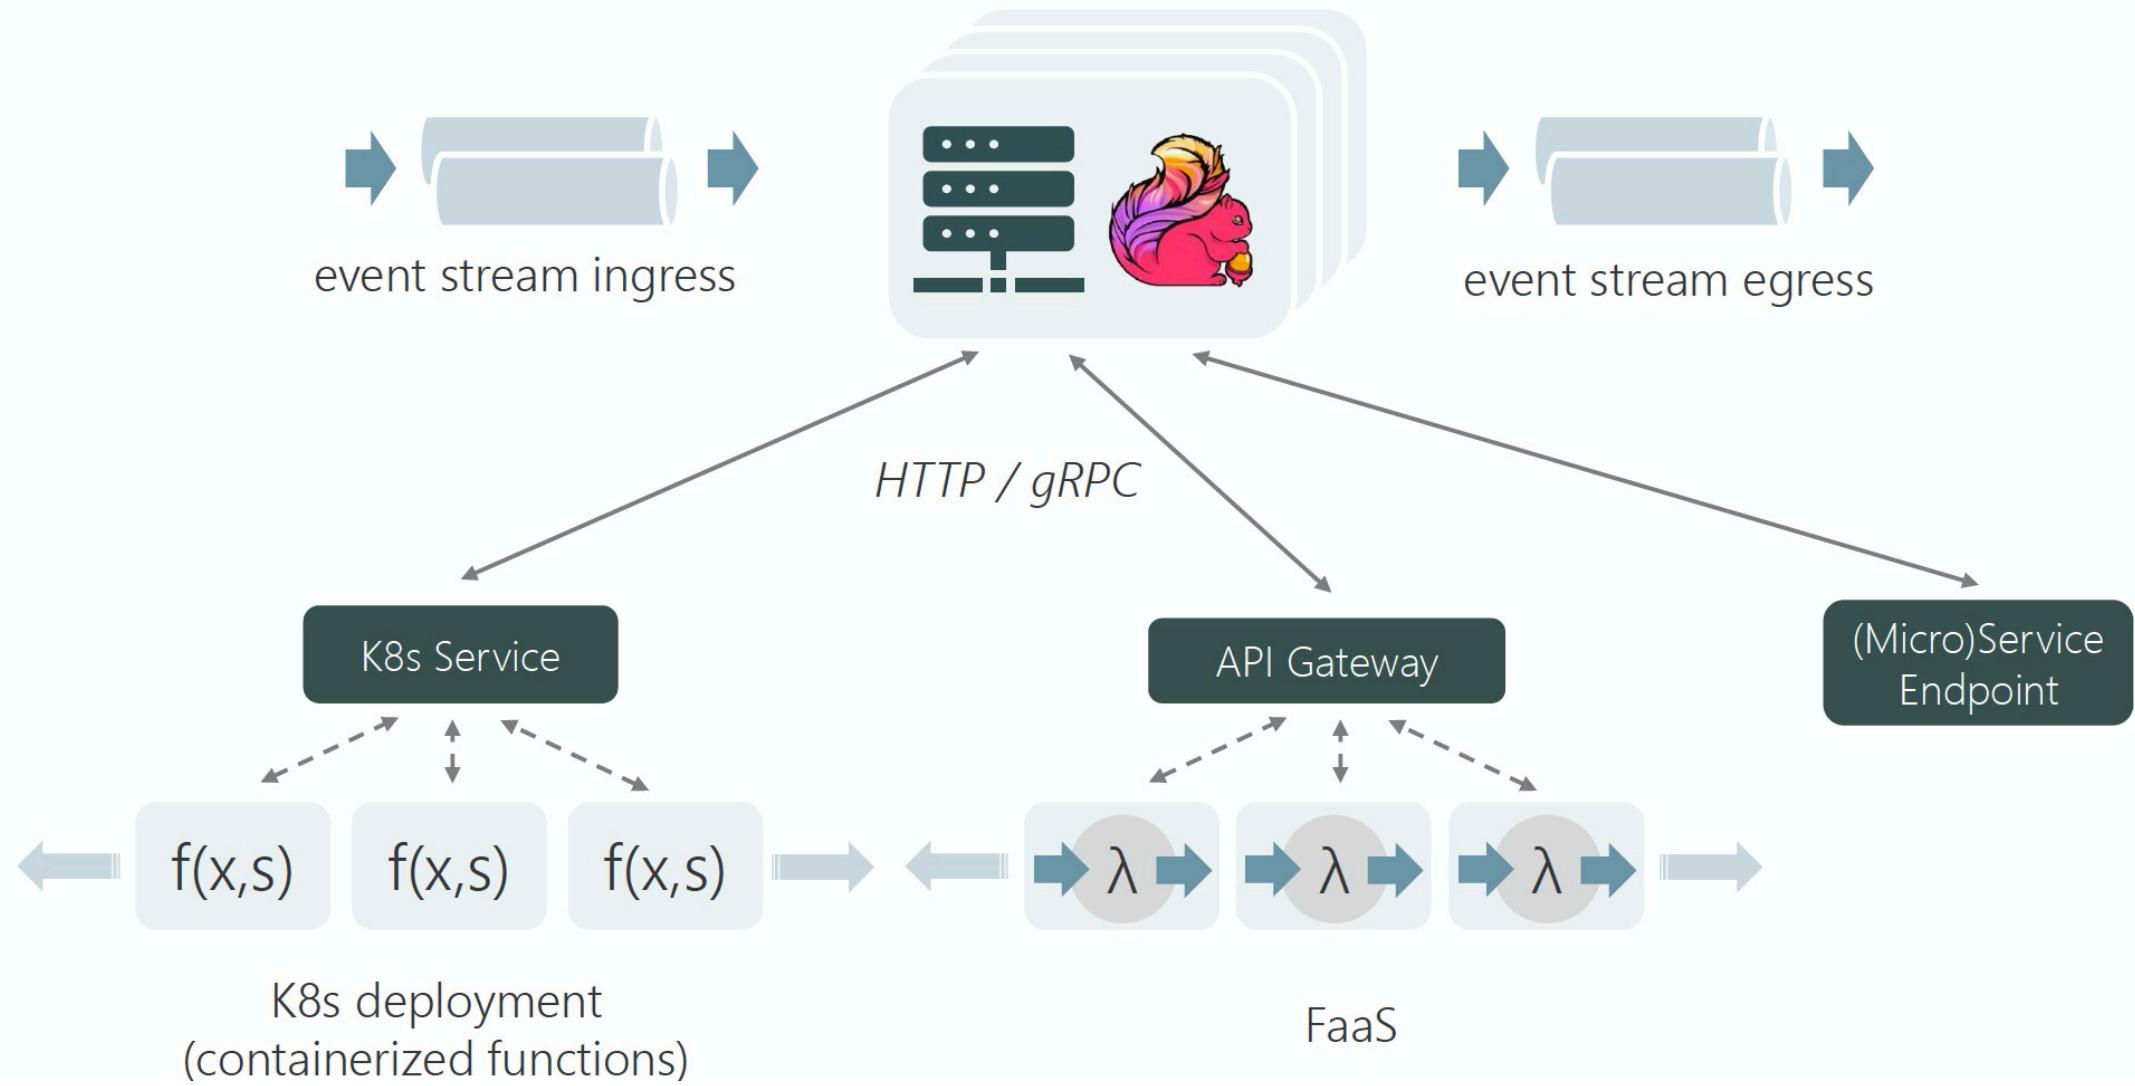


CRUD and Batch

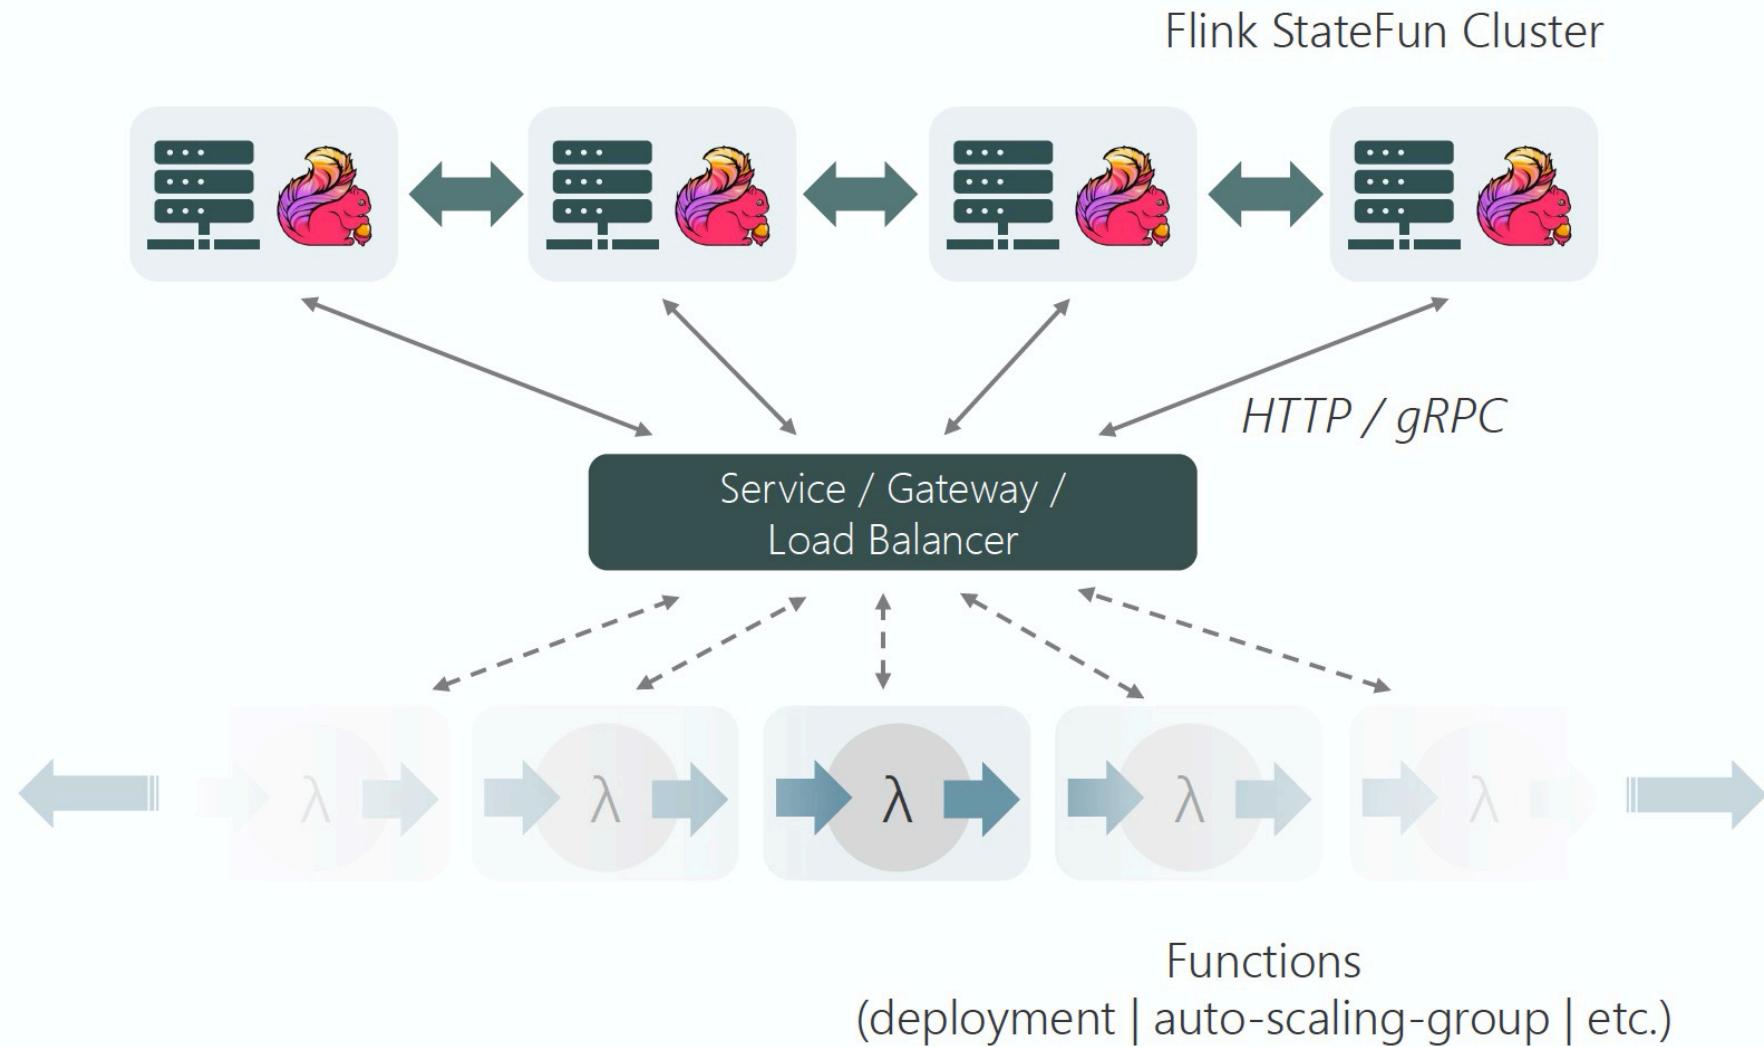
Stream Processing



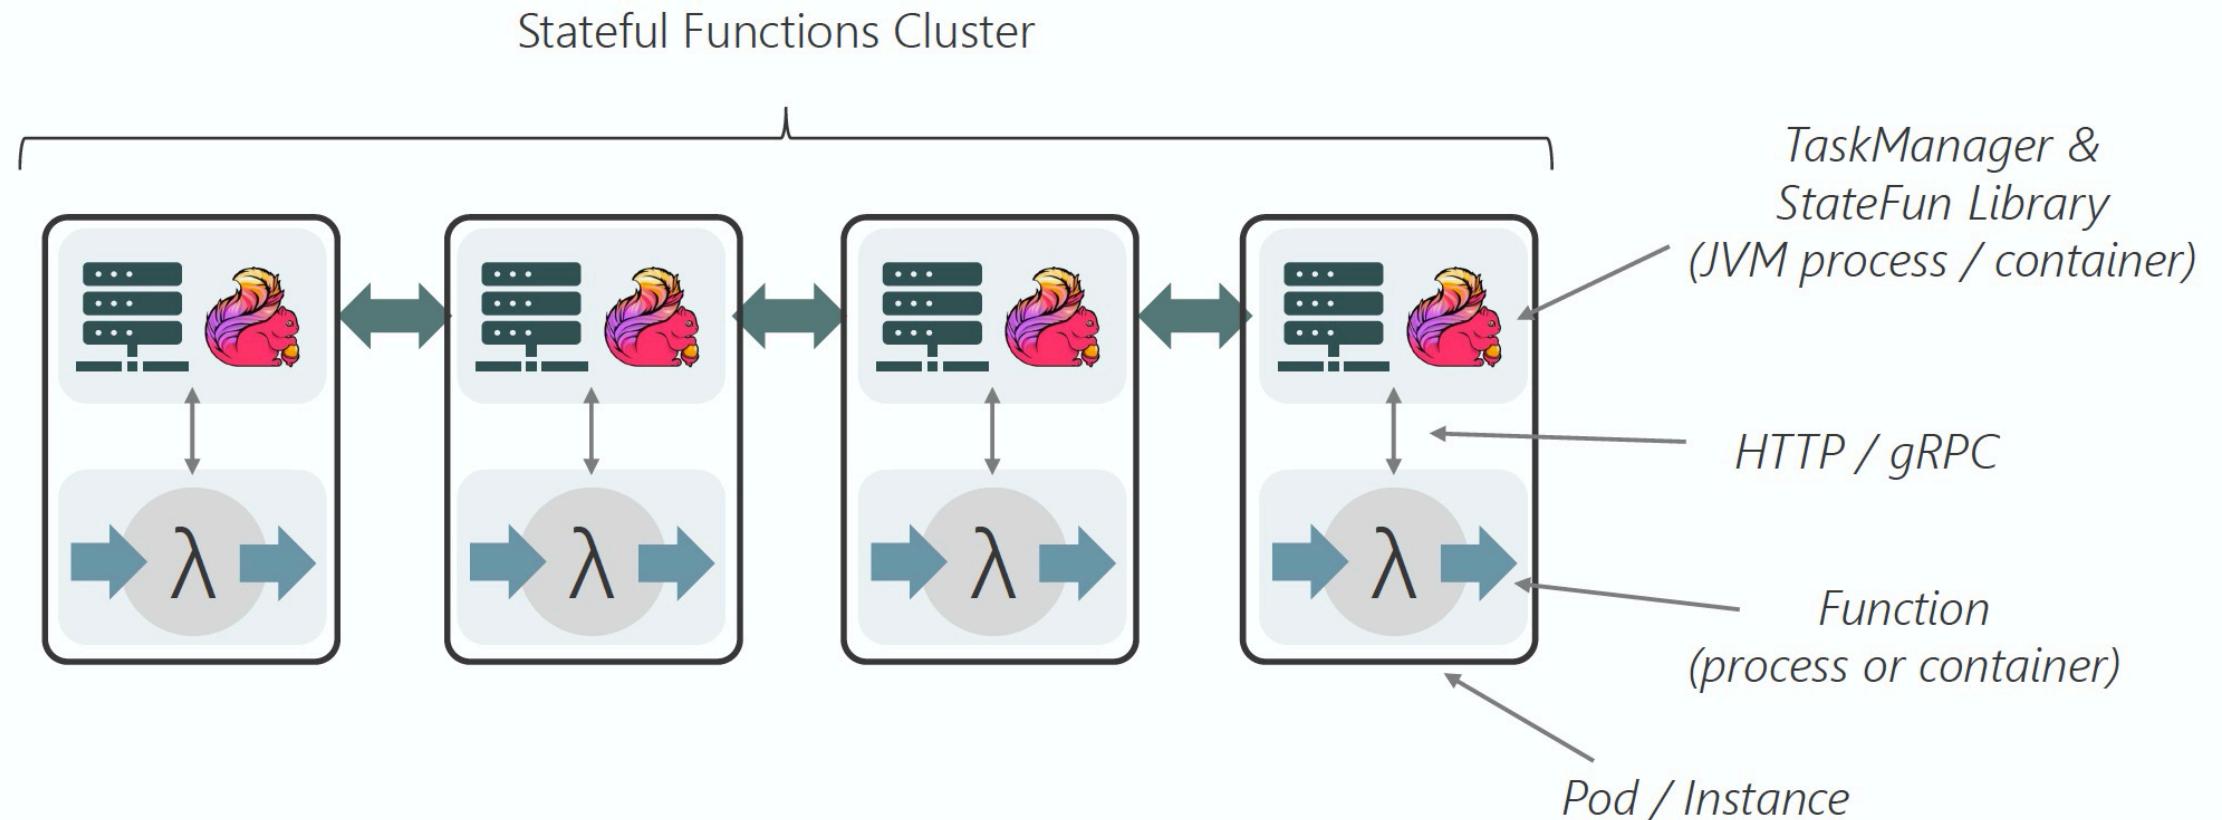
Stateful Functions 2.0



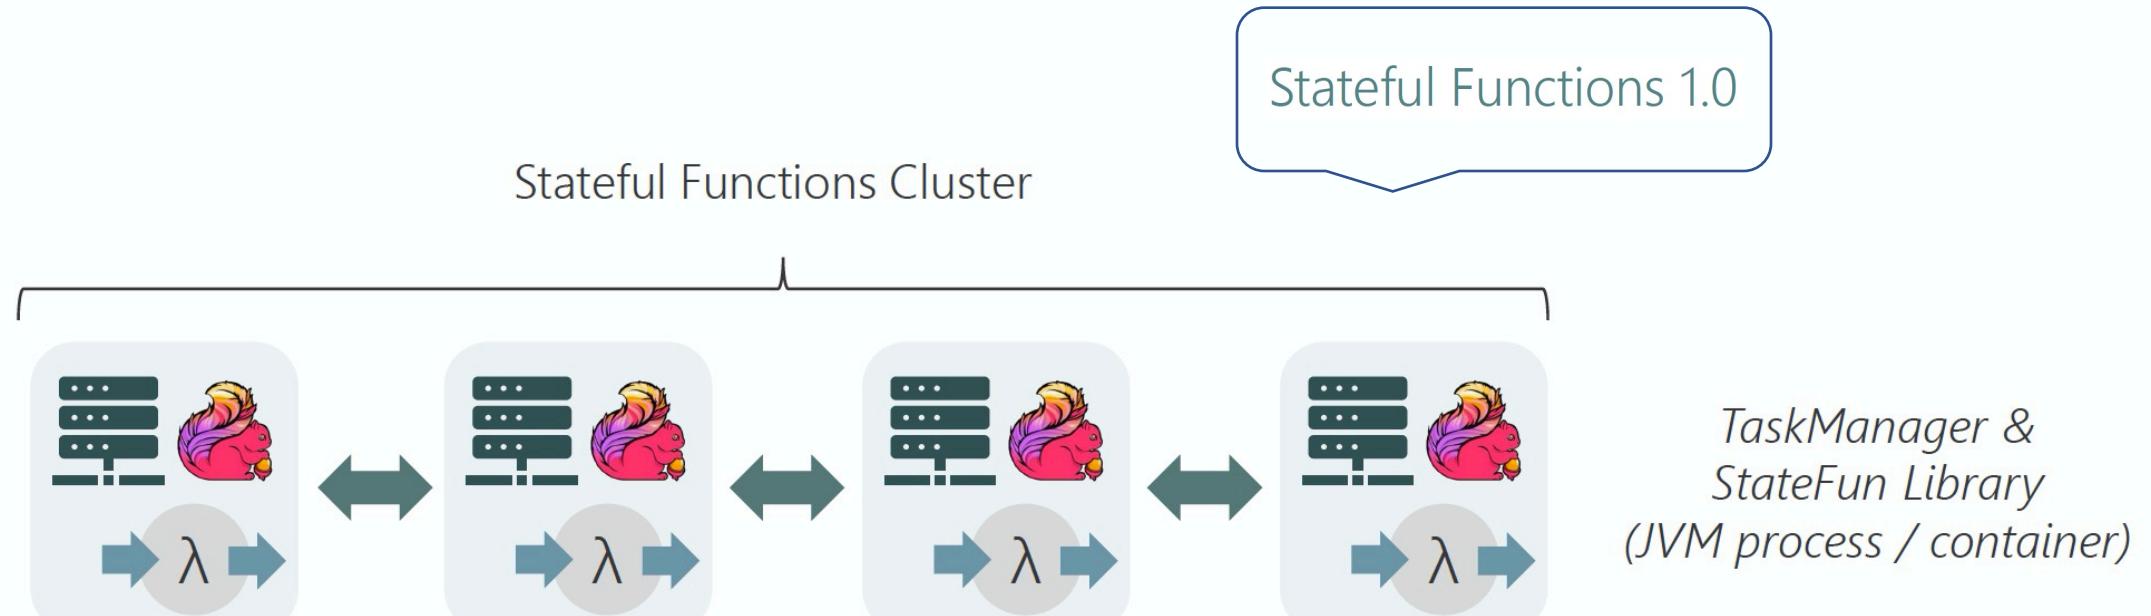
Remote Functions



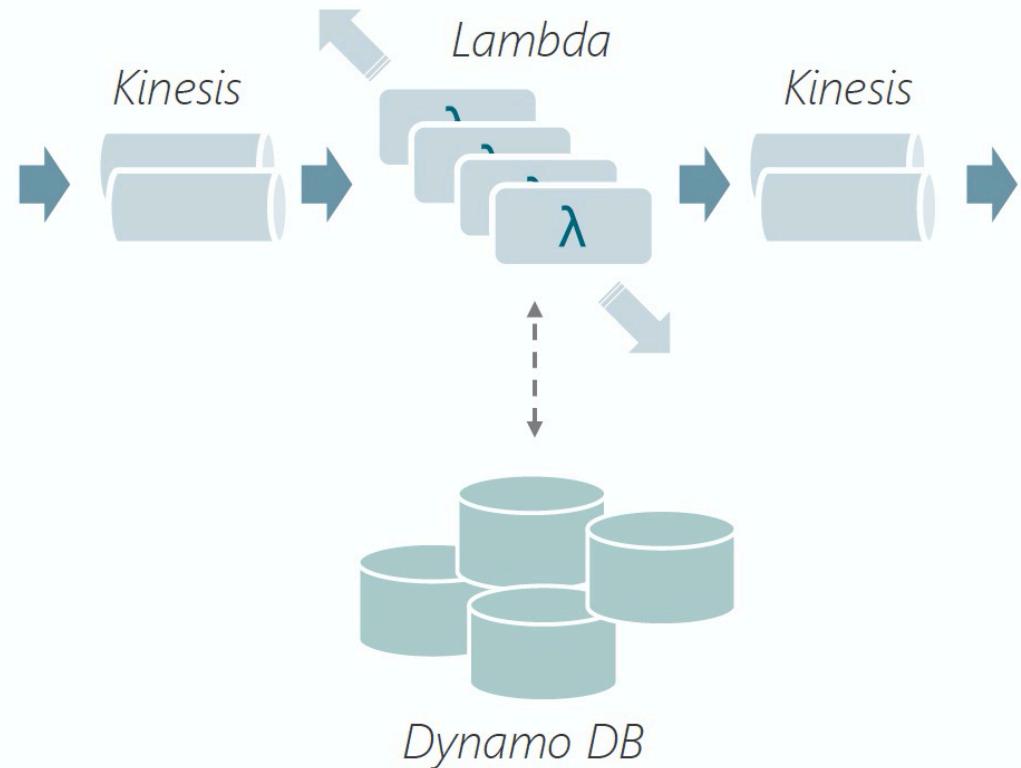
Co-located Functions



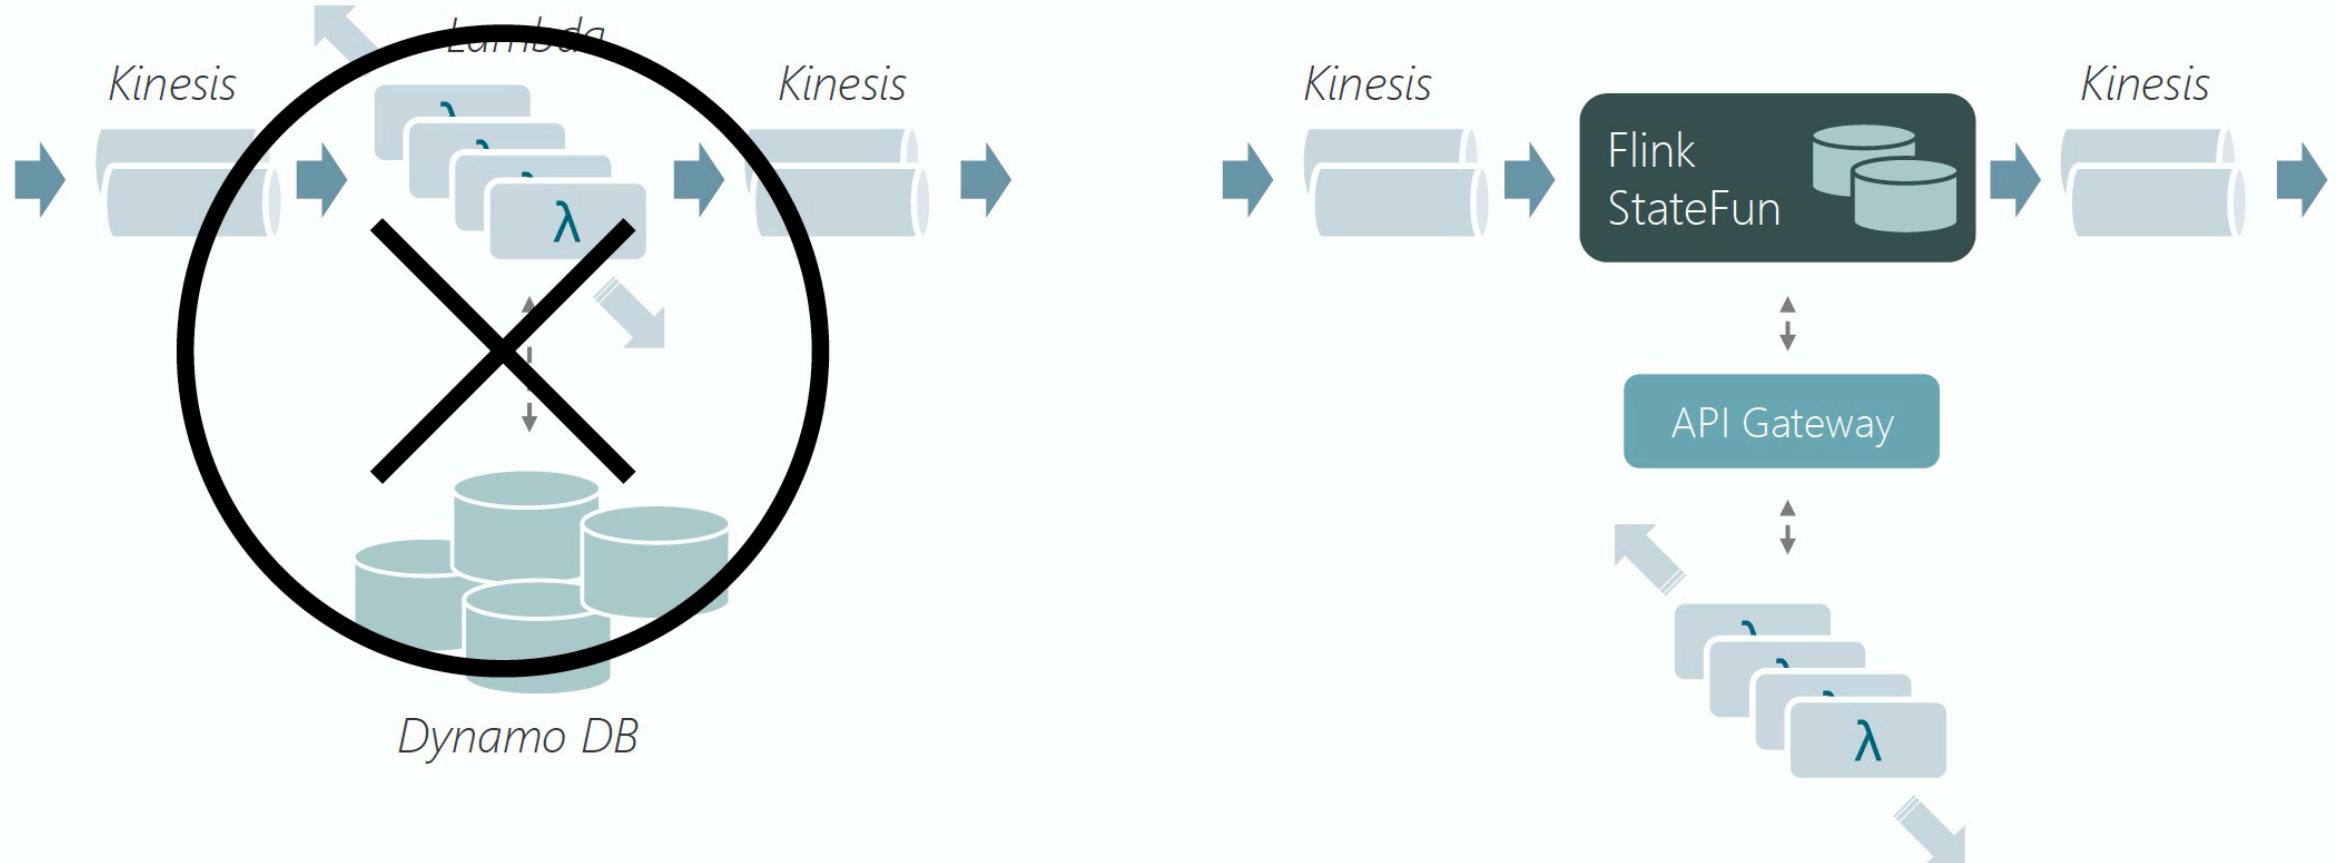
Embedded Functions



Remote Functions with FaaS



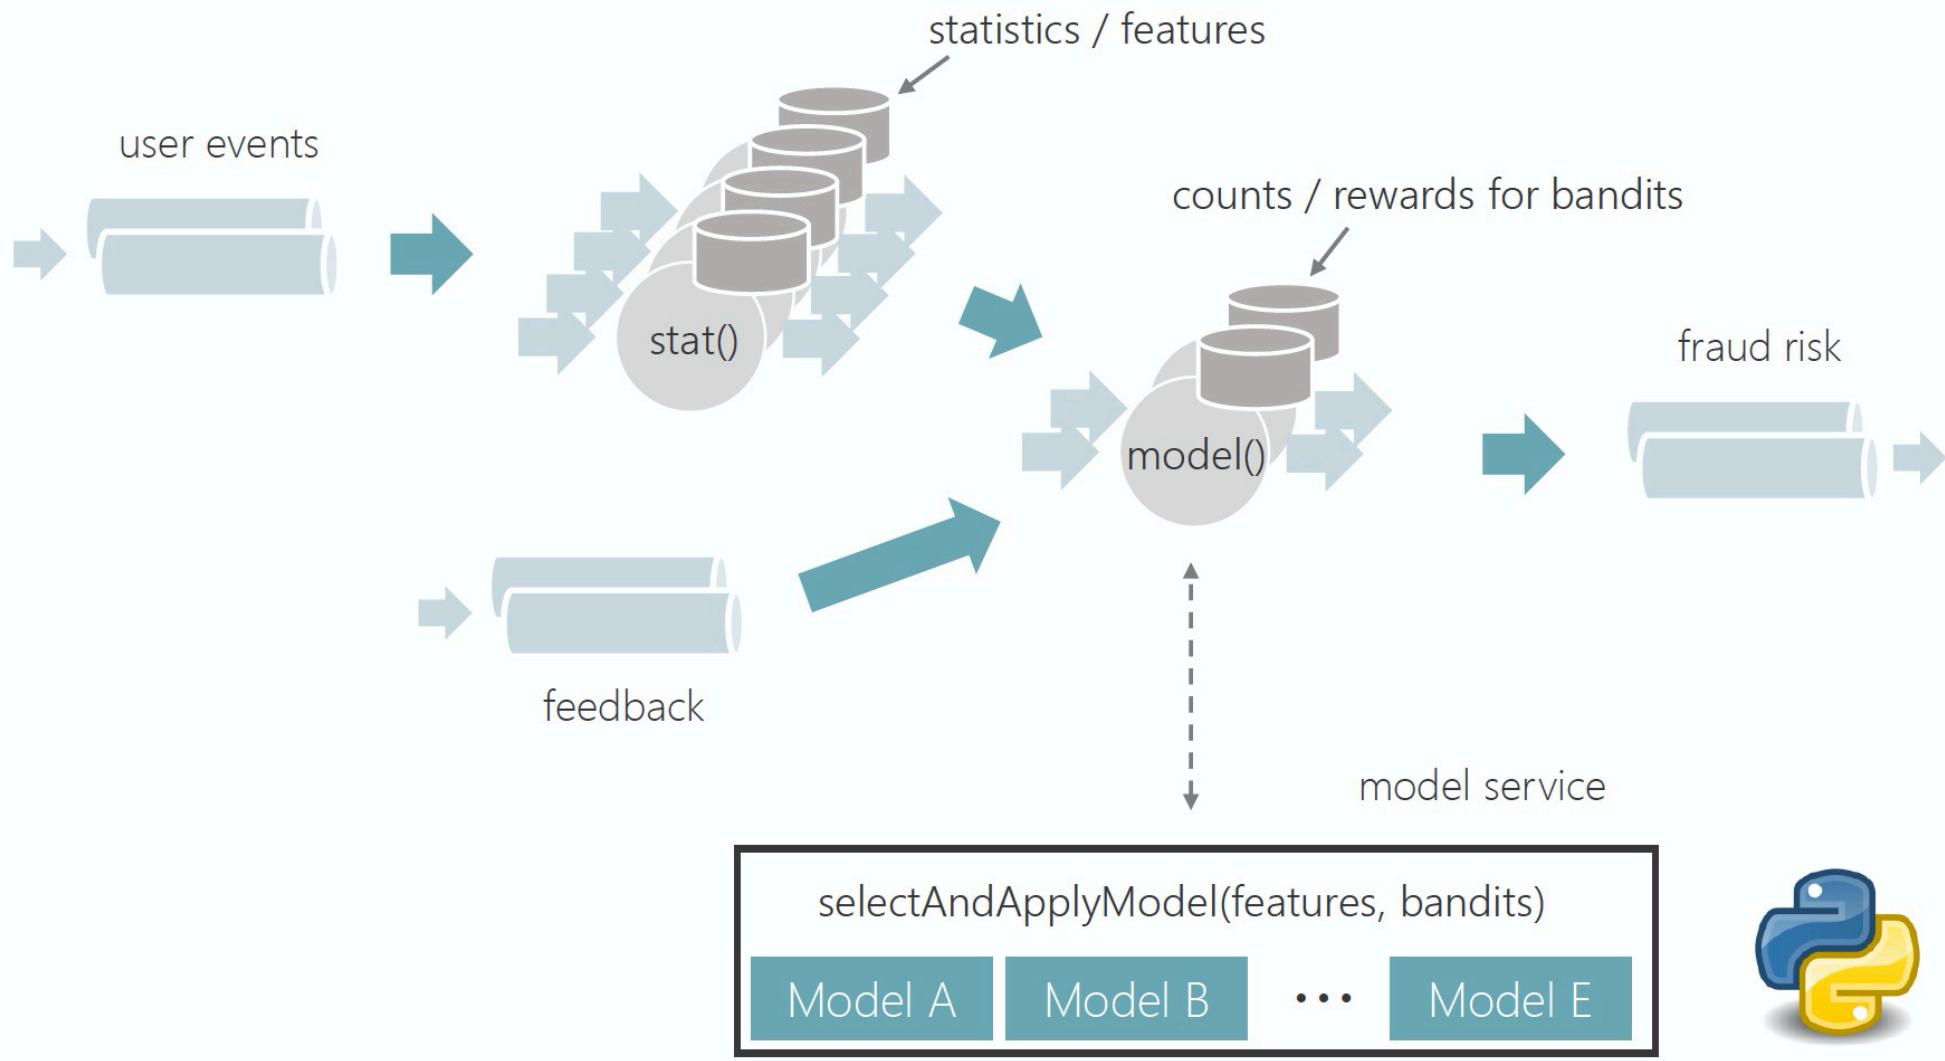
Remote Functions with FaaS



- Solves problem of consistent & scalable state
- Solves messaging / composition



Demo: Machine Learning - Feature Vectors & K-Bandits

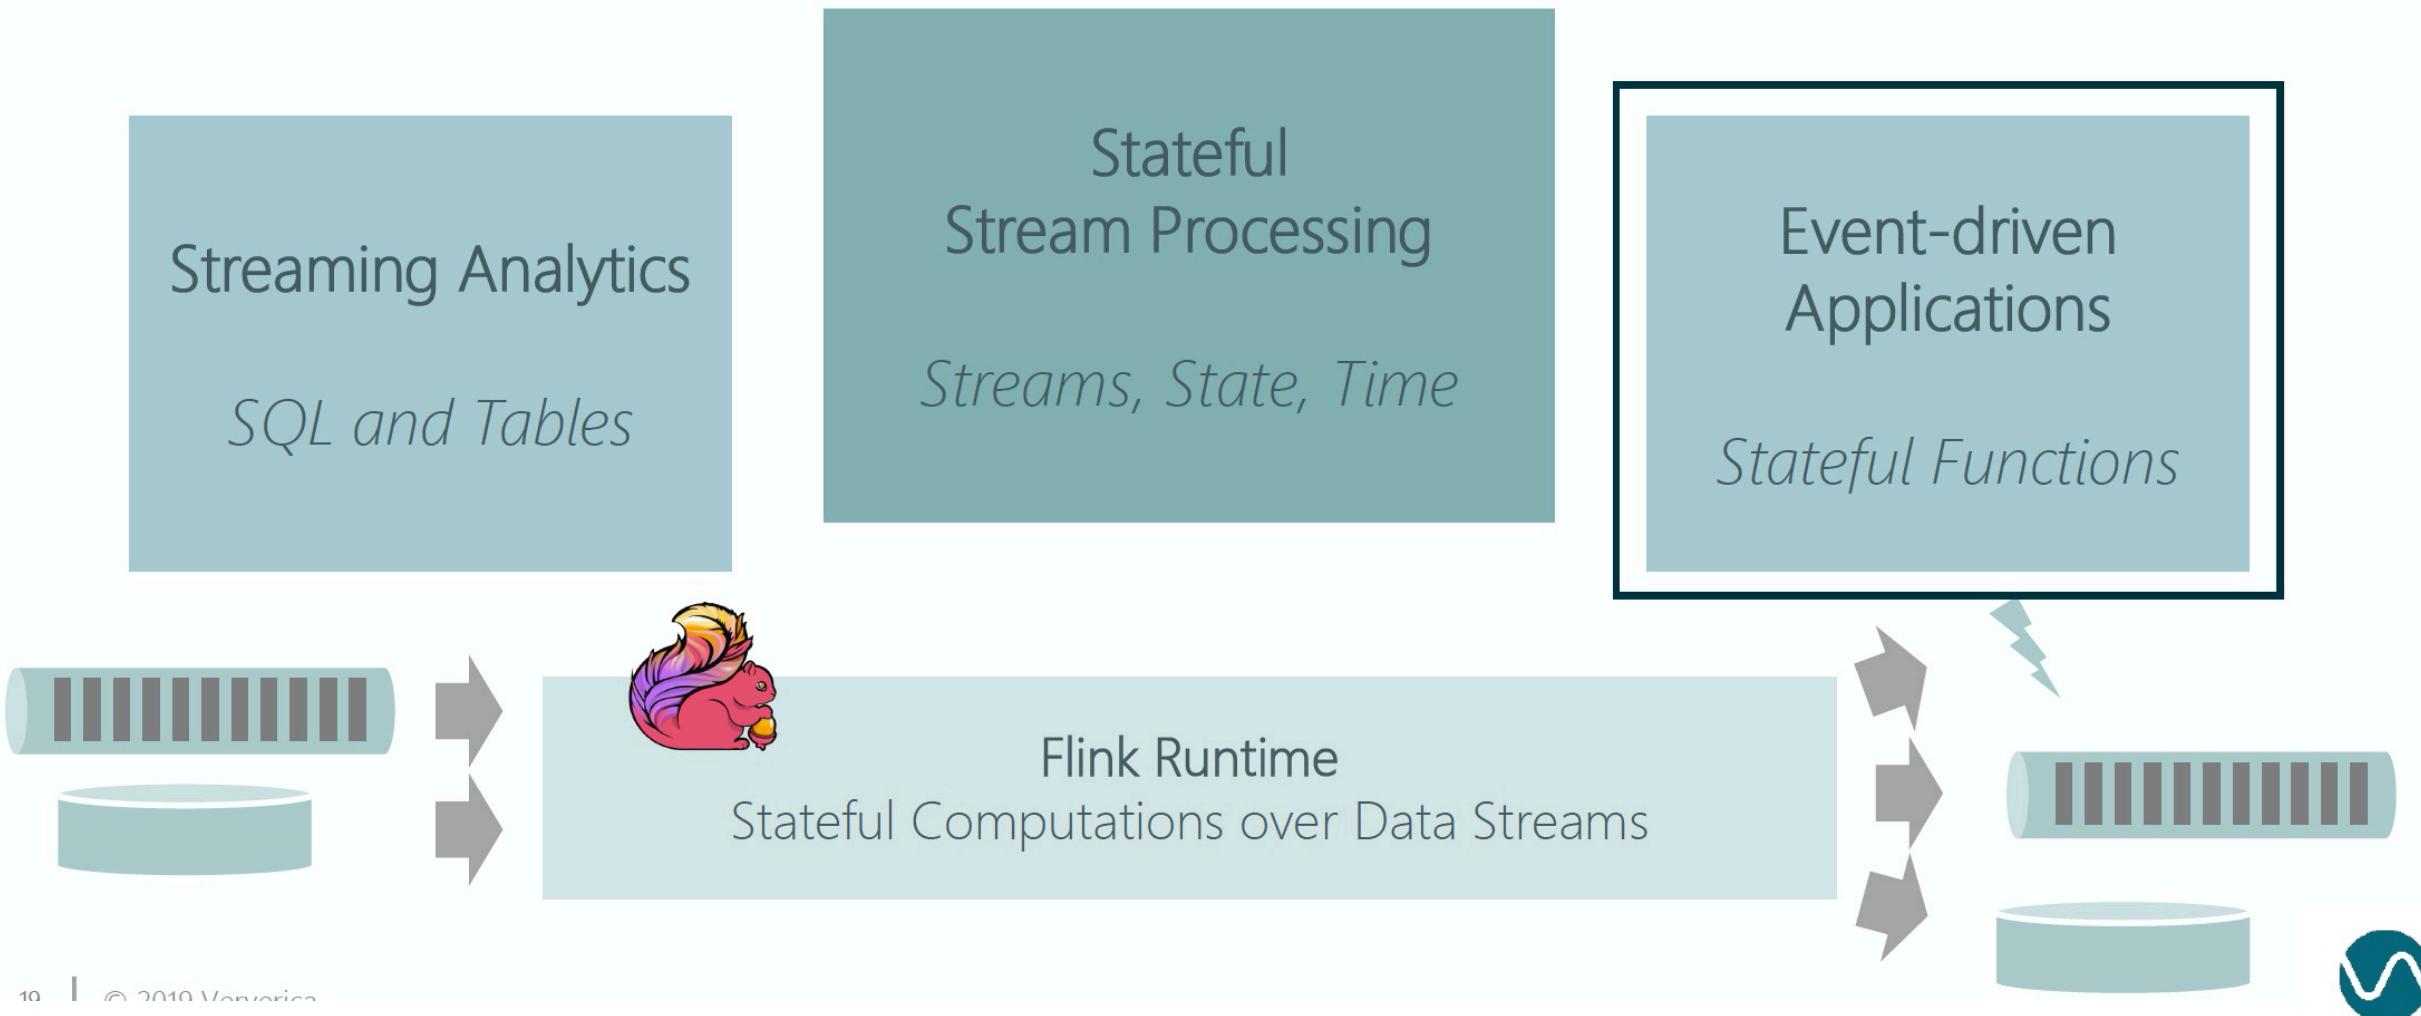


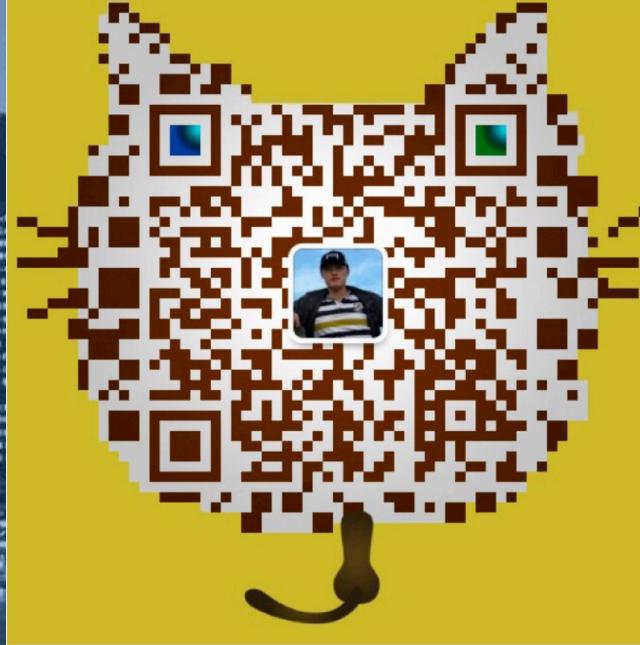
Demo: Machine Learning - Feature Vectors & K-Bandits

```
15 version: "1.0"
16 module:
17   meta:
18     type: remote
19   spec:
20     functions:
21       - function:
22         meta:
23           kind: http
24           type: demo/model
25     spec:
26       endpoint: http://python-worker:8000/statefun
27     states:
28       - arms
29     maxNumBatchRequests: 500
30     timeout: 2min
```

```
43   @functions.bind("demo/model")
44   def model(context, message):
45     if message.Is(SparseVector.DESCRIPTOR):
46       handle_prediction_request(context, message)
47     elif message.Is(PredictionFeedback.DESCRIPTOR):
48       handle_prediction_feedback(context, message)
49     else:
50       raise ValueError('unknown message type ' + message)
51
52
53   def handle_prediction_request(context, message):
54     # select the best model so far for this user
55     bandit = get_from_state(context)
56     model_index = bandit.choose_arm()
57
58     # use the selected model to make a prediction
59     result = PredictionResult()
60     result.from_model = model_index
61     result.prediction = MODELS[model_index].predict(message)
62
63     # emit the prediction result to a Kafka topic
64     out_record = kafka_egress_record(topic="predictions", value=result)
65     context.pack_and_send_egress("demo/kafka", out_record)
```

StateFun means to open Flink to Event-driven App Use-cases





We hope you enjoy the first Virtual Flink Forward.